Ferrule: Bounded Storage-Backed Weight Residency

Lu Huanbing
Independent Researcher
luhuanbing084@gmail.com

Working draft

Abstract

Storage-backed weight residency requires more than a small cache: reads, transfers, and device consumers can retain resources after a request or binding is logically finished. We present Ferrule, an architecture that connects source-bound materialization plans, stage-local admission, exact completion identity, and ownership through retirement. Its implemented expert path combines a persistent bounded read worker, optional next-expert lookahead, and grouped synchronous uploads without changing the synchronous forward interface. We evaluate a single-device, text-only Qwen3.5-35B-A3B-FP8 profile using BF16 execution weights and explicit F32 boundaries. Separate public-CLI acceptance checks complete 48 requests; a transfer-policy pilot completes 96 requests across eight fresh servers, including 40 measured, output-checked requests. Grouped uploads improve observed client latency and phase-qualified goodput with lookahead, but increase time to first content with blocking reads while end-to-end latency remains nearly unchanged. With two pairs per read schedule, these are conditional, descriptive benefits rather than a universal performance result. An earlier five-pair study separately isolates the effect of read scheduling. Quality gates, selected-process profiles, component faults, and HTTP cancellation provide complementary evidence about materialization and custody. The measurements bound selected logical host and cache accounts, not physical RSS, storage traffic, or universal device release; integrated five-class conservation and end-to-end fault coverage remain unverified.

1 Introduction

Storage-backed inference repeatedly brings selected weights into an execution cache rather than keeping every execution weight resident. This approach can reduce device residency, but a small cache alone does not bound the resources used to fill it. A read may retain host buffers, an upload may outlive its caller, and an evicted binding may still have device consumers. Releasing a logical request therefore does not necessarily make its capacity reusable. The central problem is to coordinate weight materialization, resource admission, and ownership through the last consumer.

Weight streaming and prefetch are established techniques in systems such as FlexGen, ZeRO-Inference, and PowerInfer-2 [3, 6, 5]. SGLang also provides a specialized storage-backed expert path in the inspected source [2]. Ferrule addresses a complementary architectural question: how should a materialization schedule bind source identity, stage-local accounting, and completion evidence so that capacity is not reused prematurely? We distinguish three properties that are often conflated: a configured resource budget, a measured physical working set, and evidence that a particular operation has finished using its resources. Neither storage offload itself nor performance leadership is the claimed novelty.

The design follows *plan, admit, materialize, publish, execute, retire.* A frozen plan identifies the source and destination generations. Admission acquires the claims required by each stage.

Completion validates the exact operation before publication, while the owner retains buffers and charges until the relevant consumers finish. If completion cannot be proved, the intended discipline retains uncertain resources instead of treating an error or cancellation as permission to reuse them. The resulting conservation model is conditional on complete ownership and transfer rules; the serving evidence observes selected host and expert-cache accounts, not every class in an integrated system-wide ledger.

The implemented expert path separates read scheduling from transfer policy. A persistent worker serves bounded checkpoint-read slots. The blocking schedule waits at demand; lookahead submits one next-expert read before current expert computation. Independently, grouped synchronous transfer batches the expert's weight and scale copies behind one upload completion boundary. Both choices preserve the expert order, arithmetic, and synchronous forward interface. Worker lifetime, payload lifetime, upload completion, and device-consumer retirement remain distinct: an idle live worker can still own charged metadata and stack capacity.

Contributions.

1. An architecture for bounded storage-backed weight residency that connects source identity, stage-local admission, publication, and retirement, while distinguishing logical accounting from physical memory and producer completion from consumer quiescence.

2. Reference algorithms for exact materialization and bounded cleanup, together with an implemented persistent read worker, bounded next-expert scheduling, and grouped synchronous expert uploads.

3. A profile-scoped evaluation linking numerical and HTTP output checks, logical work and resource accounts, client timing, selected-process profiling, and component-failure and cancellation behavior, including negative results.

Research questions.

RQ1: Bounded completion. Can requests complete under the declared accounting policies, and how does residency change across coupled cache budgets?

RQ2: I/O demand. What work reaches the logical reader and shared filesystem client, and which physical traffic remains unknown?

RQ3: Scheduling and transfer. How do read scheduling and upload grouping affect observed waiting, synchronization, and client outcomes?

RQ4: Failure and custody. What resources and state remain after clean abort, uncertain completion, and HTTP cancellation?

The core serving workload uses Qwen3.5-35B-A3B-FP8 through the public CLI, with BF16 execution weights and explicitly defined F32 boundaries, not a full-model W8A16 quality claim. Four acceptance servers complete 48 checked requests. A separate grouped-versus-synchronous transfer pilot completes eight fresh server lifecycles and 96 requests, including 40 measured eight-token quality checks. With lookahead, grouped uploads reduce observed client latency and increase phase-qualified goodput; with blocking reads, time to first content regresses despite nearly unchanged end-to-end latency. There are only two pairs per read schedule. An earlier five-pair read-scheduling study instead compares blocking with lookahead under its own fixed implementation, finding less demand

waiting and lower observed client latency with equal logical reads. These are separate comparisons, not a pooled population or evidence of a universal improvement.

The evaluation establishes neither a physical capacity frontier nor a matched external-system ranking. DeepSeek-V4 end-to-end execution is untested; architecture-level applicability does not qualify another model adapter. Table 5 consolidates the measurement boundaries. Section 2 defines the model and assumptions. Section 3 develops the mechanisms and algorithms, and Section 4 maps them to the implemented path. Section 5 fixes the methodology before Section 6 answers the four questions. Related work, limitations, and separate historical evidence follow.

2 Resource Model and Assumptions

2.1 Representations and physical memory

An execution profile π specifies the model, arithmetic, placement, inputs, state lifetime, and output-quality criterion. The evaluated profile performs text prefill and autoregressive decoding with resident KV state. A weight representation r includes its encoded payload, scales, decoding metadata, expansion, and required copies. Checkpoint size S_{ckpt}, complete compact-weight size W_{compact}, and full execution-weight footprint $W_{\text{full}}(r)$ describe different objects. None is the complete inference working set.

Let $\mathcal{D}_{\text{phys}}$ contain disjoint physical memory pools, with available capacity B_d in pool d. Aggregate capacity is

$$B_{\text{mem}} = \sum_{d \in \mathcal{D}_{\text{phys}}} B_d. \tag{1}$$

Discrete host and device memory can be separate pools; two aliases of unified physical memory cannot. Backing files add storage, not resident-memory capacity. Per-pool accessibility and placement constraints still apply when the aggregate capacity is sufficient.

The physical working set $M_d(t)$ includes attributable weights, expansion, activations, KV and recurrent state, staging, outstanding transfers, retained buffers, metadata, allocator and driver reserves, and resident file-backed pages. Shared allocations count once. The feasibility condition and a sampled observation are distinct:

$$\begin{aligned} M_d^{\text{peak}} &= \sup_{t \in \mathcal{T}} M_d(t) \le B_d, \\ \widehat{M}_d^{\max} &= \max_{t_i \in \mathcal{T}} \widehat{M}_d(t_i), \end{aligned} \tag{2}$$

where $\mathcal{T}$ covers the declared execution interval. Samples may miss transients or include other consumers. Independent component maxima cannot be added to obtain a simultaneous peak.

A full execution representation cannot reside in aggregate memory if $W_{\text{full}}(r) > B_{\text{mem}}$. This does not prove that storage is necessary: complete compact weights might still fit with online expansion. Conversely, a deliberately small cache can exercise bounded residency even when the compact checkpoint fits in physical memory. This paper studies that controlled-budget problem rather than compact-residency impossibility.

2.2 Accounting and stage-local admission

An accounting domain $a \in \mathcal{D}_{\text{acct}}$ may count bytes, slots, pages, or queue entries. Its configured limit b_a need not cover an entire physical pool. Accounting domains can overlap and cannot generally be added. For an additional claim c_a, the capacity condition is

$$U_a(t) + c_a \le b_a \qquad \text{for every covered domain } a. \tag{3}$$

Symbol	Meaning
π, r	Execution/quality profile and weight representation
$B_d, M_d(t)$	Physical pool capacity and attributable resident bytes
b_a, U_a, c_a	Accounting limit, current charge, and additional stage claim
δ, ρ_a	Demand class and reserve protected from speculation
F_a, R_a, I_a, L_a, T_a	Free, reserved, in-flight, live, and retained claims
$P = (K, \boldsymbol{q}, \boldsymbol{\ell})$	Frozen key, stage claims, and resident claims
K, g_s, g_x	Full resource identity, source generation, and destination generation
o, s, e, J	Operation, stage, completion evidence, and persistent owner journal
$\Omega, k, w, \mu_o[s]$	Cleanup worklist, visit limit, cursor, and cancellation marker
$H, P_{\max}, M, S$	Host limit, largest packed expert, metadata allowance, and requested stack bytes

Table 1: Core notation. Accounting quantities use their declared units; physical-memory and host-staging quantities are bytes.

Already held grants are not charged again on a revisit. The owner also validates units, overflow, source and placement compatibility, reservation authority, and demand policy. With reserve $0 \le \rho_a \le b_a$, the policy-accessible ceiling is

$$b_a^{\mathrm{eff}}(\delta) = \begin{cases} b_a, & \delta = \mathsf{Required}, \\ b_a - \rho_a, & \delta \in \{\mathsf{Warmup}, \mathsf{Prefetch}\}. \end{cases} \tag{4}$$

Required execution may use the reserve, but not exceed the hard limit.

Admission returns one of three abstract outcomes. Rejected means that the plan is invalid or a stage exceeds its policy ceiling even without competing occupancy. Wait means that an otherwise feasible stage is blocked by current occupancy; it retains existing grants and registers a revisit condition. Admitted means that the authoritative owner has validated and acquired the stage's claims. Admission is local to a stage, not an atomic reservation of all future resources.

A frozen plan $P = (K, \boldsymbol{q}, \boldsymbol{\ell})$ binds the full resource key, read/upload/install claims, and persistent residency claims. The key includes model instance, source identity, semantic content, encoding, backend, device, and generations g_s, g_x. The destination authority supplies the actual reservation and reuse generation before hard admission. Extent alignment and decoding can make source, staged, and execution byte counts different. The plan must account for those differences rather than assuming one byte count for every stage.

2.3 Completion, quality, and progress assumptions

Protocol correctness requires exact identity, authorized publication, and no reuse while physical references remain. Numerical quality is a separate contract: visible text, generated IDs, full-vocabulary logits, routes, and internal states provide different evidence. Argmax agreement alone does not establish numerical agreement, and checkpoint quantization does not specify activation or KV arithmetic.

Completion evidence e must identify operation o, key K, submitted stage s, both generations, outcome-specific bytes, and any required fence. A mismatched event has no mutation or release authority. A matching event that reports non-success can end its named stage but cannot authorize publication. Producer success and consumer retirement also differ: finishing an upload does not establish that a later kernel has stopped using the destination.

We use *custody* to mean responsibility for retaining the physical resources and evidence needed to close an obligation. Unknown denotes an obligation whose physical completion cannot be proved.

Such an obligation retains allocations and charges, potentially indefinitely. Cancellation changes demand or caller outcome; it is not itself a completion certificate.

The reference contracts assume authoritative owners, source and destination generations that identify reuse, and completion evidence supplied by the relevant provider. Conditional progress additionally requires feasible resources, acyclic waits, fair revisits, live providers, and eventually observable completion. The model does not assume that a blocking provider becomes nonblocking when called by a worker. It also does not place OS page cache, storage-server cache, or unaccounted allocator overhead under the runtime's logical credit limits.

3 Materialization, Custody, and Bounded Scheduling

3.1 Normal execution

A nonresident dependency moves through six stages. The adapter first freezes a source-bound plan and obtains an exact destination reservation. The owner then admits the read, upload, and installation claims as each stage becomes eligible. Before any preparation can submit physical work, it records the provisional operation and cleanup intent in a persistent journal.

A successful read returns read custody. The owner explicitly releases the read slot and storage-read claims while retaining the host payload charge for upload. The upload's fence ends DMA custody, but the destination remains owned. Installation validates the new binding; replacing an old binding also requires evidence that its previous consumers have finished. Publication then transfers the resident grant and resolves the exact waiters. It does not evict the newly resident bytes.

Before execution, the consumer revalidates its dependency and lease epochs and acquires any resume claims. Consumer completion and state publication have their own evidence. A request can finish while persistent KV state or cached weights remain live. Figure 1 summarizes these roles; Section 4 distinguishes the concrete operator path from the more general materialization protocol.

3.2 Safety invariants

Unique accounting ownership. For each covered domain, the intended partition is

$$\begin{aligned} b_a &= F_a + R_a + I_a + L_a + T_a, \\ U_a &= R_a + I_a + L_a + T_a, \qquad F_a = b_a - U_a. \end{aligned} \tag{5}$$

Every claim belongs to exactly one class: free, reserved, provider-in-flight, live, or retained. Different claims of one operation may occupy different classes. Unknown claims belong to T_a; uncertainty does not create a second charge. Returning custody can change ownership without reducing U_a. Release or transfer is an explicit action. This is an intended, conditional contract: it requires complete claim coverage and correct owner transitions in each declared domain. The tested logical host and device-cache bounds do not establish integrated five-class conservation or a complete physical-memory bound; the classes are not five literal implementation fields.

Exact evidence before mutation. An event must match the operation, key, stage, generations, and byte contract before it changes state. Success requires the full expected amount; a partial failure is not readiness. Stale identity cannot release a newer generation's grant. A matching event with a validated non-success outcome instead selects cleanup for its own operation. Notifications merely trigger revalidation.

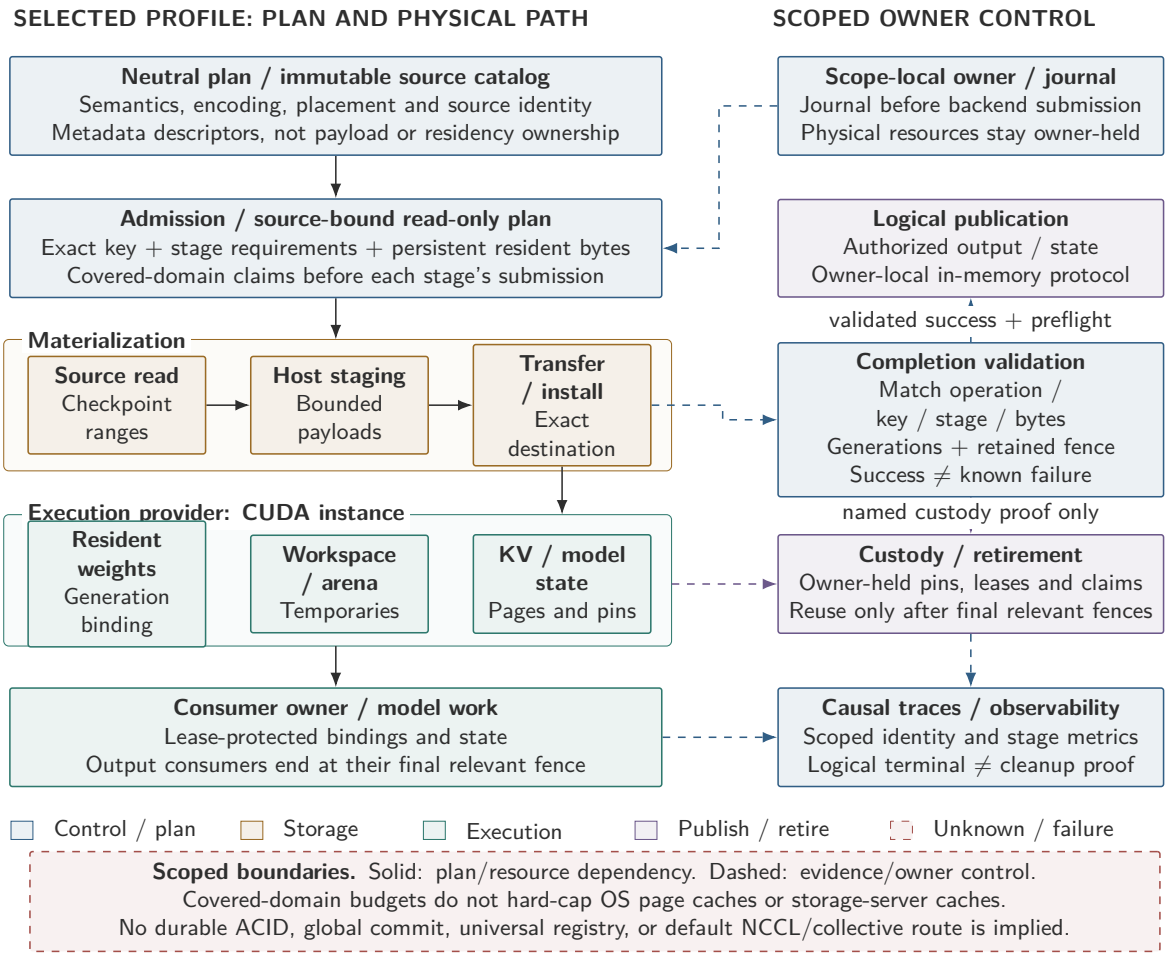


Figure 1: Profile-scoped architecture. Planning and admission precede physical materialization; the owner retains resources through publication and retirement. Blue denotes control, ochre storage, teal execution, violet publication and custody, and brick uncertainty. Solid arrows show dependencies; dashed arrows show evidence or owner control. The diagram specifies roles, not universal registry integration or measured hardware overlap.

Reachable physical custody. Each outstanding physical reference remains reachable from an authoritative owner, including after preparation or submission fails without returning a public handle. Read return, upload completion, resident transfer, and consumer retirement close different obligations. An address or read-only snapshot supplies neither allocation ownership nor completion authority.

Publication before reuse, with distinct proofs. Producer success makes a candidate ready; publication also validates the actual binding and authorization. Replacement requires old-consumer quiescence. For prepared KV state, exact installation acknowledgments precede logical publication, and retirement acknowledgments precede page reuse. This ordering is a scoped in-memory lifecycle, not distributed durable commit.

Terminal demand preserves residual obligations. Accepted cancellation detaches only the terminating demand. Independent waiters and persistent owners remain. Cancellation acceptance does not prove quiescence or reverse an irreversible publication decision. Missing physical proof retains the payload and charge as Unknown. A late exact completion can close a named stage without clearing a separately poisoned physical owner. Figure 2 makes the distinction between logical outcome and resource lifetime explicit.

3.3 Reference materialization and cleanup algorithms

Algorithms 1 and 2 express reference contracts for the relevant owners. Their helper names are protocol roles, not a claim that every runner invokes one common manager. The operation record and journal persist across calls. A pending or idle notification is pollable only while its original obligation remains known. AdmitStage checks intrinsic feasibility, canonical units, overflow, authority, compatibility, and the policy in Equations 3–4, then acquires grants with an owner-side capacity recheck. It preserves previously held grants on Wait. SelectCleanup records the primary reason once and prevents later calls from restarting normal dispatch.

Publication preflight distinguishes a live-custody block from an invalid binding. Rejection after installation preserves the journal and physical binding for cleanup; it does not undo the installation. A terminal load result likewise does not mean physical closure. Journaled returns and releases are applied once, even if a subsequent cleanup step fails.

A terminal pass operates on a finite worklist Ω of obligations reachable through J after demand detachment. The persistent cursor w and unvisited cohort prevent unresolved entries at the front from hiding later entries. New obligations cannot overtake that cohort. Removing an entry from the local worklist never removes it from its authoritative owner. A custody handoff takes effect only after acknowledgment by the new owner.

For each exact operation and stage, cancellation marker $\mu_o[s]$ is one of None, Pending, Submitted, or Uncertain. It also binds the key and generations. Known non-acceptance permits an owner-authorized retry with the original arguments; acceptance permits only polling. An ambiguous reply latches uncertainty and cannot be retried as though cancellation had been rejected. A different stage has a different marker. Safe discard of an unsubmitted reservation follows the same rule when it requires provider cancellation.

Algorithm 2 visits at most k obligations, with at most one observation and one cleanup advancement per visited item. This bounds work per pass, not provider-call latency or total cleanup time. A secondary error preserves the primary reason, original phase and arguments, and partial returns already recorded. Cleanup remains pending only while its original obligation is known pollable; otherwise it becomes Unknown. Neither a wake nor an unrelated later fence authorizes

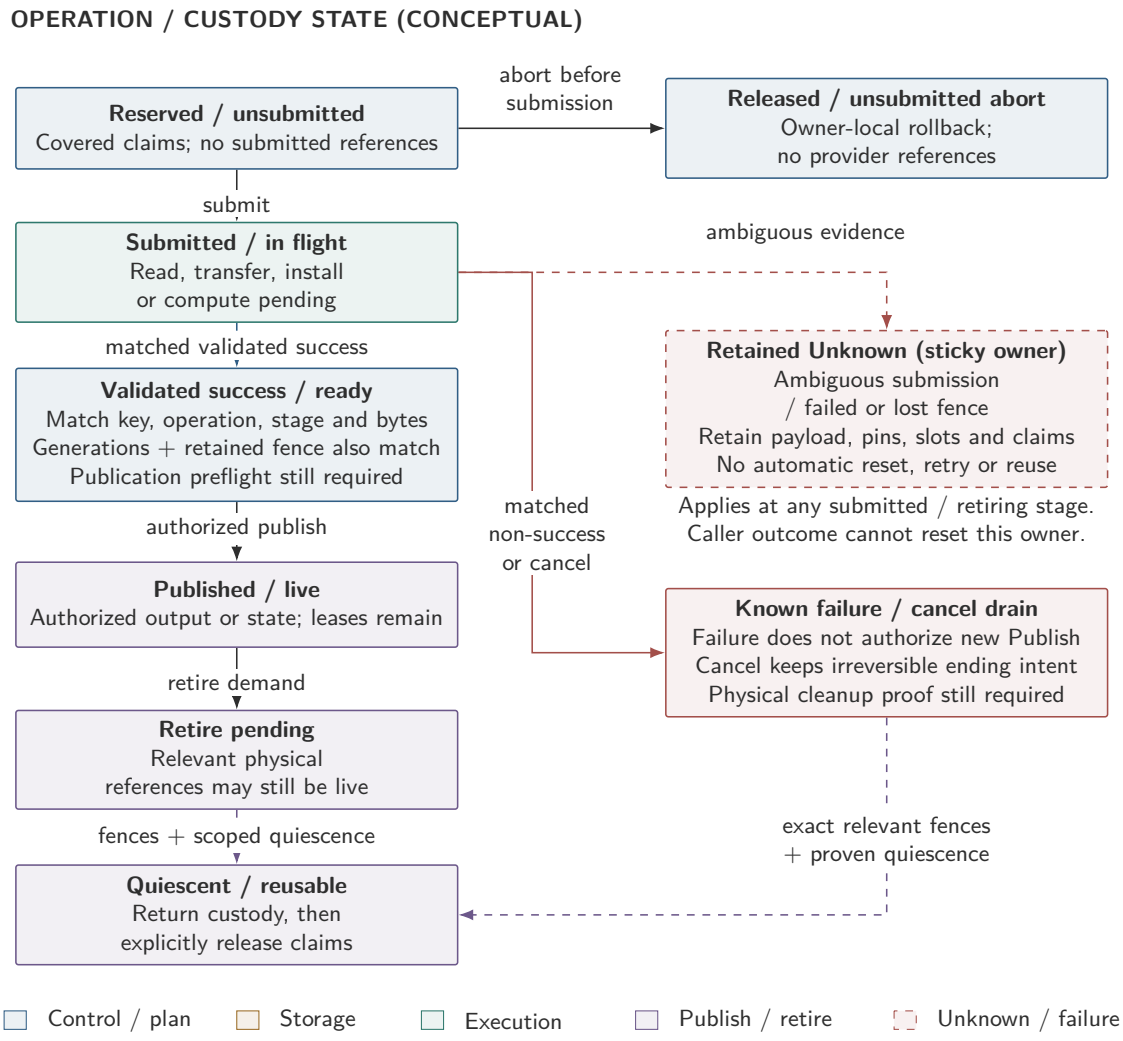


Figure 2: Logical operation and physical custody states. Exact validated success and publication preflight permit publication; known non-success selects cleanup. The dashed Unknown state retains resources and has no automatic recovery or reuse transition. Producer readiness does not establish consumer retirement, and returning custody is not an implicit credit refund.

Algorithm 1 Progress one exact materialization operation

Inputs: frozen plan P, provider v, ledger $\mathcal{L}$, demand δ, journal J
Outputs: result in {Ready, Pending, Terminal}; residual custody stays reachable

Persistent state: operation o, its stage s, grants, and original tickets
if cleanup is selected or o retains Unknown **then**
  continue only Algorithm 2; **return** Terminal
end if
while o is not resident **do**
  if o is installed **then**
    $p \leftarrow$ PublicationPreflight(o)
    if $p =$ Ready **then**
      publish exact binding; transfer resident grant; resolve exact waiters
      retire load once; mark o resident
    else if $p =$ Blocked **then**
      retain binding and grants; arm replacement-custody revisit
      return Pending
    else
      SelectCleanup(J, o, p); **return** Terminal
    end if
  else
    $s \leftarrow o.stage$
    if s is not submitted **then**
      $\alpha \leftarrow$ AdmitStage($P, s, \mathcal{L}, \delta$)
      if $\alpha =$ Rejected **then**
        SelectCleanup(J, o, α); **return** Terminal
      else if $\alpha =$ Wait **then**
        retain held grants; arm admission revisit; **return** Pending
      end if
      journal stage custody before SubmitOnce(v, o, s)
      if submission fails or its disposition is ambiguous **then**
        retain original payload and claims; select cleanup; **return** Terminal
      end if
    end if
    $e \leftarrow$ PollOriginal(v, o, s)
    if e is pending or idle and the original obligation is known pollable **then**
      arm completion revisit; **return** Pending
    else if e mismatches identity, stage, generations, or byte contract **then**
      reject without mutation; retain original obligation and revisit
      return Pending
    else if e is pending/idle without that proof, a provider fault, or lacks physical proof **then**
      retain Unknown; select cleanup; **return** Terminal
    end if
    apply exact stage-custody return once; retain residual claims
    if e reports validated non-success **then**
      SelectCleanup(J, o, e); **return** Terminal
    else if $s =$ Read **then**
      release read-slot/read-byte claims; retain host payload charge
      advance to upload; clear only completed read ticket
    else if $s =$ Upload **then**
      release upload and host claims after exact fence; retain destination
      advance to install; clear only completed upload ticket
    else
      mark o installed
    end if
  end if
end while
return Ready

Algorithm 2 One bounded terminal-custody pass

Inputs: journal J, terminal intent i, visit limit $k > 0$
Outputs: caller outcome y and custody status h; unresolved ownership persists

Persistent state: worklist Ω, cursor w, exact cancellation markers $\mu_o[s]$
if i requests cancellation and the terminal owner rejects it **then**
  return current outcome and custody status
end if
latch accepted y; detach terminating demand once; preserve irreversible intent
$\Omega \leftarrow \text{ResumeOrSnapshot}(J, w)$; $n \leftarrow \min(k, |\Omega|)$
while $n > 0$ **do**
  $o \leftarrow \text{Pop}(\Omega)$; $s \leftarrow o.stage$; $n \leftarrow n - 1$
  if another waiter or persistent owner still holds o **then**
    preserve shared load or acknowledged handoff; do not cancel or refund
  else if o has persistent physical-owner Unknown **then**
    retain payload, charge, and poison; block reuse
    apply any late exact stage completion only to its named obligation
  else if o is positively unsubmitted **then**
    attempt safe discard once under $\mu_o[s]$
    discharge eligible claims only with exact discard proof
  else
    $e \leftarrow \text{ObserveOriginalOnce}(o)$
    if e mismatches the exact completion contract **then**
      reject without mutation; preserve o
    else if e proves named stage-custody return **then**
      journal the return once; advance original cleanup at most once
      preserve primary outcome and residual custody on any cleanup error
      release only proved claims; confirm retirement before reuse
    else if e is pending and the original obligation is known pollable **then**
      if cancellation is legal and $\mu_o[s] \in \{\mathsf{None}, \mathsf{Pending}\}$ **then**
        $\gamma \leftarrow \text{CancelOnce}(o, \mu_o[s])$
        if $\gamma = \mathsf{Accepted}$ **then**
          $\mu_o[s] \leftarrow \mathsf{Submitted}$; await physical completion
        else if $\gamma = \mathsf{NotAccepted}$ **then**
          $\mu_o[s] \leftarrow \mathsf{Pending}$; retain exact retry arguments
        else
          $\mu_o[s] \leftarrow \mathsf{Uncertain}$; retain Unknown; prohibit retry
        end if
      end if
    else
      retain payload and journal; classify held claims in T_a without recharging
    end if
  end if
  advance w past o; defer unresolved entries behind unvisited entries
  arm exact completion, cleanup, or owner-progress revisit
end while
persist unvisited Ω, cursor, and all unresolved owner records in J
if no unvisited or unresolved custody remains **then**
  $h \leftarrow \mathsf{Closed}$ for this bookkeeping scope
else if physical-owner Unknown remains **then**
  $h \leftarrow \mathsf{UnknownRetained}$
else
  $h \leftarrow \mathsf{Pending}$; schedule a later pass
end if
return (y, h)

reuse. Clearing a poisoned physical owner would require a separately established teardown or isolation boundary, which these algorithms do not supply.

3.4 Bounded next-expert lookahead

Algorithm 3 describes the implemented synchronous expert-operator schedule. A bucket binds an expert to its selected token rows; the operator preserves the existing bucket order. Let H be its logical host-staging limit, $P_{\max}$ the largest selected expert's packed weights and scales, M the metadata allowance, and S the requested worker-stack reservation. The materialization owner holds one bounded worker across compatible sequential read sequences; a sequence does not own a new thread. Both schedules require

$$2P_{\max} + M + S \leq H \tag{6}$$

before worker creation or reuse. The factor two covers packed/prepared coexistence, not two permanent expert copies. During execution the charged occupancy is

$$C_{\text{host}}(t) = B_{\text{packed}}(t) + B_{\text{prepared}}(t) + M(t) + S(t) \leq H. \tag{7}$$

Packed quota covers queued, in-flight, and held completions. Prepared quota remains charged while any weight or scale alias survives. Aliases do not add copies, and dropping one alias does not release the remaining allocation. Retained Unknown is already included in these charges. One worker, a queue of depth one, and two charged read slots bound the outstanding host work.

Read completion validates the prepared source image, not a device destination generation. Destination binding and reuse checks remain with upload and consumer ownership. The residency hint is read-only: it neither pins the next expert nor alters replacement order. Speculative failure becomes relevant only if that exact expert is later demanded. Unused reads are drained and counted as wasted work; lookahead is therefore permitted to read more than blocking, although it does not do so in the paired experiment. Cancellation cannot interrupt an already executing synchronous file read. The owner still waits for its existing consumer proofs, so the complete forward remains synchronous.

Normal sequence completion returns the healthy worker to its materialization owner, not to a joined state. Only proved read slots and retired payloads are released. Metadata M and requested stack S remain charged while the worker is idle. Reuse requires the same image, reader, limits, and schedule, fresh source checks, and no surviving payload aliases or unresolved host custody. Normal final owner teardown closes the queue and joins outside the admission lock. Exceptional worker-side teardown instead retains an original join handoff for an off-worker caller. Stack credit requires an actual successful join, not thread exit or cancellation. Unknown remains sticky even after a later successful drain or fence.

Weak liveness witnesses track prepared payload lifetime without adding copies. Strong aliases retain quota, but their retirement is not a physical allocation or RSS-release proof: weak control blocks and allocator-retained storage can remain. Allocator overhead, page caches, thread-local storage, stack rounding, and driver reserves lie outside Equation 7.

4 Implementation

4.1 Capability-scoped resource owners

The model-neutral decoder binds immutable parameter descriptors to a materializer, expert provider, and execution operators. The materializer owns source validation and host quota; the CUDA leaf

Algorithm 3 Bounded next-expert reads in a synchronous forward

Inputs: materialization owner, ordered selected buckets $\mathcal{B}$, source-bound plans, mode $m \in \{\mathsf{Blocking}, \mathsf{Lookahead}\}$, limits H, M, S
Outputs: unchanged bucket order and device proofs; residual custody retained

validate three-artifact plans and Equation 6
check out the owner's compatible idle worker, or create it on first admission
State: at most one next ticket τ, exact deferred error, payload-alias charges
for current bucket b in $\mathcal{B}$ **do**
  admit and pin current device residency and scratch; validate source
  if current expert is not resident **then**
    drain unrelated speculation by its original identity
    honor any error deferred for b before new dispatch
    if no exact carried ticket for b **then**
      $\tau \leftarrow$ SubmitExact(b) under read-slot and byte limits
    end if
    await original τ; blocking waits immediately after submission
    if worker custody is Unknown **then**
      retain identity, charge, and poison; return through terminal cleanup
    end if
    take completion once
    Validate the original operation, prepared source-image identity, selected expert,
      read plan, byte count, and source freshness.
    if evidence mismatches or reports non-success **then**
      preserve primary error and residual custody; return through terminal cleanup
    end if
    reserve prepared quota while packed bytes remain charged
    construct gate/up/down artifacts without rereading; track payload aliases
    upload current expert under destination-generation checks and CUDA proofs
  end if
  drain speculation made redundant by a current resident hit
  release current host charge only after payload-alias retirement
  inspect next selected bucket without pinning it or changing replacement order
  if $m = \mathsf{Lookahead}$ and next expert is nonresident and host custody permits **then**
    try one next read before current compute submission
    defer speculative submission or read errors to that exact later demand
  end if
  submit current computation from resident device handles
  consume output under its lease; preserve existing fences and retirement
end for
on sequence end or error, cancel unused reads and drain original completions
release proved read-slot/payload claims; retain surviving alias charges
if clean, park worker at owner; otherwise preserve custody through exceptional shutdown
normal final owner teardown closes the queue and joins; retain stack charge until actual join
preserve any Unknown and its quarantine; host join does not discharge device custody

owns device-cache admission, scratch, uploads, and consumer completion. Capability checks select compatible read and transfer mechanisms rather than model-specific materialization factories. The full-model evidence covers Qwen3.5 only, not DeepSeek-V4.

This private expert path implements Algorithm 3 without instantiating the general runtime resource manager or materialization registry for these reads. Algorithms 1 and 2 are reference contracts for its responsibilities, not one universal call path.

The checkpoint catalog records tensor descriptors and ordered file extents. A read plan binds filesystem identity using device, inode, length, modification time, and change time. The worker verifies this identity before and after positioned reads; the consumer rechecks it before constructing the three compressed gate, up, and down artifacts. Descriptor hashes describe dtype, shape, and extent metadata rather than authenticating every payload byte. The worker receives neither CUDA handles nor cache leases or runtime continuations.

Both compared schedules enable the same owner-held worker. Blocking submits required demand and waits; lookahead carries one exact next ticket across current compute submission. The first admitted sequence creates the worker; clean later sequences reuse it without resetting identities or source snapshots. Sequence completion drains unused reads and releases proved payload/read-slot claims, but does not join. Timing boundaries and warmup exclusions remain those declared for each experiment; reuse implies no extra warmup procedure.

After upload, computation borrows resident device handles and immutable binding metadata, so temporary host aliases can retire before the next read. Kernel mathematics, selected expert order, and completed-only cache replacement are unchanged. A pipeline-disabled path is not an equivalent blocking control.

Responsibility	Evaluated implementation	Boundary
Source binding	Immutable catalog and ordered positioned-read plans	Metadata identity, not payload authentication
Host materialization	One persistent worker; packed/prepared quota held through aliases	Idle metadata and stack remain charged
Device residency	Private expert-cache owner, scratch admission, retained CUDA proofs	Separate from host read quota
Scheduling	One next-expert ticket before current compute submission	Synchronous forward remains unchanged
Transfer	Six pageable copies; grouped policy shares one exact stream fence	No pinned numeric-upload ticket or async guarantee
State lifetime	Prepared state, exact installation acknowledgment, retirement	Request completion can retain KV and weight residency
Observation	Passive counters, token records, owner snapshots, separate profiler	Observation does not refund or repair custody

Table 2: Mapping from the contracts to the evaluated operator's responsibilities.

The principal configuration uses a 64 MiB logical host limit, including 8 MiB metadata and 2 MiB requested stack; the remaining 54 MiB is split equally between packed and prepared quotas. With $P_{\max} = 3{,}146{,}112$ bytes, Equation 6 gives 16,777,984 bytes, matching the observed logical peak in the synchronous and grouped HTTP arms. Their idle live-owner charge remains 10,485,760 bytes, not zero. Device cache and scratch are separate domains; these charges do not bound total RSS.

Passive snapshots report reads, waits, charges, join state, and uncertainty without refunding or repairing custody. Final owner teardown, rather than sequence completion, owns the actual join. The full-model reports expose no post-owner-drop join witness. Process reaping and idle-device

observations do not supply that witness or certify physical memory release.

4.2 Synchronous expert transfer and serving policy

The baseline uploads three numeric expert projections using six pageable copies, each with its explicit completion fence. Grouped synchronous upload uses the same six copies and source representation, but one exact compute-stream fence covers the successful group. All six destinations and immutable source aliases enter custody before the first possibly submitting copy. There is no extra pinned allocation, concatenation buffer, upload event, or polling ticket. A failed proof retains the group and source charge and quarantines the owner; later proof cannot heal it. Pageable driver calls may still block or stage internally, so one explicit fence is not one total native synchronization or an asynchronous-transfer claim. These are code-level counts, not observed per-request backend-call counts.

Grouping does not remove the existing last-consumer event or owner upload-stream fence. Their redundancy has not been established for this leaf, lookahead, or other outstanding work. The complete forward remains synchronous; host lookahead alone does not prove hardware overlap.

With bounded reads absent or blocking, omitted transfer policy resolves to synchronous upload. A supported bounded-lookahead configuration on the local Mixed BF16/F32 CUDA numeric leaf instead resolves to grouped synchronous upload; explicit synchronous selection remains a baseline. Normal serving rejects grouped upload with blocking reads; that combination is confined to private experiments. These current selection rules do not relabel the frozen explicit-policy measurements on the same operator leaf or establish new numerical-quality results.

4.3 Other protocol implementations

The general materialization path installs a registry only when a runner transfers a physical provider to the resident engine. This path merges compatible exact keys, validates completion before mutation, and separates stage-custody return from credit release and resident-grant transfer. A separate routed-expert provider implements pinned reads and CUDA upload/installation. Its registered slabs retain aligned ranges and generations; tickets retain the original event and frame after notification errors. Replacement carries the old binding and consumer-completion evidence. This capability is distinct from the pageable numeric expert path, not evidence of its registry coverage.

The resource broker canonicalizes and checks claims before charging them. A later-stage budget wait retains earlier buffers and destinations. Grants have no implicit destructor refund. The five accounting classes in Equation 5 refine these ownership states; tested host/cache bounds are not an integrated five-class proof.

State publication uses a separate prepared-state protocol. A provisional journal precedes backend preparation because preparation can submit work before returning a handle. Lost installation acknowledgments cause polling of the same operation rather than redispatch. Exact acknowledgments precede logical publication; displaced-state retirement and page confirmation precede reuse. Committed KV reads revalidate owner, slot, incarnation, generation, layout, frontier, range, and page authority. Finishing a transaction consumes its handle custody, not every independent weight, KV page, or cache pin.

Other capabilities have separate boundaries. Host-KV placement exposes logical read pins and admission, not a completed production DMA path; its accepted or Unknown commands still require backend quiescence evidence. Shared CUDA transport independently retains fixed pinned chunks and device views until their own fences; it is not the pageable expert-upload mechanism. Hybrid and collective paths require their own audits, not extrapolation from single-device results. A late registry

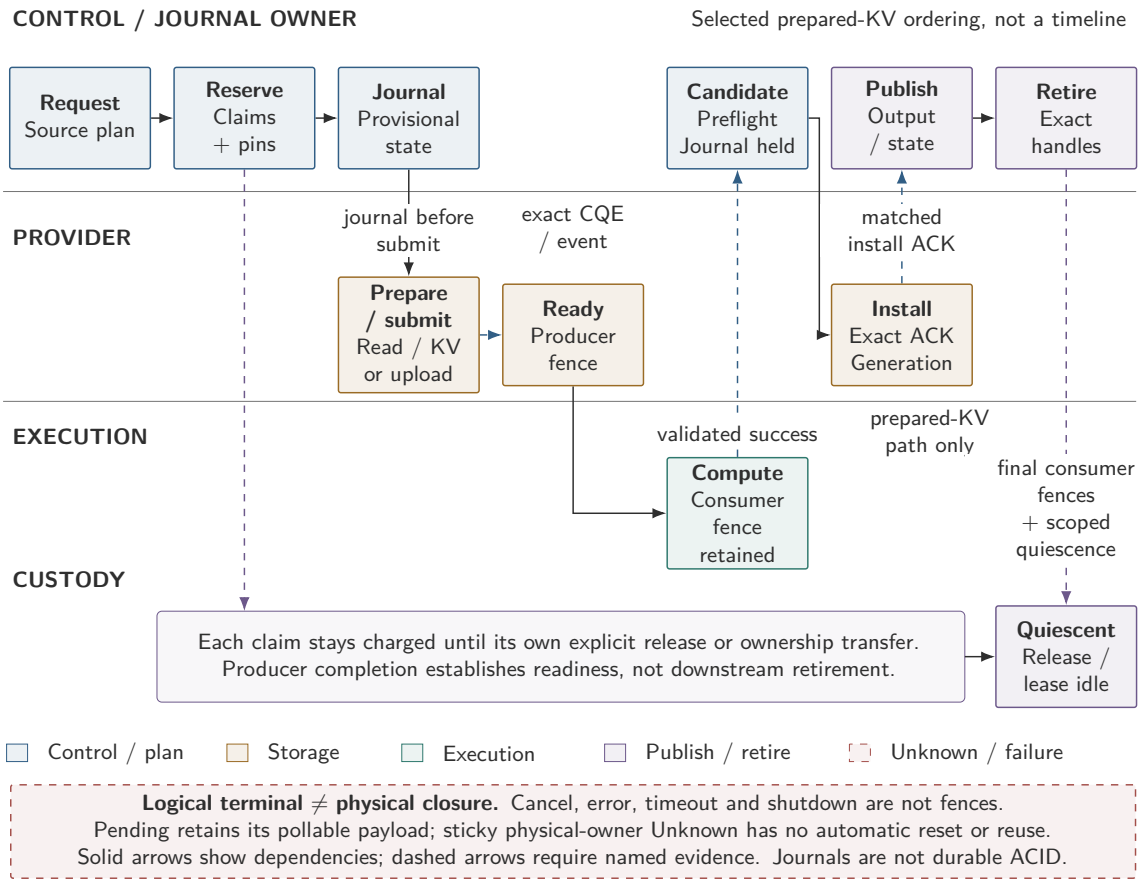


Figure 3: Prepared-state lifecycle. The provisional journal precedes backend preparation. Exact installation acknowledgments precede logical publication; consumer quiescence and retirement precede reuse. The lanes show causal ordering, not a measured overlap timeline. Cancellation preserves any irreversible ending decision and its residual cleanup obligations.

completion can close its named stage without clearing another owner's poison, and host-worker join cannot clear device uncertainty.

5 Experimental Methodology

5.1 Model, configuration, and quality criteria

The serving experiments use Qwen3.5-35B-A3B-FP8 with text-only execution on one CUDA device: pipeline, tensor, and expert parallelism are each one. Experts are materialized lazily and KV state remains resident. The Mixed profile uses BF16 activations, KV, and convolution state, with F32 accumulation and state-space-model state. Checkpoints reside on a shared NFS mount. The original five-pair read-scheduling study, its pressure sweep, read–compute captures, and HTTP-cancellation experiment share one frozen serving executable. The later persistent-worker, grouped-transfer, and HTTP-policy experiments are separate experiment families, with separately frozen application and control identities (Table 4). They are not additional runs of that executable. Component faults use a small test executable. The artifact guide records source, binary, and control identities; source identity alone does not transfer numerical or performance acceptance.

Platform and observation. The serving experiments use one NVIDIA GeForce RTX 3090 with 24 GiB of physical device memory. A same-boot post-hoc inventory associated with the paired study reported an Intel Xeon Platinum 8369B at 2.90 GHz, 128 visible logical CPUs, and 754.913 GiB of kernel-visible memory on x86-64 Linux 4.19.91; experiment-time affinity, cgroup limits, and unchanged hardware configuration were not established. The observer sequentially samples device-wide GPU memory and process-tree RSS/PSS, then waits 2 s. For samples wholly within request phases, including inter-request gaps, start-to-start intervals are 2.261–2.377 s in the paired study and 2.276–2.368 s in the pressure study; collection adds time to the nominal wait. Whole-lifecycle maxima use 49–54 samples per paired arm and 36 per pressure point, not continuous or uniform-rate observations.

Table 3 gives the original five-pair configuration. Within that older experiment, the two arms change only the read schedule. Both retain the same source checks, arithmetic, cache policy, worker, consumer proofs, and synchronous forward. Runtime profiling is disabled for ordinary timings; the separate overlap captures enable it. Logging is identical between timing arms. No experiment flushes page or NFS caches, and a fresh server is not assumed to have cold storage.

Control	Paired-study setting
Read pipeline	One positioned-read worker, queue depth one, two charged slots
Host staging	64 MiB logical ceiling, including 8 MiB metadata and 2 MiB requested stack
Expert device cache	64 entries; 320 MiB charge including 8 MiB scratch
Other device policies	7 GiB admission target, 2 GiB workspace, 16 MiB KV
Scheduling	Context 128; batch-token limit 8; active-sequence limit 2; decode cohort target 1 without deferral
Build and observation	CUDA 12.4.99, SM86; profiling annotations compiled, runtime profiling disabled
Requests	Concurrency one; fixed prompts, resets, generation, and request order

Table 3: Controls held fixed between blocking and lookahead. Limits can overlap and are not additive physical-memory partitions. MiB and GiB denote 2^{20} and 2^{30} bytes.

Independent reference. Numerical references use Hugging Face eager model forwards from Transformers 5.5.3 and Torch 2.13.0+cu129, independently of Ferrule's execution kernels. FP8 checkpoint values and BF16 scales are decoded in F32 and rounded to BF16 execution weights, with BF16 activations and KV state and explicitly preallocated F32 SSM state. This is decoded-storage BF16 execution, not the stock dynamic-FP8 quantizer or default SSM-cache precision. Exactly equal raw routing logits are resolved deterministically by lower expert ID; distinct-logit selections and native F32 routing probabilities and weights remain unchanged. Full-forward tests use a frozen two-prompt reference; retained-state tests and HTTP canaries use a separately generated three-prompt reference with the same pinned HF implementation and precision policy.

Quality is checked at two levels. Factory tests compare the model's numerical outputs with these references using tolerance $0.00025 + 0.012|\text{reference}|$, with exact generated-ID checks where observed. They cover the complete 40-layer forward, retained fresh-state execution, ordinary execution without diagnostic observers, and cold-owner execution in both schedules, plus a small constructor/provider check. Factory tests and the serving executable are distinct, so HTTP observations supply the serving-quality evidence rather than inheriting numerical acceptance from source identity alone.

For HTTP requests, we join actual worker-enqueued token IDs to request identities, emission ordinals, terminal usage, and complete server-sent-event (SSE) responses. IDs are not inferred by retokenizing text. An SSE content event is not necessarily a token boundary. In the original serving study, eight-token requests must match the frozen ID, text, and usage references; two-token smoke requests check length and usage only. The original read–compute and pressure requests use those same criteria within their stated sets. The later public-CLI acceptance and HTTP-policy pilot instead check exact IDs, text, finish reason, usage, and request/EOS semantics for all outputs, including the two-token Hello request. EOS stopping remains enabled; QA happens to produce eight actual completion tokens, not an imputed eight-token normalization.

5.2 Experiment families and estimators

Table 4 separates numerical gates, whole-test engineering timings, serving acceptance, and the later descriptive HTTP pilot. Their denominators are never pooled. The model and arithmetic policy above are shared scientific controls, not a claim that all stages execute the same source or binary.

Later HTTP-policy pilot. Each arm starts a fresh server and executes 12 serial requests: four initial checks, one excluded QA, two QA warmups, and five measured QA requests. Within each read schedule r, two paired blocks compare synchronous (S) with grouped synchronous (G) transfer; both arms retain equal observed logical work and exact outputs. Let s_{krpj} be client send time, c_{krpj} the first nonempty SSE content event, e_{krpj} client terminal time, and d_{krpj} dispatch time. Define

$$
\begin{aligned}
F_{krp} &= \operatorname{median}_{j=1}^{5}(c_{krpj} - s_{krpj}),\\
E_{krp} &= \operatorname{median}_{j=1}^{5}(e_{krpj} - s_{krpj}),\\
W_{krp} &= \max_j e_{krpj} - \min_j d_{krpj}, \qquad Q_{krp} = \frac{\sum_j n^{\text{qualified}}_{krpj}}{W_{krp}},\\
R_{kr,X} &= \frac{X_{krG}}{X_{krS}}, \qquad \widetilde{R}_{r,X} = \operatorname{median}_{k=1,2} R_{kr,X}, \quad X \in \{F, E, Q\}.
\end{aligned}
\tag{8}
$$

Here $n^{\text{qualified}}$ counts actual reference-qualified completion tokens. TTFT is client first *nonempty content*, not first GPU-token latency. Phase W includes inter-request gaps and failures; qualification requires the entire lifecycle and both paired arms, rather than dropping unsuccessful slots. All

Experiment family	Analysis population	Question and acceptance scope
Original read scheduling	Five pairs; ten lifecycles; 120 requests, 50 measured QA	Blocking versus lookahead on one older frozen executable
Persistent-worker qualification	Nine numerical/lifecycle gates	Worker ownership and reuse; not nine latency replicates
Grouped-transfer qualification	14 gates: one backend, five model, eight composition	Numerical and contract checks; profiling is separate
Grouped whole-test / profiling	Four elapsed-time pairs per schedule; separate qualified partial profile pair per schedule	Whole-test elapsed and observed process mechanisms, not HTTP TTFT
Public-CLI acceptance	Four policy/schedule cells; 48 requests	Exact serving outputs and reconciled counters; no timing aggregate
HTTP-policy pilot	Eight lifecycles; 96 requests; 40 measured QA / 320 tokens; two pairs per schedule	Grouped versus synchronous within each read schedule
Rejected retained upload	Four whole-test pairs per schedule; separate partial captures	Archived negative alternative; never pooled with grouped transfer

Table 4: Distinct experiment families. A gate, request, lifecycle, and pair are different units. Successful sub-stages do not turn a failed planned stage into a complete experiment. Detailed identities are in the artifact guide.

observed measured completions contain eight tokens. The analysis unit is the paired block, not the 40 measured requests. The median of two ratios is their arithmetic mean, not a median-log estimate. We show both pair points and no confidence intervals or significance test. Startup is excluded. The separate serialized service rate $40/\sum_j (e_{krpj} - s_{krpj})$ excludes gaps and must not be called phase-wide goodput. Neither rate measures concurrent serving throughput.

Original read-scheduling estimator. Five paired blocks use the fixed order B–L, B–L, L–B, L–B, B–L, where B is blocking and L is lookahead. Each arm starts a fresh server and executes 12 requests: four initial checks, one excluded QA request, two QA warmups, and five measured QA requests. Thus 50 measured requests are nested within ten lifecycles and five paired analysis units. The protocol fixes the workload, order, time limits, and all-five-pairs-or-no-aggregate rule before collection. It permits no retries, replacement pairs, extra warmups, or pooled profiling timings.

For block i, mode m, and measured request j, client end-to-end time t_{imj} runs from send to response EOF. Time to first content f_{imj} (often called TTFT) runs from send to the first nonempty SSE content event. The arm summaries and paired estimator are

$$\begin{aligned} E_{im} &= \operatorname{median}_{j=1}^{5} t_{imj}, \qquad F_{im} = \operatorname{median}_{j=1}^{5} f_{imj}, \\ G_{im} &= \frac{40 \text{ qualified tokens}}{\sum_{j=1}^{5} t_{imj}}, \quad \widehat{R}_X = \exp\!\left(\operatorname{median}_{i=1}^{5} \ln \frac{X_{iL}}{X_{iB}}\right), \quad X \in \{E, F, G\}. \end{aligned} \tag{9}$$

End-to-end latency is primary; TTFT and goodput are secondary. G is a qualified serial request-window rate, excluding startup and inter-request gaps, not concurrent server throughput. The descriptive resampling procedure enumerates all $5^5 = 3{,}125$ ordered paired-block resamples and

exponentiates the median log ratios at nearest ranks 79 and 3047. It does not treat individual requests as independent replicates.

5.3 Original complementary experiments and observation scope

Three fresh blocking lifecycles jointly vary cache entries and byte ceilings in fixed order: 8 entries/64 MiB, 64/320 MiB, and 128/640 MiB. Each executes one two-token smoke and three eight-token canaries: QA, arithmetic, and Chinese. This is a coupled-budget sensitivity study, not independent entry and byte ablations. The new points receive HTTP/ID checks, not a transfer of full-logit, route, or state qualification.

The overlap experiment captures one fresh QA request per schedule, separately from the timing pairs. It correlates worker read envelopes, contained positioned-read system-library calls, and actual expert kernels on the profiler's common clock. The component experiment crosses two schedules with clean failure and withheld completion proof in a two-layer numerical fixture. It executes real reads, uploads, and CUDA computation. The HTTP cancellation experiment instead uses the complete model: it disconnects after the first nonempty content event and then issues a new recovery request in the same process.

Interpretation boundary. The measurements answer different questions rather than composing into one universal guarantee. Logical charges are not RSS or certified physical peaks; shared NFS counters are not request-exclusive backend traffic; and component failure checks are not full-model fault conservation. The original scheduling estimate is descriptive at $n = 5$; the later policy pilot has only two pairs per schedule: shared cache state leaves cross-block independence unestablished, so its ranges are not guaranteed-coverage confidence intervals or significance tests. Table 5 consolidates the main availability limits; the results below retain each measurement's units and scope without imputing missing fields as zero.

Question	Available evidence	Not established
RQ1	Numerical factory checks; HTTP/ID checks; logical high-waters; device and host samples	Total physical peak, physical admission at each cache point, capacity frontier, storage necessity
RQ2	Logical reader/upload counters; request process-I/O snapshots; shared NFS deltas	Startup/shutdown process-I/O deltas, request-exclusive network or backend-disk bytes
RQ3	Older five-pair scheduling study; later two-pair/schedule HTTP policy pilot; separate scoped captures	Statistical significance, global trace completeness, physical storage-service overlap, critical-path causality
RQ4	Four real-CUDA component cells; full-model HTTP cancellation and new-request recovery	Hardware/driver faults, integrated five-class conservation, full-model Unknown recovery, final physical release

Table 5: Measurement availability and limits. A missing observation remains unknown rather than becoming a zero or a guarantee.

6 Evaluation

The current engineering evidence supports persistent read ownership and bounded grouped transfer, but not a uniform latency improvement. We report those results separately from the older read-

scheduling contrast, logical I/O observations, and failure experiments. Each comparison retains its own workload, acceptance criteria, and observation scope.

6.1 RQ1: Bounded completion and cache-budget sensitivity

Persistent ownership and grouped qualification. The later persistent-worker stage passes all nine gates. Both schedules create one worker, with 159 reuses in the full 40-layer check and 2,559 in each fresh-state and no-observer check. The observed live-idle charge is 10 MiB (8 MiB metadata plus 2 MiB requested stack). Zero joins in those still-live snapshots do not establish final-owner join or physical release. Grouped-transfer qualification separately passes all 14 gates: one backend, five model, and eight composition gates under the same numerical tolerance and required exact comparisons. These are quality/lifecycle tests, not additional serving observations or elapsed-time replicates.

Public-CLI acceptance covers all four read-schedule/transfer-policy cells: 48 of 48 requests match exact IDs, text, finish reason, and usage, with reconciled request windows. The later HTTP pilot separately completes eight lifecycles and all 96 requests, including all 40 measured QA requests and 320 actual measured completion tokens. All 96 outputs meet the exact serving reference and EOS/payload contract; none is failed, censored, or unissued. This does not imply that logits or routes were observed through HTTP.

Logical host quantity	Bytes	MiB
Original baseline / later grouped peak	16,777,984	16.0007
Rejected retained alternative peak / reservation	36,701,896	35.0017
Later persistent-worker live idle	10,485,760	10
Declared host quota	67,108,864	64

Table 6: Logical capacities, not RSS, allocator, memlock, or continuous physical peaks. Equal baseline/grouped values are separate observations, not pooled samples. The rejected alternative requires extra reservation; it is not the production design. Live idle includes retained worker metadata and requested stack, not proof of a final join.

Original bounded-completion study. All 120 requests in the original five-pair study complete under the declared accounting limits, and all three cache-budget points complete their separate four-request workloads. The numerical checks for the paired configuration also satisfy their acceptance criteria in its original nine factory tests (Table 7), distinct from the later worker gates. These observations establish bounded completion at the tested settings, not a minimum feasible capacity.

The original serving study has ten completed lifecycles, 110 exact eight-token oracle matches, and ten length/usage-only smoke checks. All 900 generated IDs join to successful HTTP terminal evidence. Failed, censored, and unissued request counts are zero. The 50 measured QA requests remain the latency subset; quality checks do not increase the number of paired replicates.

Every original paired lifecycle reaches the same logical host high-water:

$$2(3{,}146{,}112) + 8{,}388{,}608 + 2{,}097{,}152 = 16{,}777{,}984 < 67{,}108{,}864 \text{ bytes}. \tag{10}$$

This matches the packed/prepared, metadata, and requested-stack accounting in Equations 6 and 7. The selected-device sampled maximum is 4,125 MiB in every arm, based on 49–54 samples per arm, against the 7 GiB policy target. Figure 5 reports these logical and sampled quantities separately.

Factory scope	Tests	Metrics/test	Vocabulary values/test	Additional check
Full model	2	804	1,986,560	Prefill and first decode; diagnostic routes/states
Retained fresh state	2	5,427	9,187,840	64 fingerprints per schedule across four streams
Ordinary execution	2	27	9,187,840	64 fingerprints match same-schedule diagnostic outputs
Cold owner	2	1,809	3,228,160	28 fingerprints match current reference/configuration
Constructor	1	—	128 total	Four cells, each with 32 finite values and six verified reads

Table 7: Original factory quality checks for the older paired serving configuration. Two-test rows contain one test per schedule. Vocabulary counts for stateful rows describe common-history comparisons, not every checked stream. Ordinary-execution route and state metrics are unobserved.

Coupled cache sensitivity. Increasing the cache budgets increases ending residency and sampled device memory, but does not reduce the observed reader work. The three fresh points complete 12 requests, including nine exact eight-ID canaries, and generate 78 tokens in total, 72 covered by those oracles. Figure 4 shows ending entry counts of 8, 64, and 128; logical cache high-waters of 41.554, 209.574, and 401.598 MiB; and sampled device maxima of 3,997, 4,125, and 4,317 MiB. Each device maximum uses 36 retained observations. Host-tree PSS maxima are 6,417.381, 6,417.131, and 6,419.145 MiB.

Each point performs exactly 36,240,064,128 logical reader bytes, 69,114 reader calls, and 11,519 verified expert reads/uploads, with zero observed expert hits. Evictions decrease from 11,511 to 11,455 to 11,391; installs minus evictions explains the ending entry counts. The reduced eviction count is therefore additional retained residency, not evidence of cache reuse or fewer reads. Host logical high-water remains 16,777,984 bytes at every point.

6.2 RQ2: Equal logical demand and limited storage attribution

In the original five-pair study, blocking and lookahead perform identical observed logical reader work in every paired lifecycle, while the available filesystem observations do not identify physical backend traffic. For request q, let $\mathcal{V}_q$ be its anchored decoder invocations and define

$$L_q = \sum_{v \in \mathcal{V}_q} \Delta_v B_{\text{read}}, \qquad H_q = \sum_{v \in \mathcal{V}_q} \Delta_v B_{\text{upload}}, \qquad \Delta_q = H_q - L_q. \tag{11}$$

Here B_{read} counts logical checkpoint-reader bytes and B_{upload} counts bytes requested through the host-to-device copy interface, not physical bus traffic. The 1,020 invocation records reconcile with owner-lifetime totals and the shutdown ledger. High-water gauges are maxima, never summed deltas.

Table 8 shows that lookahead changes when work is issued, not its aggregate amount in this workload. Both modes repeatedly materialize weights and record zero expert hits. Equal aggregate counts do not by themselves establish an identical route-ID sequence or infer a cache-miss rate.

The paired archive contains 120 request-window and 30 lifecycle-phase NFS counter deltas. The first blocking request phase records 314,825,320 bytes in the shared client's server-read counter. Its individual request windows sum to 312,732,264 bytes, leaving 2,093,056 bytes between those windows inside the phase. The other nine request phases record zero. In the separate pressure

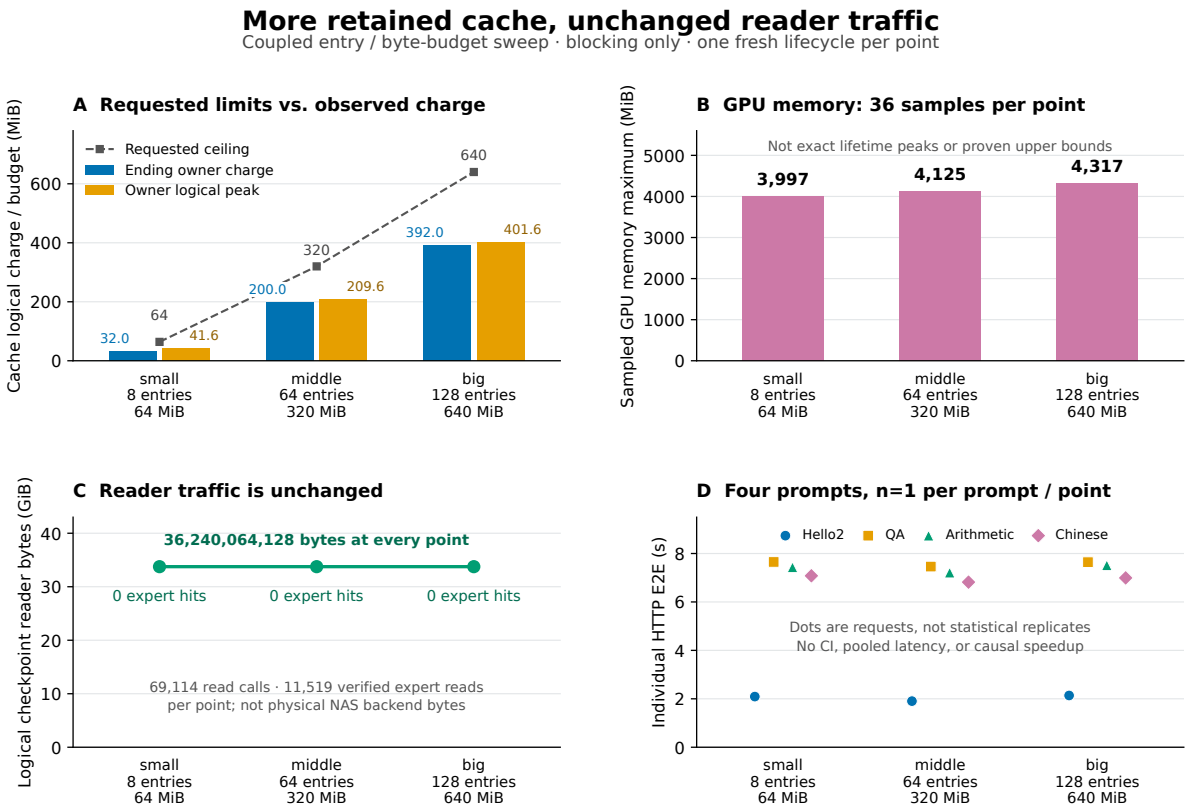


Figure 4: Coupled cache-budget sensitivity. Entry/byte limits increase together from 8/64 MiB to 64/320 MiB and 128/640 MiB. (A) Ending charge and logical high-water are separate observations under the requested ceiling; charge includes 8 MiB scratch. (B) Device-memory values are maxima of 36 samples per point. (C) Logical reader bytes remain 36,240,064,128 with zero expert hits. (D) Each latency dot is one prompt at one point, not a statistical replicate. All 12 requests complete. The sweep demonstrates bounded completion and increased residency, but neither read reduction nor a physical capacity frontier.

Per lifecycle, 12 requests	Blocking	Lookahead
Submitted and verified read operations	40,943	40,943
Checkpoint-reader calls	245,658	245,658
Logical reader bytes	128,811,263,616	128,811,263,616
Host-to-device interface bytes	128,811,927,312	128,811,927,312
Speculative reads consumed	0	36,863
Wasted read bytes	0	0
Uncertain read operations/bytes	0 / 0	0 / 0

Table 8: Identical totals in each of the five lifecycles per mode. A read operation can contain multiple positioned reads. Upload-interface bytes exceed reader bytes by 663,696 per lifecycle.

study, the shared client's normal-read counter increases by 36,491,264,808 bytes per point, while server-read and read-RPC deltas are zero. These are distinct client-level observations, consistent with cache-served logical reads but insufficient to attribute service to a physical NAS device.

Request-boundary process-I/O observations comprise 240 identity-bound snapshots and 120 deltas. Corresponding startup/shutdown process-I/O snapshots and startup logical-reader totals were not collected. Neither the zero client deltas nor fresh process lifecycles establish storage-cold behavior. The later HTTP-policy observations likewise include zero shared-client READ-RPC deltas and process `rchar` observations, not request-exclusive physical NAS traffic. OS page-cache service is not cold-filesystem or backend-disk evidence.

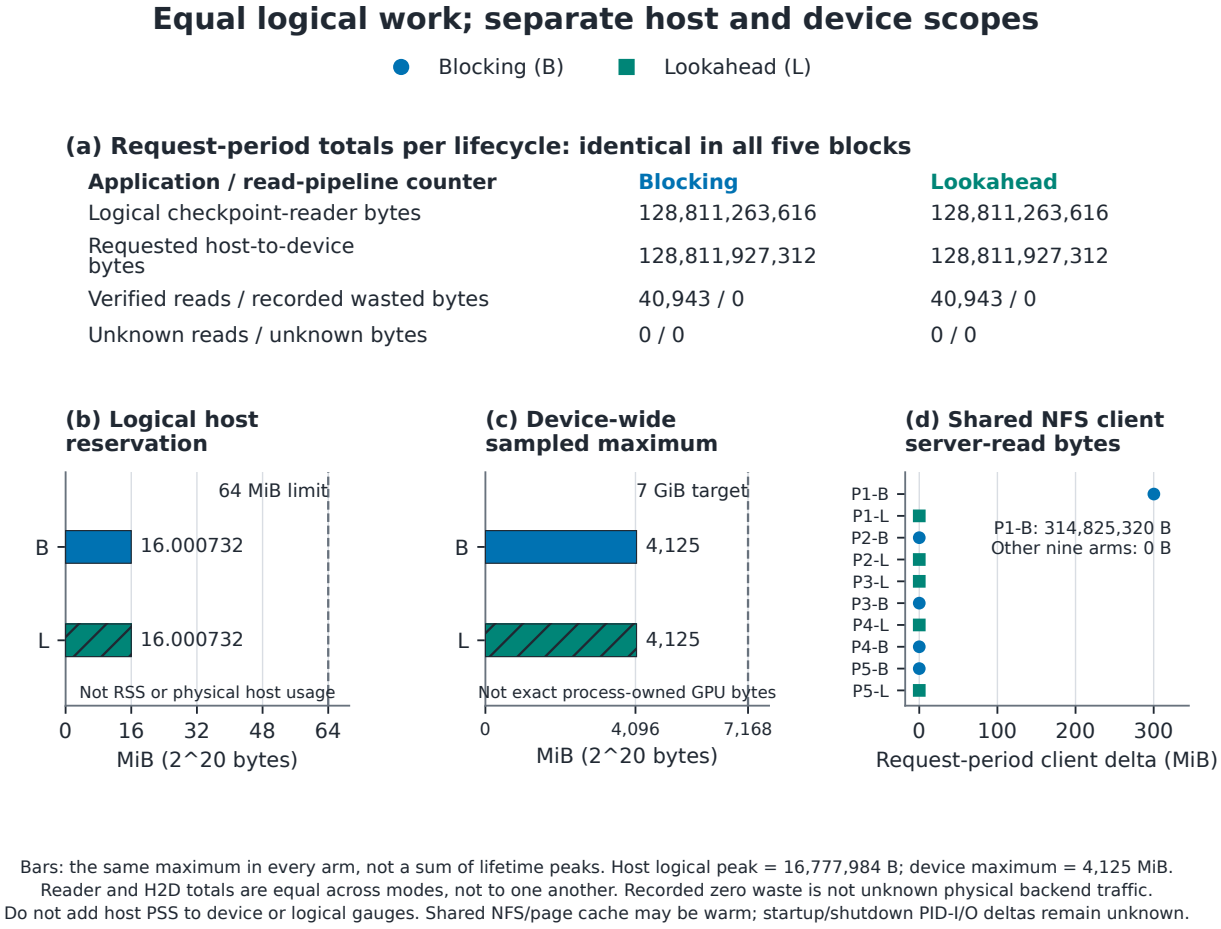


Figure 5: Logical work and memory observations for all ten paired lifecycles. Both modes read 128,811,263,616 logical bytes and request 128,811,927,312 upload bytes per lifecycle, with zero wasted read bytes. Logical host high-water is 16,777,984 bytes under 64 MiB; device-wide sampled maximum is 4,125 MiB under the 7 GiB policy target. These quantities are not additive physical partitions. Only the first blocking request phase has a positive shared-client server-read delta, with 2,093,056 bytes outside individual request windows. Shared-client counters do not measure request-exclusive backend traffic.

6.3 RQ3: Policy effects, mechanisms, and the older scheduling contrast

Current HTTP-policy pilot: a mixed outcome. Figure 6 compares grouped with synchronous transfer *within* each read schedule using Equation 8. Blocking has median paired ratios 1.049610928 for client TTFT (**4.96% regression**), 0.995661998 for E2E (0.43% lower, approximately flat), and 1.037363952 for phase-qualified goodput (3.74% higher). Lookahead ratios are 0.915281371 for TTFT (8.47% lower), 0.926627716 for E2E (7.34% lower), and 1.065190254 for phase-qualified goodput (6.52% higher). Both pair points remain visible; $n = 2$ per schedule does not support significance or a robust optimal-policy claim. The serialized-service paired ratios, 1.039273072 and 1.069921443, are different quantities and are not substituted for phase goodput.

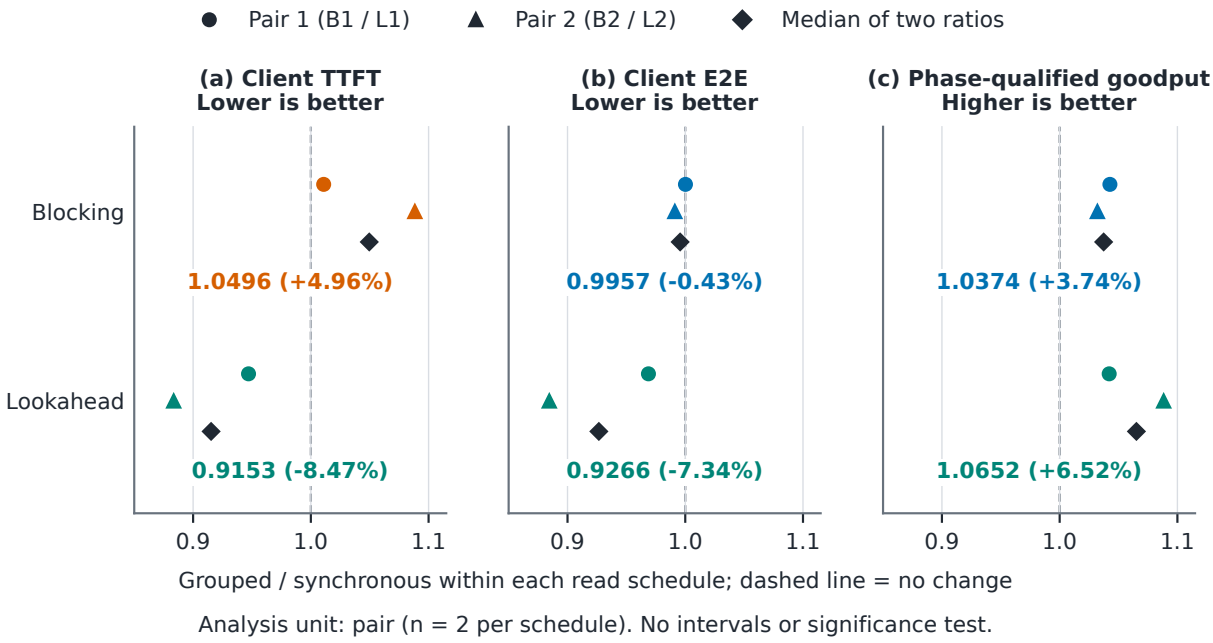


Figure 6: Later HTTP-policy pilot: grouped/synchronous ratios within blocking and lookahead, not the older lookahead/blocking comparison. Circles and triangles show the two pair ratios; black diamonds show their arithmetic median. Dashed lines mark no change. Client TTFT ends at first nonempty SSE content, E2E at client terminal, and phase-qualified goodput includes inter-request gaps. Blocking TTFT regresses by 4.96%. Eight lifecycles contain 96 requests; only 40 QA requests (320 actual tokens) are measured. The analysis unit is the pair, with two pairs per schedule; no interval or significance is claimed.

Grouped mechanism, not a causal timing decomposition. Separately qualified partial process captures for both schedules observe 382,743 native `cuStreamSynchronize` calls for synchronous transfer versus 257,163 for grouped: $-125{,}580/25{,}116 = -5$ calls per logical cache upload (Figure 7a). Each pair retains equal ordered logical work and output fingerprints. Observed H2D counts and bytes are unchanged at 195,985 copies and 82,512,273,944 bytes, and observed kernel names/counts match (409,411 kernel rows). Normalizing process totals by 25,116 logical uploads does not isolate an expert range or its elapsed cost. Native host-duration sums cannot be interpreted as causal client-latency gains.

Blocking and lookahead captures use separate observer-control stages but the same frozen

application; each qualified pair uses a common control within its own stage. The original four-capture stage remains failed; the later qualified lookahead pair does not repair its rejected capture or create a single successful four-capture stage. These are selected-model, process-wide partial observations, not full traces. Drop counts remain unknown, not zero, and absence of an API row is not evidence of zero calls.

The earlier grouped whole-test comparison has eight unprofiled pairs (four per schedule), with grouped/synchronous median elapsed ratios 0.96133 and 0.96110, approximately 3.9% lower. Its planned profiling stage failed; it is not a 24/24 success. Libtest elapsed includes construction, checks, I/O, and teardown: it is neither TTFT nor request latency, and is not pooled with the HTTP pilot or profiled durations.

Rejected retained-upload alternative. The archived alternative is slower in all eight whole-test pairs: median retained/synchronous ratios are 1.5451 for blocking and 1.5649 for lookahead (54.51% and 56.49% more elapsed time). Separate partial captures observe host-pin allocations and frees increasing from two to 25,118 each, with extra pin/free host-duration sums of approximately 29.85 s and 29.47 s (Figure 7b). These sums may overlap and are not an exclusive causal explanation of the unprofiled regression. The negative data remain archived; retained upload is removed from the production design, not recommended as its implementation.

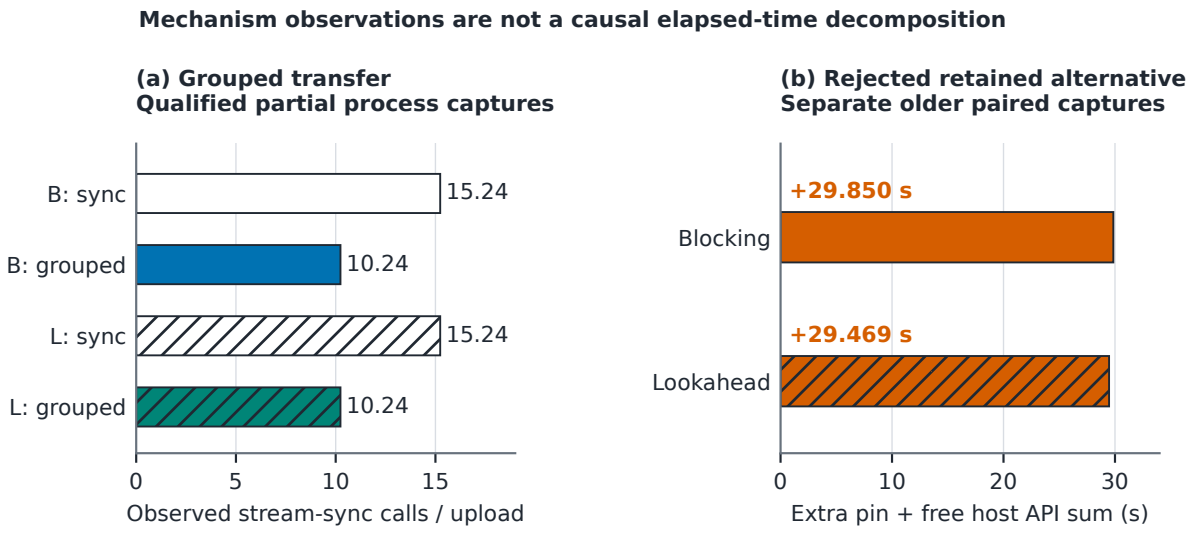


Figure 7: Separate mechanism observations. (a) Qualified partial grouped/synchronous process pairs, one per read schedule: native stream-sync counts normalized by 25,116 logical uploads fall by five calls/upload, while observed H2D work and kernel counts match. (b) The older, rejected retained/synchronous alternative adds 29.850 s (blocking) and 29.469 s (lookahead) of summed host pin/free API duration; each allocation/free count rises from two to 25,118. Panels use separate frozen experiments, not a common timing baseline or additive critical-path decomposition. Both are partial process observations; full-trace completeness and dropped-event counts remain unknown.

The post-cleanup production choice is conservative: synchronous transfer with blocking reads,

grouped synchronous with lookahead. The frozen pilot contains all four policy/schedule cells needed for the ablation; it does not measure the later cleanup or establish a causal optimum. Eight-token outputs do not establish sustained decode, token-aligned ITL, long-output, or concurrent-serving performance.

Older five-pair read-scheduling contrast. On its original frozen executable, lookahead reduces median host demand waiting and client end-to-end latency in all five paired blocks. The prespecified summaries are 8.18% lower end-to-end latency, 11.24% lower TTFT, and 10.91% higher qualified serial-window goodput (Table 9 and Figure 8).

Outcome	Aggregate L/B	Relative change	Descriptive range
Client end-to-end latency	0.91823	−8.18%	[0.89303, 0.94946]
Client TTFT	0.88755	−11.24%	[0.86004, 0.94888]
Qualified serial-window goodput	1.10908	+10.91%	[1.07393, 1.13498]

Table 9: Older five-block read-scheduling estimates from Equation 9. Ranges are descriptive block-resampling endpoints, not significance tests or guaranteed-coverage confidence intervals.

Across the five measured arm windows per mode, blocking generates 200 qualified tokens in 193.581 s of summed request time, and lookahead generates 200 in 175.393 s. The resulting 1.03316 and 1.14029 tokens/s are separate accounting summaries, not replacements for the paired estimator. Launch to readiness takes 12.847–13.247 s and is excluded from request latency.

Host demand-wait reductions range from 0.398 to 0.703 s per arm-median request; end-to-end reductions range from 0.384 to 0.827 s (Figure 9). A request's wait is the sum of anchored per-invocation demand-wait durations, followed by the median within each arm. Equal logical work and zero wasted bytes support interpreting this as a scheduling contrast rather than a smaller workload. The association does not decompose the critical path: host waiting cannot be added to client latency or treated as a measurement of GPU activity.

Original direct overlap observation. The separate profiled lookahead request has 174.409210 ms of reader-envelope intersection with expert kernels, but only 0.751507 ms of positioned-read wall-time intersection; blocking has zero observed intersection in both measures. Table 10 keeps these quantities separate. Let $\mathcal{R}$ contain allocation-plus-read worker envelopes, $\mathcal{P}$ the contained `pread` wall intervals, and $\mathcal{K}$ actual expert kernels identified by unique launch correlations. For half-open intervals on the common profiler clock, write $\mathcal{U}(X)$ for their merged union. We report $|\mathcal{U}(\mathcal{R}) \cap \mathcal{U}(\mathcal{K})|$ and $|\mathcal{U}(\mathcal{P}) \cap \mathcal{U}(\mathcal{K})|$, not intersections with a first-to-last kernel envelope.

Whole-request union duration (ms)	Blocking	Lookahead
Reader envelope	2,982.515538	2,989.969321
Positioned-read wall time	2,299.019830	2,295.077102
Actual expert kernels	194.741187	194.217266
Reader envelope ∩ kernels	0	174.409210
Positioned-read wall time ∩ kernels	0	0.751507

Table 10: One separately profiled QA request per schedule. Union lengths preserve gaps between kernels and do not represent physical storage-service time.

Each arm completes the same eight-ID QA output over nine decoder invocations, with 3,678 worker read envelopes, 22,068 contained positioned reads, and 33,403 attributed expert kernels.

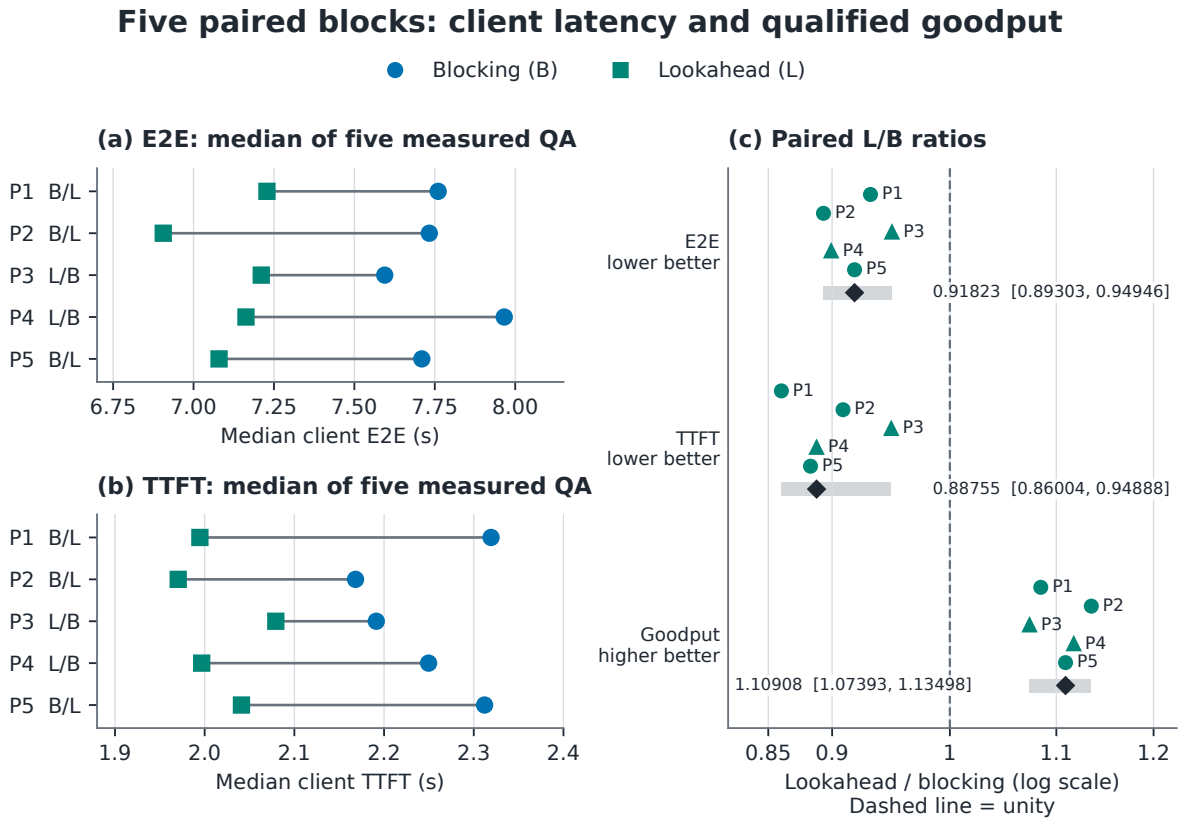


Figure 8: Paired client outcomes. Connected values are arm medians of five measured QA requests; goodput is 40 qualified tokens divided by summed client duration. Blue circles denote blocking and teal squares lookahead. Block order is retained. Diamonds show the exponentiated median paired log ratio; brackets show descriptive endpoints over all 3,125 ordered block resamples. The analysis unit is the paired block ($n = 5$), not the 50 measured requests. Startup is excluded, and serial-window goodput is not concurrent serving throughput.

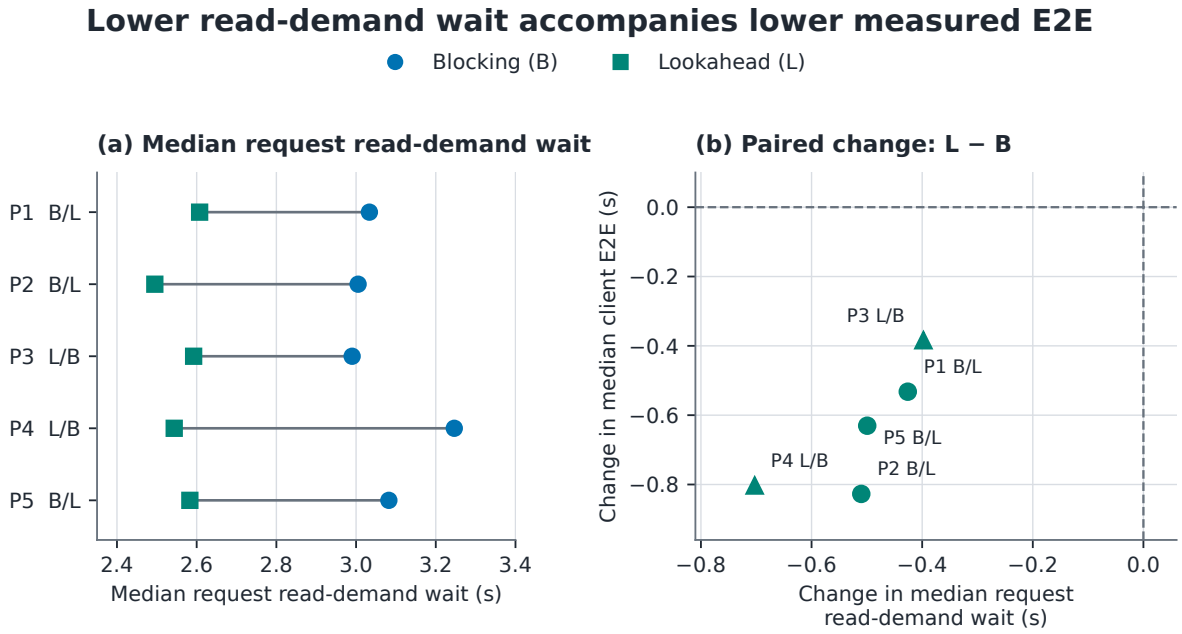


Figure 9: Host demand waiting and client-latency changes in the five paired blocks. Each arm summarizes five measured requests. The difference panel subtracts blocking from lookahead arm medians, rather than taking a median of requestwise differences. All five blocks show lower wait and end-to-end latency with lookahead. These host measurements establish an association, not a causal critical-path decomposition or hardware-overlap result.

Each records 11,571,399,936 source-verified requested bytes. The profiler's file-call records do not expose an interpreted byte count or per-call success value, so those requested bytes come from the operation ledger, not the profiling API. The read envelopes cover allocation and the source read call, excluding session opening and source verification.

The lookahead intersections are 5.83314% of the reader-envelope union and 0.0327443% of the positioned-read union. Target-process checks pass, but global trace completeness remains unverified: each capture retains three warnings concerning a non-CUDA wrapper, neither is accepted as a globally complete trace, and the dropped-event count is unknown. Observed launch/kernel correlations close without orphan or duplicate keys, which cannot exclude simultaneous loss of both records. These two profiled requests are not timing replicates and are not pooled with the five ordinary pairs.

6.4 RQ4: Abort, uncertain custody, and HTTP cancellation

The original component and cancellation experiments retain their own frozen configuration; they do not qualify later transfer-policy fault paths. Clean component failures preserve committed state and permit a checked retry; withheld completion proof instead retains device charges and blocks unsafe reuse. Separately, a real HTTP disconnect reaches a cancelled worker outcome, after which a new request in the same process matches its eight-token oracle. These results test different boundaries of the lifetime contract.

Real-CUDA component failures. Both blocking and lookahead preserve committed state after clean abort and retain device charges after completion-proof discard. These four component cases use a small fixture with two layers, four experts per layer, hidden and expert-intermediate dimensions of 16, and a vocabulary of 32. It uses FP8 weights with BF16 scales and the Mixed arithmetic profile. Its cache has one entry and a 65,536-byte logical limit; host staging has a separate 2 MiB limit.

The injected boundary follows a real routed-expert output submission. For the clean case, a scheduled software error is handled with valid completion evidence. For the Unknown case, the fixture first successfully waits for the actual compute event and upload synchronization, then deliberately discards that proof. This tests the response to missing authority, not a failed driver, hung device, or physically lost event.

After clean abort, committed recurrent-state and KV contents remain unchanged, transient resources are released, and retrying the same continuation succeeds before its test transaction is aborted. After proof discard, repeated abort and independent waits do not clear the poisoned owner. Its logical cache retains 6,918 bytes of the 65,536-byte limit: 774 pending-upload bytes and 6,144 scratch bytes, all classified within the retained Unknown subset. That subset is not an additional charge. Host cleanup reaches zero logical charge in both schedules, illustrating that host and device custody are different domains. Device state contents after Unknown are unobserved rather than assumed equal. Table 11 summarizes the distinction.

Full-model HTTP cancellation. The client closes its connection after receiving the first nonempty SSE content event. The matching worker terminal is observed 0.6680 s later with cancellation accepted and submitted, rather than natural maximum-length completion. The worker enqueues one token for the cancelled request. This deliberately censored request has no complete client usage or successful QA outcome. Only after cancellation acceptance does the client issue a new request and session in the same resident process. That request delivers all eight expected IDs,



Figure 10: Scoped read–compute overlap. The upper tracks show deterministic 20-ms excerpts starting at the first source-qualified worker read in each arm: reader envelopes, positioned-read wall intervals, and actual expert kernels. Start symbols do not enlarge durations. The lower panels use whole-request unions, not estimates from the excerpts. Lookahead intersects kernels for 174.409210 ms of reader envelopes but only 0.751507 ms of positioned-read wall time; blocking has zero observed intersection. Target-process checks pass, while wrapper warnings and an unknown drop count leave global completeness unverified. File-call wall time does not establish physical NAS-service overlap.

Experiment	Outcome	Resource/state evidence
Clean component failure, both schedules	Abort completes; same continuation retry succeeds	Committed state unchanged; transient charge released; 4,096-byte scratch reservation remains
Proof-discard component, both schedules	Unknown remains; unsafe reset, publication, and release rejected	Host charge zero; 6,918 device-cache bytes retained; post-Unknown contents unobserved
HTTP disconnect, complete model	Cancelled after first content; new request passes exact eight-ID check	Worker terminal observed in 0.668 s; final host charge zero; device residency retained

Table 11: Scoped failure results. Component faults are synthetic interventions after real CUDA work; HTTP cancellation is an ordinary client disconnect. Neither supplies a complete serving fault ledger or final physical-release proof.

text, usage, and terminal evidence. It is fresh-request recovery, not continuation of the cancelled request.

The owner records 5,116 completed and verified read operations with zero cancelled or uncertain reads. Thus cancelling the HTTP request does not imply that an individual synchronous source read was interrupted. At shutdown, logical host charge and observed pending operations are zero; device-cache charge is 209,739,776 bytes with residency retained. Per-request KV/frontier snapshots and a before/after pending-operation ledger are unavailable.

Across the ten paired serving lifecycles, the owners also report known local closure with retained residency and zero final logical read charge. Process drain and idle observations complete supervision, but do not certify final physical release. The component cache equations are not measurements of the runtime broker's five-class invariant in Equation 5; integrated serving conservation under faults remains unverified.

6.5 Reproducibility

The supplementary materials contain family-specific frozen configurations, request references, provenance inventories, public measurement projections, and offline checkers. They reconstruct all paired arms, token-to-response joins, logical counters, shared-client observations, and descriptive resamples, as well as the separate pressure, overlap, and cancellation results. Figure inputs derive from those records. Appendix A explains why earlier incomplete studies and unmatched supplements do not contribute to the reported denominators. The accompanying artifact guide maps each result to its inputs and checks, distinguishing public offline verification from rebuilding or rerunning with external artifacts. The current CPU figure renderer checks pinned canonical JSON, exact HTTP references and raw receipt hashes, the four acceptance stats snapshots, pair denominators, and the recorded rendering/font environment without changing the older five-pair renderer. Raw local traces, binaries, and checkpoint payloads are not made public merely by publishing their hashes.

Limits and conditional claims. The evidence supports these tested workloads and source-bound configurations, not a universal resource or performance guarantee. Logical charge is not physical peak memory, a fresh server is not cold NFS, and process drain is not final-worker-join or universal CUDA-release proof. Conditional lifetime proofs retain their stated source-identity and completion-authority assumptions. Component withheld-proof checks do not establish full-model Unknown recovery, and acceptance does not transfer to an unmeasured source configuration. Neither

unmatched resident observations nor the HTTP pilot close the matched-vLLM comparison gap or establish a state-of-the-art ranking. Other models remain untested.

7 Related Work

Storage-backed inference. FlexGen [3] distributes tensors across GPU, CPU, and disk and optimizes throughput-oriented schedules with compression and attention-cache management. ZeRO-Inference [6] describes CPU or NVMe weight hosting, layer streaming, prefetch, and release. PowerInfer-2 [5] targets smartphone inference beyond resident capacity using neuron clusters, segmented caching, and an I/O/computation pipeline. These systems establish weight streaming and transfer/computation scheduling as prior work. Ferrule focuses on the contract linking each materialization stage to its source identity, resource claims, completion evidence, and eventual reuse.

Demand-loaded expert caches. The inspected SGLang source includes a specialized expert-pack loader and runtime cache [2]. Cache misses issue storage reads into staging slots, followed by event-gated host-to-device copies; preparation of the packed representation does not eliminate those runtime reads. The loader and argument checks restrict this path to DeepSeek-V4 Flash and text-only Kimi-K3, require unit tensor, data, and expert parallelism, and disable CUDA graphs and shared-expert fusion. These are boundaries of the inspected path, not claims about every SGLang configuration or future version. We neither ran this path nor established its release inclusion, numerical qualification, or relative performance. Its source demonstrates that storage-backed expert loading is not unique to Ferrule; equivalent failure-custody guarantees have not been exhaustively audited in either direction.

Resident offload and state management. PowerInfer [4] exploits activation locality by executing hot neurons on a GPU and cold neurons on a CPU. This resident-domain hybrid differs from PowerInfer-2's storage pipeline. The inspected CPU-weight-offload paths in vLLM use host-resident weights, including unified virtual addressing; they should not be equated with disk-backed weight materialization. vLLM's PagedAttention [1] addresses serving KV memory through paged allocation. KV offload, whether to host memory or storage, concerns a different object from checkpoint-weight demand. A fair performance comparison must match representation, arithmetic, physical memory pools, workload, and quality criterion. The unmatched resident-system observations preserved with the artifacts do not satisfy those conditions and support no comparative speedup or ranking here.

Scheduling and lifetime. An asynchronous interface, an independent progress worker, and hardware overlap are distinct properties. The earlier read-scheduling comparison holds worker, representation, and resource limits fixed, then changes when one next-expert read is submitted. The transfer-policy comparison instead holds the read schedule fixed and changes upload grouping. Client timing measures the resulting behavior; correlated file-call and kernel intervals separately measure observed read–compute overlap. Selected-process API counts quantify observed synchronization calls, not a complete causal decomposition. Neither a worker, a broad host-range intersection, nor fewer synchronization calls alone establishes device overlap or faster serving.

Ferrule contributes a profile-scoped account of source binding, stage-local admission, publication, and retirement, together with measurements that keep logical and physical observations separate. We have not exhaustively audited the custody or failure semantics of the systems above and do not claim that they lack equivalent mechanisms. Model-registry entries and architectural capability are not substitutes for qualified execution: in particular, Ferrule has no tested DeepSeek-V4 serving

System	Resource focus	Principal mechanism
FlexGen	GPU, CPU, disk	Tensor placement, batching, compression
ZeRO-Inference	CPU or NVMe to GPU	Layer streaming and prefetch
PowerInfer	GPU–CPU residency	Hot/cold neuron execution
PowerInfer-2	Smartphone memory/storage	Neuron clusters and I/O pipeline
SGLang expert pack	Storage to GPU expert cache	Demand reads, staging, event-gated copies
vLLM	Serving KV residency; host-weight offload	PagedAttention; separate resident offload paths

Table 12: Mechanism-level positioning from the cited papers and inspected source. The SGLang entry concerns its specialized expert-pack path, not all adapters. These are not matched performance measurements.

result. The empirical contribution is controlled internal comparisons and their explicit resource and completion scope, not a first-system or state-of-the-art claim.

8 Limitations

Resource and deployment scope. The experiments bound selected logical accounts, not the complete physical working set. The host limit includes packed and prepared payloads, metadata, and requested worker stack; OS stack rounding, TLS, allocator overhead, weak-reference backing, and file-cache residency require additional headroom. Persistent-worker observations retain 10 MiB of logical metadata and stack charge while the owner is idle and alive. Zero joins in those snapshots are not evidence of a missing required live join or of a completed final-owner join; the latter witness remains unobserved. Device samples can miss transient peaks and lack exact process attribution. The three cache points vary entry and byte limits together, so they neither isolate a replacement-policy effect nor locate a physical capacity frontier. The complete compact checkpoint may fit in host memory. Logical reads do not establish physical NAS traffic or physical storage necessity.

Numerical and model scope. The evaluated Qwen3.5-35B-A3B-FP8 text profile decodes storage values to BF16 execution weights and uses explicit F32 accumulation and state boundaries. Its numerical and HTTP checks do not establish full-model W8A16 quality, stock dynamic-FP8 behavior, or same-precision equivalence to another engine. DeepSeek-V4 end-to-end execution remains untested. Other precisions, adapters, topologies, long contexts, shared-memory devices, and multimodal workloads remain outside the evaluated profile.

Statistical and measurement scope. The earlier read-scheduling study has five paired blocks, not 50 independent request repetitions. The transfer-policy pilot has two pairs per read schedule; its 40 measured requests are nested observations, not 40 independent policy replicates. Schedules, source cuts, and experiments are not pooled. Shared page and NFS caches can carry state across fresh servers. The older descriptive resampling ranges provide neither guaranteed coverage nor a significance or power result; no intervals or significance claim are made for the newer pilot. Older serial-window goodput excludes inter-request gaps, whereas the newer phase-qualified goodput includes them between the first measured dispatch and last measured terminal. Neither measures concurrent-server capacity or startup-inclusive throughput. Client time to first nonempty content is

not first-device-token latency. Instrumentation overhead remains unmeasured.

The read–compute overlap study concerns two separate requests. Its scoped checks pass, while wrapper warnings and an unknown dropped-event count leave full-trace completeness unverified. The grouped-transfer profiles are also partial selected-process observations: qualified pairs exist for both read schedules under separately disclosed collection controls, but full-trace acceptance is false and event loss remains unknown. Failed captures retain their failed status. API and interval sums may overlap and cannot be added into an exclusive explanation of elapsed time. Neither fewer synchronization calls nor the retained-upload pin/free overhead identifies the exact cause of unprofiled timing differences.

Failure and verification scope. The component experiment runs real CUDA operations but injects a software failure or discards proof after valid fences. It does not exercise a driver failure, lost hardware event, or full-model uncertain completion. The HTTP experiment tests a client disconnect followed by a new request, not resumed execution or recovery from a poisoned owner. Neither experiment establishes the integrated five-class conservation invariant across a serving request. Logical host/cache observations do not supply that missing ledger. Final physical release remains unobserved; process drain, device idle, and local lock release are not universal CUDA-release certificates. Unknown custody can reduce availability indefinitely; the design supplies no general in-place reset, durable crash recovery, or arbitrary cross-rank atomic commit.

Reproducibility scope. Public archives support offline reconstruction of the earlier paired results and selected figures. The newer worker and transfer summaries instead retain bounded derived evidence and hashes of local originals; their raw reports, traces, binaries, and build inputs are not a complete portable public archive. Receipt checks do not replay private authentication or recover omitted native dependencies, model payloads, or live mapped-memory contents. Integrity hashes are not external attestations. Historical rejected or incomplete studies remain separate rather than becoming successful repetitions, and measured source cuts do not automatically qualify a later cleaned implementation.

9 Conclusion

Bounded storage-backed weight residency requires a lifetime contract as well as a cache policy. Ferrule connects source-bound planning and stage-local admission to exact completion, publication, and ownership through retirement. The implemented expert path combines a persistent bounded read worker, optional lookahead, and grouped synchronous uploads while preserving a synchronous forward interface. Quality-checked serving and component experiments support this architecture within one explicit numerical profile and selected logical accounting domains.

The measurements also show why mechanism and performance claims must remain separate. The earlier five-pair experiment isolates read scheduling; newer whole-test and HTTP comparisons isolate transfer policy within each read schedule. In the HTTP pilot, grouping improves observed latency and phase-qualified goodput with lookahead but regresses time to first content with blocking reads. A retained-upload alternative has a substantial negative whole-test result. Partial profiles expose synchronization and host API costs without proving a complete critical-path explanation. These findings support conditional engineering choices, not universal superiority. Integrated fault conservation, a physical capacity frontier, final physical release, and broader model generality remain outside the demonstrated result.

References

[1] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient Memory Management for Large Language Model Serving with PagedAttention, 2023. arXiv:2309.06180.

[2] SGLang contributors. SGLang expert-pack runtime, loader, and configuration checks. Official source repository, 2026. Source snapshot; accessed October 18, 2026.

[3] Ying Sheng, Lianmin Zheng, Binhang Yuan, Zhuohan Li, Max Ryabinin, Daniel Y. Fu, Zhiqiang Xie, Beidi Chen, Clark Barrett, Joseph E. Gonzalez, Percy Liang, Christopher Ré, Ion Stoica, and Ce Zhang. FlexGen: High-Throughput Generative Inference of Large Language Models with a Single GPU, 2023. arXiv:2303.06865.

[4] Yixin Song, Zeyu Mi, Haotong Xie, and Haibo Chen. PowerInfer: Fast Large Language Model Serving with a Consumer-grade GPU, 2023. arXiv:2312.12456.

[5] Zhenliang Xue, Yixin Song, Zeyu Mi, Xinrui Zheng, Yubin Xia, and Haibo Chen. PowerInfer-2: Fast Large Language Model Inference on a Smartphone, 2024. arXiv:2406.06282.

[6] ZeRO-Inference: Democratizing massive model inference. DeepSpeed technical web page, September 2022. September 9, 2022; technical web page, not a conference paper.

A Additional Quality Detail and Separate Evidence

A.1 Original numerical coverage

For each retained-fresh and ordinary schedule in Table 7, common-history and free-feedback streams each cover 27 calls and 9,187,840 vocabulary values; reset covers nine calls and 3,228,160 values, and fork covers one call and 248,320 values. For each cold-owner schedule, common history, feedback, and reset each cover nine calls and 3,228,160 values; fork covers one call and 248,320 values. Fingerprints identify compared outputs rather than adding logit observations. Ordinary execution does not observe diagnostic route or state fields. The fork check replays a same-history committed frontier; it does not test divergent KV copy-on-write or broad task accuracy.

The pressure study supplies narrower evidence at its three new cache points: complete visible outputs and actual generated IDs. It does not establish full-vocabulary, route, or internal-state equivalence at those settings. Likewise, the tiny component fixture tests lifecycle behavior, not the numerical quality of the complete 35B model.

A.2 Incomplete and non-pooled observations

Earlier rejected studies remain excluded from successful-result denominators, with their original outcomes retained in the supplementary materials. Table 13 summarizes the exclusions by scientific scope rather than by implementation chronology. Later corrections establish new measurements; they do not retrospectively complete a missing request or turn a rejected capture into an accepted repetition.

Separate evidence	Observed outcome	Treatment
Two earlier paired studies	Each attempts one lifecycle but completes none; all 120 HTTP requests and 50 measured QA slots unissued	Each retains five incomplete pairs and no aggregate
Earlier pressure study	Three requests complete, one is censored, eight unissued; only two complete HTTP canaries	No qualifying cache point; raw worker IDs do not repair missing HTTP termination
Pressure admission rejection	No server starts; all 12 requests unissued	No latency, memory, or I/O observation
Earlier profiler captures	Two collected captures rejected; other preparations stop before collection	Neither contributes an accepted arm; only the two fresh scoped captures support RQ3
Earlier numerical configurations	Initialization failures and observation-policy rejections coexist with some numerically accepted tests	Outcomes stay configuration-specific and are not added to current tests

Table 13: Non-pooled evidence. Missing execution or observation is not reported as zero cost or numerical failure. Detailed identities and diagnostics remain in the artifact documentation.

A.3 Earlier logical-demand observation

A separate, heavily instrumented single-lifecycle observation with the read pipeline disabled completed four serial requests. One was a two-token length/usage smoke; three passed visible-text and eight-token usage canaries. Generated IDs, logits, routes, and states were not observed. Its 30 decoder invocations recorded 36,240,064,128 logical reader bytes, 36,240,248,688 upload-interface

bytes, 11,519 expert uploads, 11,455 evictions, and zero expert hits. The final observed cache held 64 entries. These values describe an earlier logical-demand workload, not additional repetitions of either schedule in the paired experiment.

At the observed invocation boundaries, entry counts reconcile as

$$N_j = N_{j-1} + I_j - E_j, \tag{12}$$

where N_j is resident entries and I_j, E_j are installation and eviction deltas. The initial zero is implied by the first boundary, not directly measured. This count identity does not establish byte conservation or continuous occupancy. Historical plots and the detailed timing and memory samples remain in the supplementary materials rather than extending the current performance evidence.

A.4 Excluded unmatched resident-system observations

Separate tensor-parallel resident observations with vLLM use unmatched arithmetic and working-set budgets. They are retained only in the artifact documentation and are excluded from all demand-serving denominators and comparative speedup claims. Neither the current HTTP pilot nor those observations close a matched-baseline gap or establish a state-of-the-art ranking.

The independent W8A16 strict numerical comparison fails logits, ordered routes, and state checks despite argmax agreement. That negative result remains a failure of its own quality contract and cannot be replaced by the successful Mixed-profile observations.