Ferrule: Bounded Storage-Backed Weight Residency

Lu Huanbing
Independent Researcher
luhuanbing084@gmail.com

Working draft — not a submission

Abstract

Storage-backed inference needs bounded materialization and safe reuse, not merely a small device cache. Ferrule connects source-bound plans, stage-local admission, exact completion identities, and physical custody to bounded weight residency. Its next-expert pipeline and safe blocking control use one positioned-read worker, a depth-one queue, identical host/device limits, and the existing synchronous forward. Lookahead moves the next selected read ahead of current compute submission; GPU overlap requires correlated kernel activity, not host spans. Nine exact-once GPU qualification gates pass on one frozen source cut, covering constructor/provider cells and both schedules' full40, retained-fresh, ordinary and cold-owner factory paths. Each full40 gate compares 1,986,560 vocabulary values; each ordinary run matches all 64 actual output fingerprints of its same-schedule diagnostic primary, without observing routes or states. Observed logical host charge and allocator-accounting high-water remain within the declared limits, not certified physical peaks or release proofs. Separately, one historical DEBUG run records four HTTP successes, three limited canaries and 36,240,064,128 logical reader bytes; it is not current CLI quality or scheduling evidence. We separate factory qualification, bounded accounting, I/O observability and physical custody. Capacity-frontier, CLI/HTTP, paired performance and GPU overlap qualification remain open; physical-memory necessity, nonblocking forward and performance leadership are not claimed.

1 Introduction

Storage-backed weight residency lets an inference owner repeatedly materialize selected weights into a bounded execution cache. The immediate problem is to bound those copies and reuse their capacity safely, whether or not the complete compact checkpoint could fit in host memory. Exceeding a device's memory, exceeding aggregate physical residency, and choosing a controlled cache budget are different conditions. A checkpoint's file size, an expanded tensor's size, and a request's working set are not interchangeable. The working set also includes activations, model state, staging, outstanding transfers, allocator/driver reserves, metadata, and resident file-backed pages.

Ferrule studies bounded storage-backed residency through *plan, admit, materialize, publish, execute, retire*. A source-bound plan fixes what is read; admission bounds covered claims; completion validates the requested generation; custody prevents premature reuse. Its concrete next-expert pipeline separates the read window from owner compute submission. A safe blocking control waits on the same worker immediately, whereas lookahead submits the next selected read after current upload and before current compute submission. The whole forward remains synchronous. These contracts make scheduling and resource reuse testable; they do not themselves prove storage necessity, hardware overlap, model quality, or a speedup.

Research questions. We organize the paper around four questions, each with a concrete observable:

RQ1 – capacity. Can a complete, quality-accepted request finish with bounded repeated materialization while the measured working set stays within declared per-domain limits? A successful configured-budget run is not, by itself, evidence that the model exceeds physical memory.

RQ2 – request I/O. What request-time demand is visible at the logical-reader and NFS-client levels, and what attribution or backing-store traffic remains unknown after separating startup and cache effects?

RQ3 – safe scheduling. Under the same representation, quality contract, and budgets, what differs between the implemented safe blocking and bounded lookahead schedules in demand waiting, speculative work, and correlated read/transfer/kernel activity?

RQ4 – custody and failure. Do exact identities, credits, and custody prevent stale publication and premature reuse, while retaining capacity when completion is Unknown and releasing it only after evidence of quiescence?

Contributions and boundaries. This manuscript makes four scoped contributions:

1. a resource model that separates disjoint physical domains, configured accounting domains, representations, and the measured time-varying working set;
2. a normal-path contract that connects source-bound planning and stage-local admission to generation-qualified completion, publication, and retirement;
3. structured success/terminal-custody algorithms and an implemented bounded next-expert schedule with an identical-budget blocking control; and
4. an evidence-oriented evaluation layout that ties each RQ to tables, traces, memory peaks, and plots rather than promoting design intent to a result.

The historical serving evidence is deliberately narrow. TP1 attempt `-O2` is one DEBUG-instrumented exploratory run ($n = 1$): four HTTP requests completed, the three limited visible-text/eight-output-token-usage canaries passed, and Hello2 passed HTTP and usage checks only. It records 30 decoder invocations and 36,240,064,128 logical checkpoint-reader bytes; these are logical bytes, not physical NAS/network or block-device reads. Generated IDs, full numerical quality, process-owned GPU attribution, and a complete resource contract are absent or unknown. Attempt `-O1` failed before server launch and supplied no request or read result. Consequently, TP1 does not establish RQ1 capacity, storage necessity, RQ3 scheduling, or RQ4 custody closure.

Separate TP2 receipts include numerical/lifecycle observations for a test binary and descriptive resident comparisons. They are not interchangeable with TP1: the TP2 test binary is not a public serving binary, the arithmetic and topology contracts are profile-specific, and the two resident comparison blocks are not a fair speedup proof. In particular, the vLLM supplement is not evidence of storage necessity or a matched total working-set budget. Both opt-in schedules now have same-cut real-GPU factory qualification: nine exact-once gates pass on `033519c`. This closes their selected numerical/provider gate, not CLI/HTTP, a capacity frontier, paired performance, or physical release. Historical failures and wrapper outcomes remain separate (Appendix A); no old PASS is borrowed. Multimodal, unified-memory, universal-registry, ACID, and default NCCL/device-collective claims

Symbol	Meaning and units
π, r	Execution/quality profile and weight representation/placement
$S_{\text{ckpt}}, W_{\text{compact}}, W_{\text{full}}(r)$	Checkpoint, complete compact, and full-resident weight bytes
$d \in \mathcal{D}_{\text{phys}}, B_d$	Disjoint physical domain and available physical byte capacity
$a \in \mathcal{D}_{\text{acct}}, b_a$	Covered accounting domain and configured limit in that domain's unit
$M_d(t), M_d^{\text{peak}}$	Deduplicated physical working-set bytes at time t, and its true peak
$\widehat{M}_d(t_i), \widehat{M}_d^{\max}$	Reported sample and maximum over stated sampling times; not a true peak
$U_a(t), c_a$	Charged occupancy and additional proposed claim, in domain a's unit
$\delta, \rho_a, b_a^{\text{eff}}(\delta)$	Demand class, execution reserve, and its policy-accessible accounting ceiling
F_a, R_a, I_a, L_a, T_a	Free, reserved, provider-in-flight, live, and retained claims
$P = (K, \boldsymbol{q}, \boldsymbol{\ell})$	Frozen plan: exact key, stage claim vectors, resident claim vector
K, g_s, g_x	Full resource key, source generation, destination reuse generation
o, e, J	Owned operation, completion evidence, provisional owner journal
$\Omega, k, w, \mu_o[s]$	Post-detachment worklist, per-pass visit limit, persistent cursor, and exact-stage cancellation marker
$\mathcal{V}_q, L_q, H_q, \Delta_q$	Request q's invocations, logical reader bytes, facade H2D bytes, and $H_q - L_q$
N_j, I_j, E_j	Expert entries at invocation boundary j, installs/uploads, and evictions
$X_j, Q_j, S_j, C_j^{\text{cache}}, Z_j$	Resident, pending-upload, scratch, cache charge, and retained Unknown subset gauges
$C_{\text{peak}}^{\text{cache}}$	Lifetime cache-accounting high-water; not a total physical-owner peak
$P_q^{\text{net}}, P_q^{\text{blk}}$	Physical network and backing-block bytes attributed to request q
$b_{\text{cache}}, b_{\text{scratch}}$	Configured cache-credit and scratch limits; neither is a physical-domain bound

Table 1: Notation used by the contracts, algorithms, and evaluation. Accounting limits are never substituted for physical capacity.

are outside the demonstrated evidence. Storage inference and prefetch are prior art; no first-system or state-of-the-art performance claim is made.

The rest of the paper follows the execution argument. Section 2 defines representations, physical and accounting domains, quality, and failure contracts. Sections 3, 4, and 5 follow a normal profile, state its invariants, and give structured algorithms. Section 6 identifies the code paths and their limits. Section 7 starts with the frozen setup and then reports RQ1–RQ4 in order. Sections 8, 9, and 10 place the design in prior work, collect the remaining evidence gates, and conclude.

2 Resource Model and Execution Contracts

2.1 Representations and one notation

The target is a complete request or forward output under a declared model, arithmetic, placement, and quality profile π. Text decoding is the current instance: prefill, multiple decode steps, and KV/state lifetime. Vision, audio, and combined modalities require their own complete-output contracts and remain future implementation and evaluation work.

A representation r identifies payload, decoding metadata, expansion, and required physical copies. Checkpoint bytes S_{ckpt} include the specified files and container layout. Complete compact weights W_{compact} include all encoded weights, scales, and decoding metadata. Full-resident weight demand $W_{\text{full}}(r)$ is the footprint of all weights under representation and placement r; it is weight-only, not the complete inference working set. Expanded execution weights may be larger than compact weights and need not all be live at once. Aliased views are not extra copies; genuinely copied payloads are.

2.2 Physical capacity and measured working set

Physical capacity counts disjoint resident-memory pools once:

$$B_{\text{mem}} = \sum_{d \in \mathcal{D}_{\text{phys}}} B_d. \tag{1}$$

For a declared representation and placement, consider

$$W_{\text{full}}(r) > B_{\text{mem}}. \tag{2}$$

The inequality rules out aggregate full residency of *that* representation/placement, not streaming feasibility. Disk is excluded. Discrete host and GPU pools may be added; host and accelerator aliases of one physical pool may not. Unified virtual addressing over separate physical pools does not merge those pools. Report installed capacity, fixed exclusions, co-tenants, copies, and placement/migration policy for each evaluated domain. Per-domain accessibility and placement remain constraints even when the aggregate sum is sufficient.

Define $M_d(t)$ as physical bytes resident in domain d that are attributable to the selected inference deployment at time t. Count resident weights, expansion/workspace, activations, model state, staging, outstanding and retained buffers, metadata, allocator/driver reserves, and attributable file-backed pages, deduplicating shared physical allocations. A full attribution requires fixed and external consumers to be measured or excluded explicitly; a process heap or device counter alone does not supply it. The desired feasibility condition and the reported sampling statistic are different:

$$\begin{aligned} M_d^{\text{peak}} &= \sup_{t \in \mathcal{T}} M_d(t) \le B_d \quad \text{for each } d, \\ \widehat{M}_d^{\max} &= \max_{t_i \in \mathcal{T}} \widehat{M}_d(t_i) \qquad \text{over the declared samples.} \end{aligned} \tag{3}$$

$\mathcal{T}$ is the declared startup/request/cleanup interval. A sampled maximum can miss a transient peak; selected-device telemetry also includes unattributed processes, and tree PSS does not cover every machine consumer. Thus those proxies are not silently identified with $M_d(t)$. Independent component maxima must not be added into a simultaneous working-set peak. RQ1 uses Table 5; the storage-demand figures in Section 7.3 diagnose cache and transfer counters without turning sample maxima into certified physical bounds.

The stronger, optional storage-necessity claim additionally asks whether *all compact weights and metadata* can remain resident with online expansion under the quality contract. It is not a prerequisite for the bounded-residency mechanism studied here. Failure to fit expanded weights is insufficient if that compact alternative fits. A low configured cache budget is likewise not physical host-memory insufficiency. Storage necessity is scoped to evaluated representations, not a proof against every possible compression scheme.

2.3 Accounting and admission contract

A covered accounting domain a may count bytes, slots, pages, or queue entries. Its limit b_a is configured, not installed physical capacity; accounting domains can overlap or omit parts of a pool and cannot in general be summed. A cache limit covers its cache, not the whole working set. The necessary covered-domain capacity condition is

$$\forall a \in \mathcal{D}_{\text{acct}}, \qquad U_a(t) + c_a \le b_a. \tag{4}$$

Here c_a is the additional charge to acquire, not a second charge for an already held grant. A revisit preserves that grant and rechecks only unmet claims. The inequality is not sufficient for admission: the owner must also validate canonical claims and units, arithmetic overflow, source/placement/backend compatibility, actual reservation and grant authority, and the selected demand policy. With execution reserve $0 \le \rho_a \le b_a$, the policy-accessible ceiling is

$$b_a^{\mathrm{eff}}(\delta) = \begin{cases} b_a, & \delta = \mathsf{Required}, \\ b_a - \rho_a, & \delta \in \{\mathsf{Warmup}, \mathsf{Prefetch}\}. \end{cases} \tag{5}$$

A non-required demand cannot consume the execution reserve even if Equation 4 holds. Required execution may use that reserve, not exceed the hard ceiling. The reserve protects accounting capacity from speculation, not eventual progress against a stalled provider.

The abstract admission result has three cases. Rejected includes invalid plans and an intrinsically impossible stage: for example, a required claim exceeding b_a, or a non-required claim exceeding its unchanged policy-accessible ceiling even with no other occupancy. Waiting cannot fix such a stage; it needs rejection or an explicitly changed, revalidated plan or demand policy. Wait means the otherwise admissible stage is blocked by current occupancy, including reserve-protected capacity. It retains existing claims and registers a revisit condition rather than recharging or redispatching them. Admitted requires all authority, compatibility, and policy checks as well as capacity. These are explanatory result classes, not a new backend API, and do not imply that all future stages can be admitted atomically or that their resource waits are acyclic.

Each claim has an authoritative owner, unit, and end condition. Canonicalizing same-kind claims is not deduplication of arbitrary aliased allocations. A bounded request channel does not bound all downstream buffers. Map claims to the actual ledger and physical allocations on the selected path; do not infer one global broker from a common interface.

The plan P binds K, read/upload/install claim vectors $\boldsymbol{q}$, and persistent residency $\boldsymbol{\ell}$. K includes model instance, source identity, semantic content descriptor, encoding, backend/device, and g_s, g_x. Alignment and multiple extents may change staged byte counts; expansion may change execution bytes. The destination authority supplies the actual reservation and g_x before hard admission. Publication consumes validated bindings, not requested addresses. Algorithm 1 uses the capacity and policy checks in Equations 4 and 5 before each physical dispatch.

OS page cache, storage-server cache, metadata, and driver reserves are not implicitly under runtime credit control. Resident file-backed pages count as memory; the backing file adds no capacity. Swap must be disabled or explicitly accounted for. A logical read, page fault, or migration between physical pools is not evidence of a physical storage read. RQ2 consequently distinguishes $L_q, H_q, \Delta_q, P_q^{\mathrm{net}}, P_q^{\mathrm{blk}}$ rather than translating them into one bandwidth number.

2.4 Quality, completion, and progress contracts

Protocol exactness means compatible sharing, exact completion identity, authorized publication, and no reuse before physical references end. It is separate from numerical/task quality. A quality contract fixes complete outputs, independent reference, arithmetic/state precision, inputs, and acceptance thresholds. Visible text, usage, generated IDs, full-vocabulary logits, ordered routes, and internal states are different checks; argmax agreement does not establish numerical agreement. Quantized checkpoint storage alone does not specify kernel or KV arithmetic.

Completion evidence must name its operation, full key, submitted stage, generations, outcome-specific byte count, and any required physical fence. Rejecting an event with mismatched identity, stage, generation, or byte contract is not the same as accepting a matching event whose validated

`CompletionOutcome::Stale` is a non-success result. The former supplies no mutation/release authority; the latter selects terminal cleanup of its named operation, never readiness or publication of a newer generation. Returning named read custody is separate from releasing the read-slot and storage-read claims; pinned host capacity remains charged through upload or proven cleanup. Publication preflight similarly distinguishes Ready, Blocked by live replacement custody, and Rejected by an invalid binding/conflict. Rejected publication retains its journal and remaining custody for cleanup.

Only validated producer success proves readiness of that stage, not quiescence of a later consumer. Only owner-accepted cancellation selects a canceled caller outcome and detaches demand; a cancellation request or provider acceptance alone does not prove physical completion. Shared loads retain their remaining owners. Pending/Idle is retained as Pending only for a positively known original pollable obligation; otherwise it retains Unknown custody. Unknown means remaining physical references cannot be proved. Unknown retains allocation and credits, possibly indefinitely. A late exact scoped completion can close its named registry obligation but cannot clear physical-owner poison or quarantine. Bounded cleanup passes retain unresolved obligations and a traversal cursor across wakes and revisits, so the same first items do not hide unvisited obligations. Ambiguous cancellation replies remain non-retryable without exact owner evidence; they are not treated as known non-acceptance. Neither that scheduling contract nor Unknown retention promises eventual reclamation or availability.

Conditional liveness requires feasible per-domain resources, no cyclic resource waits, fair progress and completion handling, live providers, and eventually observable completion. A blocking provider or stalled progress worker can violate these conditions. An asynchronous API, event-driven progress, and hardware overlap are separate properties. Neither algorithms nor current receipts establish global nonblocking execution, lock freedom, durable crash recovery, or arbitrary multi-owner atomic commit.

3 A Normal Profile Walkthrough

Consider a profile π with a nonresident weight dependency. The following walkthrough composes protocol roles; it is not a claim that every TP1 runner uses the general registry. It connects the resource model to the algorithms and to the observables evaluated in Section 7.

Plan and admit. The adapter binds source descriptors, representation, placement, and workspace. The destination authority supplies the actual reservation and reuse generation. The resulting $P = (K, \boldsymbol{q}, \boldsymbol{\ell})$ is frozen before hard admission. Equation 4 is a necessary covered-domain check, not complete admission or a sum of heterogeneous budgets. The stage also passes authority, compatibility, intrinsic-capacity, and execution-reserve checks in Section 2. An impossible stage is rejected; only otherwise admissible demand blocked by current occupancy waits. Before preparation can submit physical work, the owner records provisional state and cleanup intent in journal J.

Read, transfer, and install. A new load admits read/staging claims and submits once. Exact read completion returns named read custody. On read success, the owner then explicitly releases held ReadSlot and StorageReadBytes credits while retaining the pinned buffer and PinnedHostBytes charge for upload. Custody return alone is not that credit discharge. Upload needs its own admission and retained fence; its completion ends DMA custody but not destination ownership. Installation admits its own claims. Replacement also needs evidence that old consumers have stopped using

the old binding. An ambiguous submission retains its ticket, allocations, and claims even when no public handle is returned.

Publish and consume. Install success is validated against operation, key, stage, bytes, and both generations. Publication preflight checks the actual binding and replacement rules. Ready permits resident-grant transfer and exact waiter resolution; Blocked waits for replacement leases/custody to end; Rejected preserves the original journal, grants, and physical binding for terminal cleanup instead of publishing or refunding residency. Only Ready retires the load operation by transferring its resident grant. This is not eviction: resident bytes remain charged. A consumer revalidates its dependency and lease epochs and obtains resume admission before executing. A compatible resident hit skips materialization; a compatible in-flight load can be joined only on paths that implement sharing.

Bounded next-expert scheduling. On the concrete routed-expert path, the owner first uploads the current expert and retires its temporary host payload aliases. It can then submit one read for the next selected, nonresident expert before submitting current compute. A read-only residency hint neither pins that next expert nor changes cache replacement order. Blocking uses the same worker and waits at demand instead. This split is a scheduling candidate, not a measured GPU overlap or an asynchronous forward API (Algorithm 3).

End and retire. Consumer completion returns only the custody discharged by its adapter. Ending/publication and post-ending cleanup are separate phases. Prepared-KV owners require exact install acknowledgments before logical publication and retirement acknowledgments before displaced pages are reused. Independent cache pins and persistent state survive request completion. Cancellation changes the caller outcome, not these end conditions; Unknown may retain capacity indefinitely.

The walkthrough supplies three different evidence gates: startup final-claim completion, materialization producer/replacement completion, and consumer/ ending completion. None substitutes for another. The experiment should record those transitions alongside complete-output quality, physical memory samples, and request I/O. The arrows in Figure 1 are conceptual dependencies, not measured hardware overlap.

4 Design Invariants

These are intended contracts for the implemented boundaries, not completed formal verification or independently measured properties of every profile.

I1: bounded, uniquely owned claims. Admitted claims must preserve $U_a + c_a \le b_a$, together with the demand policy and authority checks; this inequality alone is not admission. For each accounting domain, the intended semantic partition is

$$\begin{aligned} b_a &= F_a + R_a + I_a + L_a + T_a, \\ U_a &= R_a + I_a + L_a + T_a, \qquad F_a = b_a - U_a. \end{aligned} \tag{6}$$

Each claim occupies exactly one class: free, reserved, provider-in-flight, live, or retained. Different claims of one operation can occupy different classes. Unknown is classified in T_a, not charged again as $I_a + T_a$. Returning provider custody does not necessarily decrease U_a: an explicit release or

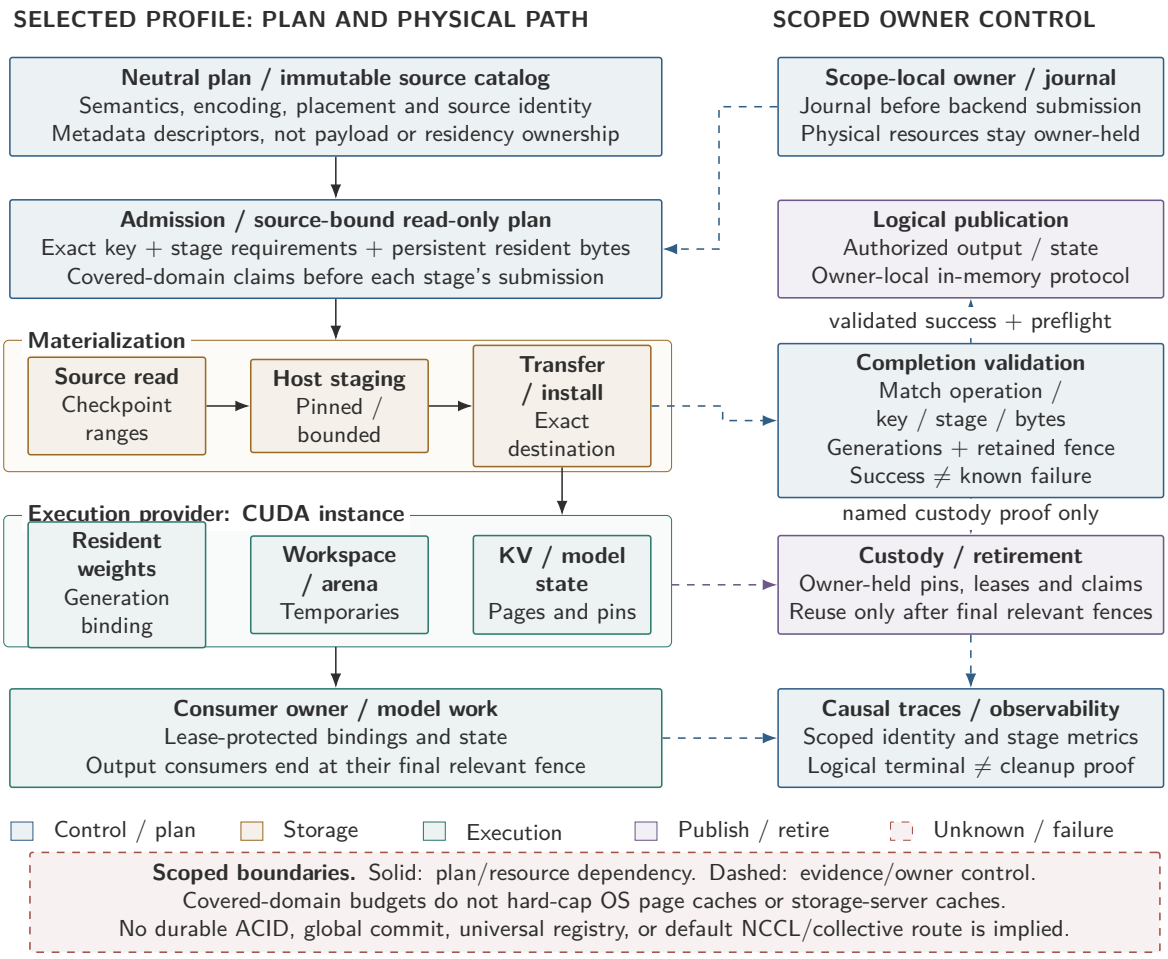


Figure 1: Conceptual, profile-scoped architecture. The shared palette uses blue for control/planning/validation, ochre for storage/materialization, teal for execution/provider resources, violet for publication/custody/retirement, and brick for Unknown/failure. Named nodes and labeled swatches carry the meaning independently of color. Solid arrows denote dependencies; dashed arrows name evidence or owner-control gates. Source planning and covered-domain admission precede materialization; providers retain physical resources. CUDA is one provider instance. No universal registry, OS-cache hard cap, durable/cross-rank ACID, or default NCCL/device-collective route is implied.

ownership transfer is separate. The broker records capacity, occupancy, authenticated owner, and a submitted bit, not five literal bucket fields. Equation 6 is an intended invariant; the current TP1 receipts do not independently measure its partition or final closure. RQ4 specifies the missing owner-ledger and fault evidence.

I2: exact evidence before mutation. A completion must match o, K, submitted stage, g_s, g_x, and outcome-specific bytes. Success requires the full expected amount; partial failure is not readiness. Identity-stale, retired, unowned, wrong-stage, or generation-mismatched events are rejected before mutation and cannot release a newer generation's grants. In contrast, a matching completion whose validated `CompletionOutcome::Stale` reports non-success ends the named stage according to its proof and selects terminal cleanup; it is not an identity rejection and does not authorize publication. Waiter/request/dependency epochs and mapping epochs have different scopes; a wake notification is not validation of any of them.

I3: custody follows physical work. Every outstanding physical reference is required to remain reachable from an authoritative owner, including on submission, notification, preparation, or cleanup errors. Read return, upload-fence return, resident transfer, and consumer retirement close different obligations. An address descriptor is not allocation ownership or a completion certificate. A read snapshot is not mutation authority.

I4: publication and reuse have different proofs. Producer success makes a candidate ready. Publication additionally checks the actual binding and authorization; old-consumer quiescence and retirement permit replacement reuse. Typed committed-KV sources and append witnesses are revalidated against the current page authority, and zero-reference pages are reused only after retirement confirmation. Several participant ACKs establish this scoped lifecycle order, not arbitrary distributed atomic commit.

I5: terminal demand does not erase obligations. Owner-accepted cancellation detaches only its demand, preserving other waiters and persistent owners. An unsubmitted shared load is preserved under its remaining authority or handed off with acknowledgment, never discarded because one demand ended. Provider cancellation acceptance is recorded once per exact operation/stage/generation; it is not a quiescence certificate. It cannot reverse an irreversible install/publish intent. Pending cleanup polls the original identity and arguments without redispatch. Unknown retains physical payload and credits and preserves applicable owner poison or quarantine. A matching later registry completion may close a named stage; it cannot heal a separately poisoned physical owner. No general in-place reset/ replay or unconditional eventual reclamation is asserted.

The testable consequences are no stale publication, no duplicate dispatch, no premature slab/page reuse, and retained capacity when proof is missing. RQ1 checks bounded completion, RQ2 checks the load demand, RQ3 checks progress and physical intervals, and RQ4 checks these invariants under faults. A process exit or idle sample does not independently prove any final owner ledger.

5 Lifecycle and Bounded Scheduling Algorithms

The pseudocode names protocol roles and abstract result classes, not new provider APIs. Its records and journal persist across calls. AdmitStage validates intrinsic feasibility, canonical claims, authority, compatibility, and demand/reserve policy before acquiring that stage's grants. An added claim $c_a > b_a^{\mathrm{eff}}(\delta)$ is intrinsically impossible under that policy; an otherwise feasible claim blocked

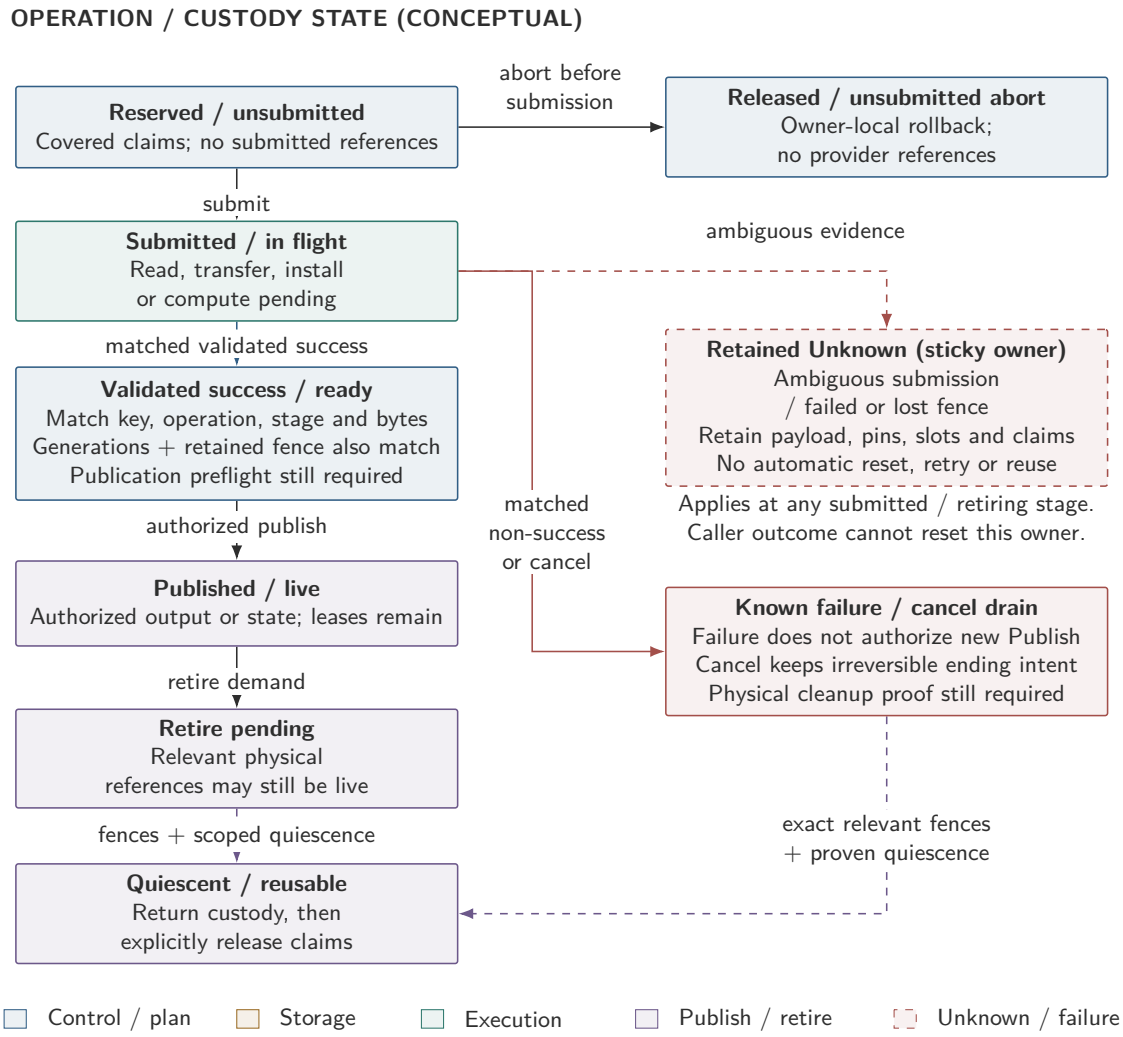


Figure 2: Conceptual operation/custody states, not one Rust enum. Labeled roles retain the shared blue planning/validation, ochre storage, teal execution, violet publication/retirement, and brick Unknown/failure palette. Only exact validated success plus publication preflight authorize new publication. Known non-success drains the named stage, not readiness; cancellation keeps irreversible ending intent. Dashed-border physical-owner Unknown retains payload, pins, slots, and claims and has no automatic recovery/reuse edge. Closing a matching registry-stage obligation does not heal a separately poisoned owner. Logical outcomes and charged resources are separate axes.

by current $U_a + c_a > b_a^{\text{eff}}(\delta)$ waits. The authoritative owner rechecks capacity on acquisition, without implying a cross-owner atomic reservation. Rejected means no further dispatch under the unchanged plan; Wait preserves already held custody and arms a budget/owner-progress revisit. RetainCleanup records the primary reason and original payload, then selects the cleanup protocol in Algorithm 2; it does not recreate a ticket or discard other demand. Publication Rejected is the abstract class for preflight validation/conflict errors, distinct from a live-custody Blocked result. Rejection after install does not undo the physical installation. The owner journal records applied custody returns and releases so later cleanup advances only residual obligations, without applying a return twice. Cleanup selection is latched: subsequent calls cannot restart normal dispatch or replace the primary failure reason. Pending/Idle requires a positively known original pollable obligation; otherwise its custody is Unknown and retained. Unknown follows cleanup without clearing physical-owner poison.

A Terminal load result selects cleanup, not physical closure. Identity rejection leaves the original operation owned and pollable; a validated Stale outcome instead follows its terminal retirement reason. A validated non-success can return only the custody named by its stage proof, with remaining credit discharge and physical retirement left to cleanup. No read, upload, install, or publication gate substitutes for a later consumer fence.

Algorithm 1 advances a load, not an entire request. Its record persists across Pending returns; fresh records start in the read-admission phase, with the frozen plan and exactly-once retirement authority. When every dependency is ready, the consumer revalidates the exact lease set and any typed read snapshot, acquires ready-cohort/resume claims, and submits execution. Pending execution retains its original payload. Only validated execution success enters the existing ending phase. Prepared-KV profiles preflight and revalidate the participants, dispatch install once, and poll that identity for exact ACKs. Other profiles advance their existing provider ending. Logical publication follows those ACKs; post-ending cleanup cannot resubmit ending. Displaced state retires on its own ACKs, and page retirement is confirmed before reuse. Figure 3 shows this scoped sequence.

Persistent cleanup and cancellation state. The terminal authority first accepts or rejects the requested caller outcome. Only an accepted terminal selection detaches that request's demand; shared, prefetch, and persistent owners survive. After detachment, Ω is a finite snapshot by reference of the affected obligations reachable through J: shared reservations, submitted tickets, ending/cleanup payloads, and retained claims. Removing an item from this snapshot does not remove it from its authoritative owner. Sharing and cancellation legality are rechecked at action time, not inferred from the earlier snapshot. A handoff removes local custody only after another authoritative owner acknowledges it; otherwise the original record stays reachable and charged. On revisits, the selected caller outcome and demand detachment are not applied again. A persistent cursor w starts the next snapshot after the last visited identity, including when that item is unresolved, and wraps at the end. Retired or newly added obligations are handled by the authoritative worklist, not by restarting every pass at its head or permanently skipping items before the cursor. ResumeOrSnapshot first resumes the persisted unvisited cohort; new obligations cannot overtake it, and visited unresolved items move behind unvisited ones. Repeated passes over a finite reachable set must examine every item. This is a scheduling requirement, not measured starvation freedom.

The persistent marker $\mu_o[s]$ records None, Pending, Submitted, or Uncertain for operation o and exact stage s. Its identity also includes the full key and source/destination generations. CancelOnce may call the provider only if cancellation is legal, no independent demand remains, no irreversible ending is reversed, and $\mu_o[s] \in \{\mathsf{None}, \mathsf{Pending}\}$. Provider acceptance records Submitted;

Algorithm 1 Progress one exact materialization operation

Inputs: P : Plan, v : Provider, $\mathcal{L}$: Ledger, δ : Demand, J : Journal
Outputs: h : {Ready, Pending, Terminal}, o : LoadRecord; residual custody stays reachable

State: persistent o : LoadRecord; s : Stage; e : Evidence; p : PublicationResult; α : AdmissionResult
Preconditions: frozen plan, exact destination, authoritative journal; residency/sharing checked
if cleanup is selected or o has retained Unknown **then**
  preserve J, o; continue only Algorithm 2
  return (Terminal, o)
end if
while $o.phase \neq$ Resident **do**
  if $o.phase =$ Installed **then**
    Publication: $p \leftarrow$ Preflight(o)
    if $p =$ Ready **then**
      publish exact binding; transfer resident grant
      resolve exact waiters; retire once; $o.phase \leftarrow$ Resident
    else if $p =$ Blocked **then**
      retain J, o; arm replacement-custody revisit
      return (Pending, o)
    else if $p =$ Rejected **then**
      RetainCleanup(J, o, p); **return**(Terminal, o)
    end if
  else
    $s \leftarrow o.stage$
    if stage s is not submitted **then**
      Admission: validate/canonicalize units, overflow, authority,
      source/placement/backend compatibility, and demand policy
      if validation fails or $\exists a : c_a > b_a^{\text{eff}}(\delta)$ **then**
        $\alpha \leftarrow$ Rejected with invalid-plan/intrinsic-capacity reason
        RetainCleanup(J, o, α); **return**(Terminal, o)
      else if $\exists a : U_a + c_a > b_a^{\text{eff}}(\delta)$ **then**
        $\alpha \leftarrow$ Wait; retain held claims; arm admission revisit
        return (Pending, o)
      else
        acquire owner-rechecked stage grants; $\alpha \leftarrow$ Admitted
      end if
      journal stage custody; SubmitOnce(v, o, s)
      if submission fails or is ambiguous **then**
        RetainCleanup(J, o, SubmissionError); **return**(Terminal, o)
      end if
    end if
    Completion: $e \leftarrow$ PollOriginal(v, o, s)
    if e is Pending/Idle and the original obligation is positively known pollable **then**
      retain o; arm completion/owner-progress revisit
      return (Pending, o)
    else if e is an identity/stage/generation/byte-contract mismatch **then**
      reject event without mutation; retain o and revisit
      return (Pending, o)
    else if e is Pending/Idle without that proof, a fault, or lacks physical proof **then**
      retain Unknown; RetainCleanup(J, o, e); **return**(Terminal, o)
    end if
    apply exact validated stage-custody return; retain residual claims
    if $e.outcome \in$ {Failed, Cancelled, Stale} **then**
      validated non-success: RetainCleanup($J, o, e.outcome$)
      return (Terminal, o)
    else if e is read success **then**
      release held ReadSlot/StorageReadBytes credits; retain pinned host charge
      $o.stage \leftarrow$ Upload; clear only the completed read ticket
    else if e is upload success **then**
      discharge upload/pinned claims and slab after exact fence; retain destination
      $o.stage \leftarrow$ Install; clear only the completed upload ticket
    else
      exact install success: $o.phase \leftarrow$ Installed
    end if
  end if
end while
return (Ready, o)

later passes poll/drain rather than resubmit that stage's cancellation. Known non-acceptance retains Pending with exact arguments for an owner-authorized retry. An ambiguous return latches Uncertain and retains Unknown custody: later passes cannot retry it as though the provider had rejected cancellation. Uncertain is a cancellation-marker state, not proof that physical-owner Unknown has ended. A new stage does not inherit an old stage's submitted marker. For a positively unsubmitted reservation, DiscardOnce follows the same marker rule if provider cancellation is needed, or the owner's exact safe preparation-discard protocol otherwise. Neither marker nor acceptance proves that submitted physical references have ended.

A pass visits at most k obligations from its finite snapshot, with at most one observation and one cleanup advancement per visited item and no redispatch of an accepted or ambiguous stage cancellation. The remaining snapshot is persisted before return; unvisited and unresolved obligations stay in persistent owner records. Cursor advancement keeps earlier Pending/Unknown items from hiding later items on every pass. Matching completion, accepted handoff, budget or lease release, and owner progress wake a later pass; without a notification, the known pollable path requires a scheduled owner revisit rather than a busy loop. These are bounded-work scheduling obligations, not a claim that provider calls are nonblocking, passes have bounded wall-clock latency, or every cleanup finishes. A stale identity event cannot mutate current custody, while a matching validated Stale outcome is terminal cleanup evidence only for its exact scope.

A secondary cleanup error preserves the primary outcome, current cleanup phase, exact identity and arguments, and residual custody; the same rule applies to discard and retirement errors. It cannot restart irreversible install/publication/ending. Revisit the original cleanup only if it remains positively known pollable; otherwise retain Unknown. Journaled partial returns/releases are not applied twice, even when a later cleanup step fails.

Physical-owner Unknown is never cleared by this algorithm. A late exact completion can close its named registry-stage obligation without resetting independent owner poison or freeing still-retained allocations. Such a registry closure alone cannot make the cleanup result Closed while that physical owner remains retained in J. An acknowledged custody handoff changes the responsible owner, not the retained charge or poison. Replacing a poisoned owner requires a separately established physical teardown/isolation boundary and fresh admission; logical shutdown, a new generation, or an unrelated later fence is insufficient. Cleanup success is not successful inference, and scoped bookkeeping closure does not flush live caches or prove system-wide quiescence.

5.1 Bounded next-expert lookahead

Algorithm 3 is the synchronous expert owner's schedule, not assumed general-registry use. Immutable plans select canonical gate/up/down artifacts; tickets bind exact operations, generations and extents. Matching demand validates source and binding before consuming a completion.

Symbol	Host-staging accounting meaning (bytes)
H	Combined configured host staging ceiling
$P_{\max}$	Largest selected expert's packed compressed weights and scales
M, S	Charged metadata allowance and requested worker-stack reservation

Table 2: Scheduling byte bounds; $P_{\max}$ is distinct from resource plan P.

Before spawning the worker, both modes require

$$2P_{\max} + M + S \leq H. \tag{7}$$

Algorithm 2 One bounded terminal-custody pass

Inputs: J : Journal, i : TerminalIntent, $k : \mathbb{N}_{>0}$
Outputs: y : Outcome, $h : \{\mathsf{Closed}, \mathsf{Pending}, \mathsf{UnknownRetained}\}$; unresolved custody persists

State: Ω : Worklist; $n : \mathbb{N}$; o : Obligation; s : Stage; e : Evidence; γ : CancelReply; β : CleanupReply; persistent w : Cursor, $\mu_o[s]$: CancelMarker
Preconditions: owner authority and exact identities; independent owners, irreversible intent, and live pins preserved
if $i = \mathsf{Cancel}$ and the terminal owner rejects cancellation **then**
  return $(\mathrm{CurrentOutcome}(J), \mathrm{CustodyStatus}(J))$
end if
Outcome: latch owner-accepted y; detach terminating demand once
preserve irreversible intent; suppress success only for accepted cancellation
$\Omega \leftarrow \mathrm{ResumeOrSnapshot}(J, w)$ after detachment; $n \leftarrow \min(k, |\Omega|)$
while $n > 0$ **do**
  $o \leftarrow \mathrm{Pop}(\Omega)$; $s \leftarrow o.stage$; $n \leftarrow n - 1$
  if another waiter/prefetch/persistent owner still holds o **then**
    preserve shared load or ACK-handoff; no cancel/discard or credit refund
  else if o has sticky physical-owner Unknown **then**
    retain payload, charge, and poison; block unsafe reuse
    if late exact scoped registry-stage completion is available **then**
      close only that named registry obligation; never clear owner poison
    end if
  else if o is positively unsubmitted **then**
    Discard: $\mathrm{DiscardOnce}(o, \mu_o[s])$
    if exact safe-discard proof is present **then**
      discharge only eligible held claims; consume retirement authority once
    end if
  else
    Drain: $e \leftarrow \mathrm{ObserveOriginalOnce}(o)$
    if e mismatches identity/stage/generation/byte contract **then**
      reject event without mutation; preserve original o
    else if e is validated and proves named stage-custody return **then**
      apply return; validated Failed/Cancelled/Stale selects terminal cleanup only
      Cleanup: $\beta \leftarrow \mathrm{AdvanceOriginalCleanupOnce}(o)$
      if $\beta = \mathsf{Error}$ **then**
        record secondary error; retain primary/phase/arguments; no ending redispatch
        keep custody Pending only if original cleanup is known pollable; else retain Unknown
      end if
      discharge only proved claims; account residency; confirm retirement before reuse
    else if e is Pending/Idle and the original obligation is positively known pollable **then**
      if cancellation is legal/unshared and $\mu_o[s] \in \{\mathsf{None}, \mathsf{Pending}\}$ **then**
        $\gamma \leftarrow \mathrm{CancelOnce}(o, \mu_o[s])$
        if $\gamma = \mathsf{Accepted}$ **then**
          $\mu_o[s] \leftarrow \mathsf{Submitted}$; still await physical evidence
        else if $\gamma = \mathsf{NotAccepted}$ **then**
          $\mu_o[s] \leftarrow \mathsf{Pending}$; retain exact retry arguments
        else
          $\mu_o[s] \leftarrow \mathsf{Uncertain}$; retain Unknown custody; no retry
        end if
      end if
    else
      Unknown: classify held claims in T_a once, without a second charge
      retain physical payload/journal and applicable quarantine; no reset/reuse
    end if
  end if
  $w \leftarrow \mathrm{AdvanceCursor}(w, o)$; defer unresolved behind unvisited
  arm exact completion/cleanup/owner revisit; never release on a wake alone
end while
persist unvisited Ω and cursor in J; retain unresolved owner records
if J has no unvisited or unresolved custody **then**
  finish only scoped terminal bookkeeping; $h \leftarrow \mathsf{Closed}$
else if J contains physical-owner Unknown **then**
  retain poison, payload, and charge; $h \leftarrow \mathsf{UnknownRetained}$
else
  retain J; $h \leftarrow \mathsf{Pending}$; arm a later pass
end if
return (y, h)

This covers packed/prepared coexistence, not two permanent expert copies. For admitted host scope, the owner additionally maintains

$$C_{\text{host}}(t) = B_{\text{packed}}(t) + B_{\text{artifact}}(t) + M(t) + S(t) \le H. \tag{8}$$

B_{packed} includes queued, in-flight and held-completion quota; B_{artifact} retains the artifact reservation while any exact weight/scale payload alias survives. Aliases do not add copies or earn a refund; independent copies need their own charge. $M(t)$ and $S(t)$ are the metadata and requested-stack reservations of live/retained host scope. Unknown is already retained in these charges, not added twice. One worker has queue depth one and two charged slots. The measured owner-transition high-water tests this logical contract, not continuous RSS: OS stack rounding, guards, TLS, weak-reference backing and allocator headroom remain distinct.

Algorithm 3 One bounded next-expert read window in a synchronous forward

Inputs: ordered selected buckets $\mathcal{B}$, source-bound plans, mode $m \in \{\mathsf{Blocking}, \mathsf{Lookahead}\}$, limits H, M, S
Outputs: unchanged bucket order and device proofs; residual custody retained

Preflight: validate canonical three-artifact plans and Equation 7; start the same bounded worker in both modes
State: at most one next ticket τ; exact deferred error; packed/prepared payload charges and allocation-alias witnesses
for current bucket b in $\mathcal{B}$ **do**
  owner admits/pins current device residency and scratch; validate source
  if current expert is not physically resident **then**
    drain any unrelated speculation using its original identity
    honor only b's deferred error before any new dispatch
    if no exact carried ticket for b **then**
      $\tau \leftarrow$ SubmitExact(b) under shared read-slot/byte limits
    end if
    Demand: await original τ (immediately after blocking submission)
    if worker custody is Unknown **then**
      retain original identity/charge/poison; do not take or fabricate completion
      return with residual custody reachable
    end if
    take completion once; validate ticket, plan, image/layer/expert and source
    if non-success or mismatched evidence **then**
      preserve primary error and residual custody; select terminal cleanup
      return without fabricated completion or publication
    end if
    reserve prepared-payload quota while packed bytes remain charged
    construct exact gate/up/down artifacts without reread; track all payload aliases
    upload current expert using existing owner admission and CUDA proofs
  end if
  Read window: drain speculation made redundant by a current resident hit
  release current host charge only after alias retirement
  inspect next selected bucket with read-only residency hint; no next pin/LRU mutation
  if $m = \mathsf{Lookahead}$ and next is nonresident and host custody permits **then**
    try to submit that one next read before current compute submission
    retain speculative submission/read errors for that exact later demand
  end if
  Compute: owner submits current compute from resident device handles
  consume output under current lease; preserve existing fences and retirement
end for
Scope end/error: cancel unused reads, drain and join original worker custody
retain Unknown/poison and alias-held quota; host join never heals GPU Unknown

Speculative errors wait for the matching demand, not the current bucket. Unused reads are canceled/drained and counted; lookahead may read more bytes than blocking. Hints do not pin a later hit, and alias/budget checks never refund live custody. OS reads are not interruptible by cancellation; host join proves no GPU quiescence. Whole forwards remain synchronous.

Weak witnesses track exact weight/scale allocations without adding payload copies. Strong aliases retain quota. After their last retirement, weak backing/control allocations and allocator-

Contract	Implementation anchor (under `crates/`)	Evaluation obligation
Source-bound P	`ferrule-model/src/checkpoint/`: catalog, source, read plan	Freeze source/representation; descriptor hash is not payload authentication
Exact K, e	`ferrule-common/src/io_protocol.rs`: keys, completion, slabs, leases	Reject stale operation/stage/generation/byte evidence
Covered claims	`ferrule-runtime/src/scheduling/resources.rs`: `PhysicalResourceBroker`	Map ledgers to real allocations and final retained custody
Load publication	`ferrule-runtime/src/io/registry.rs`: reserve, completion, preflight, resident transfer	Verify provider transfer installs this registry on the evaluated path
Physical completion	`ferrule-model/src/moe/cuda_materialization.rs`: routed-expert provider	Trace pinned read, upload, install, replacement, and Unknown ownership
State publication	`ferrule-runtime/src/parallel/pipeline/owner.rs`; decoder KV and page manager	Journal before preparation; exact ACKs and retirement before reuse

Table 3: Contract-to-code mapping. Common interfaces do not imply universal registry integration or measured storage traffic.

retained memory may still exist: quota release is not RSS reduction. Passive observations preserve Unknown and poison without refund, recovery or fabricated Closed status.

6 Implementation Mapping and Integration Boundaries

The implementation separates common protocol identities, model semantics, backend physical work, and runtime coordination. CLI/HTTP interfaces deliver caller-visible results; they are not the model driver's publication authority. Table 3 maps the preceding contracts to concrete modules and names the coverage that must be checked before interpreting a run. These are local implementation anchors, not a profile-wide correctness result.

Source binding and startup preparation. Checkpoint catalogs and read plans hold semantic tensor descriptors, ordered file extents, and filesystem identity snapshots, not payload bytes. In this adapter, `ContentHash` hashes dtype/shape/byte descriptors; source identity additionally binds paths, offsets, and file metadata. Neither is a read-and-hash authentication of every payload byte. The physical provider resolves the plan against the residency authority and obtains the actual destination generation. A startup CUDA image claim is a different authority from a per-request load ticket: actual image grants, compiled context, image fence, KV/initial-state creation, and the final claim fence are checked before claim consumption. Eligible raw execution bindings may then be replaced by runtime descriptors; surviving shared source aliases are not thereby freed.

6.1 Implemented bounded expert-read path

The opt-in checkpoint pipeline is a real positioned-read worker, not a task that reports completion without performing I/O. Its immutable plan preserves ordered extents and original filesystem snapshots. Unix physical-source identity version 2 includes device/inode, length, mtime, and ctime coordinates; ctime strengthens change detection but is not payload authentication. The worker

verifies the source before/after reading, and the consumer rechecks it before constructing the canonical three compressed artifacts. No CUDA handle, cache lease, or runtime continuation enters that worker.

The integration has three local anchors: checkpoint `read_pipeline.rs` implements tickets, reads, completion custody and join; transformer `expert_reads.rs` selects gate/up/down plans and accounts payload aliases; transformer `standard/cuda.rs` separates the read window from compute submission. CLI options flow through a prepared, source-bound admission plan and lower into the existing numeric BF16 factory/materializer. The selected path remains PP1/TP1/EP1, resident KV and lazy experts; enabling host staging does not enable unsupported topology or a different forward API.

Both *blocking* and *lookahead* explicitly enable the same pipeline; the default disabled CLI path has no worker-read evidence and is not the control. Blocking submits required demand and waits immediately. Lookahead carries one next ticket across the current owner compute call, then waits/takes it only when that exact expert is demanded. After current upload, temporary host artifacts leave the compute path: compute borrows resident GPU handles and immutable binding metadata. A next-residency probe is read-only and occurs after current admission/upload; it neither pins the next expert nor changes completed-only LRU order, routes, or kernel mathematics. Submission and source errors from speculation are deferred to the matching demand; abandoned work is canceled/drained and counted as wasted read activity.

The study configuration fixes $H = 64$ MiB, $M = 8$ MiB, and $S = 2$ MiB, with one worker, queue depth one, and two read-operation slots in both arms. These host charges are separate from the 64-entry, 320 MiB device cache charge and its 8 MiB scratch component. Metadata and requested stack are *inside* the host ceiling; actual OS/TLS/allocator overhead requires additional unmeasured headroom. Packed completion and prepared payload quota stay live through their actual aliases, including escaped weight/scale `Arc`s. Weak liveness witnesses do not add a payload copy, but retained allocation backing and allocator RSS are distinct from quota (Section 5.1).

Passive snapshots expose requested/effective mode, worker/read evidence, operation, image, layer and expert identity, demand waits, completed and wasted bytes, charged payload/metadata/stack, join status, and retained Unknown. Reading a snapshot does not repair poison, join a worker, or refund custody. A scope-ended status is not a final physical-resource closure certificate. Cancellation/drain/join proves only named host obligations, never GPU quiescence or healing of a poisoned GPU owner. Whole forwards still block on their existing consumer proofs. These mechanisms and CPU test fixtures exist. The frozen `033519c` factory/profile passes the nine GPU gates in Table 6; the enabled CLI pilot stopped pre-Ready. Its source-plan query fix is committed at `86c58a0`, but new CLI/HTTP qualification is pending (Appendix A). Passive resource snapshots leave allocator state unobserved because its getter can poll retirement fences; separately reported allocator high-water is not cost-free passive telemetry or physical release evidence.

6.2 Other lifecycle integration boundaries

General materialization path. The registry is installed only when a runner transfers a physical materialization provider to the resident engine. It merges compatible exact keys, validates completion before mutation, and separates return of stage custody from credit release and resident-grant transfer. The concrete routed-expert provider uses Linux pinned reads/io_uring and CUDA upload/installation. Registered slabs retain aligned ranges and exact generations; upload tickets retain the original event and frame, even after notification errors. Replacement passes the previous binding and old-consumer compute proof to the install adapter. A generic lease completion contract is owner metadata, not an independently recorded global GPU fence. Some preparation helpers use compute

and upload stream events; synchronization is not excluded.

The broker canonicalizes claims and validates their limits before charging. Required execution can use a reserve unavailable to warmup/speculation, but that does not guarantee progress if a provider stalls. Load admission is stage-local, not an atomic reservation of all future stages. A later-stage budget wait leaves earlier slabs/destinations charged. Returned provider custody clears the submitted obligation, not occupancy; grants have no implicit destructor refund. The five classes of Equation 6 are therefore semantic refinements of owner state, not five measured fields.

Owner-local state protocols. Ordinary TP1 pipeline segments use startup preparation, persistent stage owners, and a prepared-KV journal; they do not automatically inherit registry single-flight or I/O overlap. The provisional journal precedes backend preparation because preparation can submit work before returning a handle. Readiness validates existing evidence, not new submission. Lost install ACKs poll the same identity without redispatch. Exact participant ACKs precede logical publication, then displaced-state retirement and zero-reference-page confirmation precede reuse. Finish consumes transaction-handle custody, not live weights, KV, or independent pins. HostKV decode/idle reclaim use typed read and placement/page-retirement protocols rather than the expert registry. Committed-KV source snapshots include owner, slot, incarnation, generation, layout, frontier, range, and pages; acquiring a read lease revalidates them, and physical quiescence precedes pin release. Append witnesses are also revalidated against the live reservation. These are not database snapshots.

Unknown and collective scope. Provider Unknown retains native tickets, allocations, claims, and cleanup intent, and may poison slot layers, devices, or executors. A matching registry completion can close a named obligation without resetting independent poison. There is no general in-place reset/replay or device-reset recovery claim. Prepared-KV coordination is a narrow state lifecycle, not arbitrary distributed atomic publication. Current collective exchanges have separate host-payload bounds; Mixed device participation defaults to disabled/host. A typed device permit is not evidence of a default GPU reduction or cross-rank fence. For the vLLM supplement, Socket-plugin initialization is not proof that local GPU payload used Socket: the scoped logs record `NCCL_P2P_DISABLE=1` and local channels `via SHM/direct/direct`. No NCCL capability or transport conclusion is generalized beyond the recorded profile.

What must be traced. A full path audit names each progress owner, queue/limit, dependencies, submission behavior, completion evidence, blocking calls, locks, and synchronization from request admission through storage, expansion, transfer, execution, publication, retirement, cancellation, and shutdown. Include model-specific and host-KV bypasses. A worker wrapping a blocking call is not proof of event-driven progress or hardware overlap. Current BF16/numeric and Mixed TP profiles have distinct shard/cache/expansion policies; actual storage integration and residency require profile-specific evidence. Shared/unified physical pools and other modalities remain design scope, not evaluated support. Rust ownership assists local reasoning but does not discharge asynchronous physical references owned by providers.

7 Evaluation: Setup and Four Research Questions

We report an exploratory historical TP1 logical-demand observation, current same-cut GPU factory qualification, and separate TP2 numerical and resident-serving records. The nine current gates are a numerical/provider qualification, not paired performance, CLI/HTTP, or physical release. We

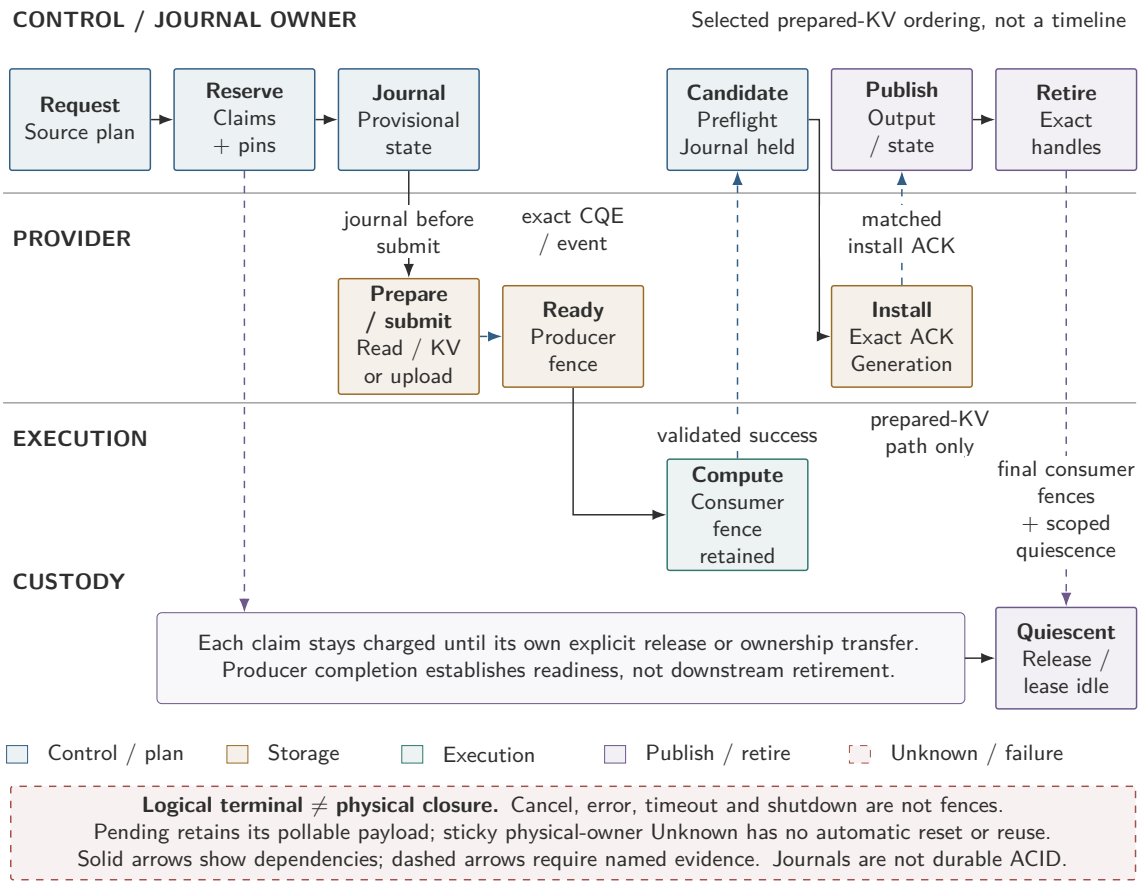


Figure 3: Conceptual successful lifecycle for the selected prepared-KV owner, not a measured overlap timeline. Named lanes and swatches retain blue control/validation, ochre materialization, teal execution, violet publication/custody, and brick Unknown/failure roles. The provisional journal precedes backend preparation even if submission fails without a handle. Candidate readiness validates existing evidence; exact generation-matched install ACKs precede logical publication on this path. Producer readiness is not consumer retirement. Final relevant fences and scoped quiescence precede custody return, with explicit release/transfer still required. Cancellation cannot reverse irreversible ending. Neither durable ACID, universal registry integration, nor general cross-rank atomicity follows.

Configured field	Value / scope
Device admission / workspace	7,168 / 2,048 MiB; not a measured total working-set cap
Expert cache / scratch	64 entries; $b_{\text{cache}} = 335{,}544{,}320$ bytes (320 MiB, shared cache-accounting ceiling for resident, pending, and scratch); separate $b_{\text{scratch}} = 8$ MiB
KV / maximum tensor	16 / 1,024 MiB; distinct covered limits
Host expert retention	Zero; not zero OS page-cache residency
Request queue / inflight	2 / 2; actual workload concurrency 1
Physical host limit / caches	Host cap unknown; page cache uncontrolled, not storage-cold
Logging / generation	`ferrule_performance=DEBUG`; temperature 0, seed 1

Table 4: TP1 middle configuration. Values from the frozen case manifest; heterogeneous or overlapping limits cannot be added as physical memory.

do not combine profiles to create a quality, capacity, or performance result. E0 (RQ1), E3 (RQ2), the safe blocking/async part of E1/E4 (RQ3), and integrated E5 (RQ4) retain their historical experiment identifiers. Broad serving sweeps (E2), further ablations (E4), and protocol overhead (E6) remain secondary evidence plans.

7.1 Setup, controls, and evidence units

Historical TP1 middle `-O2` uses Qwen3.5-35B-A3B-FP8. Its single thread-owned CUDA runner has PP1, TP1 and EP1, the resident engine, lazy experts and resident KV, context 128, batch-token limit 8, active-sequence limit 2, and prefill chunks of 8. Its Mixed arithmetic uses BF16 activation/KV/convolution with F32 accumulation/SSM; the FP8 checkpoint is not a complete arithmetic contract. The sealed binary is `36e6aaa5...70a58c`. Table 4 reports configured limits, not physical partitions.

The historical tables and all four data figures below retain this frozen `-O2` identity. They use DEBUG invocation counters from the disabled read-pipeline path, not enabled-pipeline statistics. The current scheduling configuration adds the separate 64 MiB host ceiling in Section 6.1; its implementation is not new execution data and does not update these plots.

Here b_{cache} is a configured cache-accounting ceiling, not a physical-memory bound. The plot field `cache.peak_bytes` (raw receipt key `expert_owner_peak_bytes`) reports the lifetime high-water of cache-charged bytes: resident plus pending-upload reservations plus scratch. That charged peak is governed by the 320 MiB configured cache ceiling. It is not a complete physical-owner memory peak or an independent measurement of allocator retirement; a total physical-owner ceiling remains unknown. The configured cache flag includes scratch according to the sealed startup plan, while b_{scratch} is a separate component limit. A shared cache byte ceiling constrains resident, pending, and scratch charge; it is not a separate resident-only budget.

Attempt and request accounting. The initial CPU-prepared entry was unexecuted after its final validation was canceled. Authorized attempt `-O1` then failed before server launch with a harness `AttributeError` for missing `subprocess.DEVNULL`: planned 4, dispatched 0, not-run 4. No model/OOM, read, request-latency, or quality result was obtained. Authenticated drain/post-idle checks passed, while physical I/O and the resource contract remained unknown and GPU quiescence unverified. This attempt is retained, not pooled with `-O2`.

Attempt `-O2` dispatched four serialized requests: Hello2 and the QA, arithmetic, and Chinese canaries. All four reached successful HTTP termination; failed and not-run counts were zero. They returned 26 usage output tokens in 30 decoder invocations. Hello2 has no reference and

Observable	Reported value	Scope
Selected GPU0 sampled maximum	4,125 MiB	32 GPU samples
Server-tree PSS sampled maximum	7,610.864 MiB	32 complete host samples
Server-tree RSS sampled maximum	7,613.770 MiB	Same sample population
Expert resident entries / bytes	64 / 201,351,168	Last invocation gauge
Scratch bytes	8,388,608	Last invocation gauge
Cache charged-byte high-water	219,754,496	Lifetime accounting gauge
Pending-upload / quarantine bytes	0 / 0	Invocation gauges only

Table 5: TP1 `-O2` sampled memory and cache-accounting gauges. Sampled memory maxima are independent, not additive or certified physical peaks. Cache charged-byte high-water covers resident, pending, and scratch; its ceiling is 320 MiB, not a physical-owner bound or final ledger proof. GPU values are selected-device aggregate telemetry, not process attribution.

passed HTTP/usage only. The other three passed their frozen visible-text/eight-output-token-usage checks (3/3); QA is one canary, not an extended QA workload. Generated IDs, logits, routes, and full current-model quality were not observed. This is one independent DEBUG-instrumented run ($n = 1$); observer/logging overhead was not measured. Its request timings are exploration, not a fast baseline or steady-state comparison.

TTFT is send-to-first-nonempty SSE content and E2E is send-to-EOF. SSE events and usage are not exact token timestamps; no exact-token TPOT is reported. Named request/time/log-byte anchors attribute serialized invocation deltas; a transaction ID is not a request ID. Memory samples, logical counters, physical traffic, and owner-ledger occupancy are distinct measurement units. Freeze snapshot/dirty manifest, binary, model/tokenizer, workload, state precision, backend, storage, topology, caches, and interval boundaries for any future comparison. Offered/accepted/completed/failed/rejected/canceled/ timeout counts must remain in the denominator record.

7.2 RQ1: Factory quality qualified; capacity frontier open

Equations 1 and 2 define deduplicated physical capacity and the full-residency test for a declared complete weight representation; Equation 3 asks whether the actual streaming working set fits each domain. The configured limits in Table 4 answer neither question alone. Table 5 reports the available memory proxies without promoting them to continuously measured M_d^{peak}. The storage-demand cache figure in RQ2 is a counter/gauge reconciliation, not a physical-working-set certification.

Host capacity, attributable page-cache residency, continuous peaks, and process-owned GPU occupancy remain unknown. The OS cache was uncontrolled; zero host expert retention does not imply zero file residency. Complete compact weights plus scales/metadata have not been ruled out as a feasible resident alternative. Nor do four successful HTTP responses constitute the independent complete-output quality gate. Therefore RQ1 is open: quality-qualified configured-budget feasibility is not established by this TP1 record. Physical-memory insufficiency and storage necessity are stronger, optional claims and are not established either.

Current enabled-pipeline factory qualification. The receipt records **9/9 exact-once GPU gates PASS** on `033519c`, with 1,484 frozen source members. Constructor and composition use the same cut, not pooled historical passes. Both schedules retain Table 4's composition limits and Section 6.1's 64/8/2 MiB host/metadata/stack contract. Table 6 closes RQ1's selected factory numerical subgate, not the capacity frontier or full physical working set.

Phase	Gates	Metrics/gate	Vocab values/gate	Additional evidence
Full40	2	804	1,986,560	2,560 ordered route IDs per schedule; diagnostic state checks
Retained fresh	2	5,427	9,187,840	64 actual fingerprints per schedule; diagnostic routes/states
Ordinary	2	27	9,187,840	All 64 fingerprints match own-schedule fresh primary; routes/states unobserved
Cold owner	2	1,809	3,228,160	All 28 fingerprints match own current fresh primary
Constructor	1	—	128 total	Four boundary/schedule cells; 32 finite vocabulary values and six verified reads each

Table 6: All nine current factory gates pass once. Each two-gate row has one gate per schedule, not performance repetitions. Fresh/ordinary/cold vocab counts are common-history only; constructor is tiny provider evidence. Profile-specific metric counts are not summed.

The tolerance remains $0.00025 + 0.012|\text{reference}|$, with exact ordered route/greedy-ID checks where observed, not universal bit-exactness. Ordinary runs have no diagnostic route/state observer: those fields remain null, not zero or inferred observations. Their actual logit-word/shape/greedy fingerprints match the pinned diagnostic primary of the *same schedule*. All 28 cold fingerprints match that schedule's current primary, with the same config/oracle, unlike the old source mismatch in Appendix A. The main metric counts do not mix ordinary output checks with diagnostic state/route metrics.

Separate prediction-stream denominators. For each schedule's fresh and ordinary runs, common history and free feedback *each* have 27 calls/9,187,840 vocabulary values; reset has 9/3,228,160 and fork replay 1/248,320. For each cold run, common history, free feedback and reset *each* have 9/3,228,160, and fork has 1/248,320. The 64 or 28 fingerprints cover these named streams, not additional logits. These denominators are kept separate across stream, observer mode and schedule; common-history values are not counted repeatedly as one larger independent-quality result. Fork scope is same-history committed-frontier replay, not divergent KV copy-on-write or broad task-quality evidence.

Real worker and bounded-accounting observations. Each schedule records 2,089 verified reads in full40, 25,116 in each fresh/ordinary run, and 11,354 in cold; constructor records six per cell. These are last-sampled owner-lifetime verified reads, not additive invocation deltas, sampled ticket counts, physical NAS bytes or throughput observations. The reported logical host high-water is 16,777,984 bytes, with $P_{\max} = 3,146,112$ bytes:

$$2(3{,}146{,}112) + 8{,}388{,}608 + 2{,}097{,}152 = 16{,}777{,}984 \leq 67{,}108{,}864. \tag{9}$$

This is alias-aware logical charge under Equations 7 and 8, not RSS. Separately, allocator reserved-byte high-water is 4,106,223,616 bytes, below the 7 GiB policy. Its getter can poll fences/release blocks; immutable snapshots leave allocator unobserved. These qualified point observations certify neither continuous total physical peaks, cost-free observation, physical release, a capacity frontier nor a statistical guarantee.

Actual PID/birth mappings record CUDA 12.4.99 with SM86 compiled code and NVTX compiled but profiling OFF (`FERRULE_PROFILE_NVTX=0`). All nine gates ran once with no probes, warmups,

TP2 test record	Calls	Comparator metrics
greedy4-02	10	10,890
fork2-03	6	6,534
canary2-01	5	5,445
recreate1-01	4	4,356
cold-owner1-01	4	4,356
cold-owner1-02	4	4,356
Total	33	35,937

Table 7: Separate TP2 Mixed test-binary numerical/lifecycle evidence from `experiments/tp2_resident/model-runtime-quality.json`. Exactness applies to the frozen comparator metrics, not TP1, vLLM arithmetic, all workloads, or the current working tree.

retries or captures. Natural exit/authenticated `ECHILD`, absent owners, all-eight stable-idle v2, final three zero-memory/zero-utilization/no-app samples and released/reprobed local UUID locks close process supervision, not CUDA physical release. Distinct CLI/constructor/composition executables with common source/native/ runtime provenance are not mathematical CLI/HTTP or machine-identity equivalence. No paired timing, NFS or GPU-overlap result is inferred.

The portable receipt archive, committed at `17b16fa`, is relative to `papers/arxiv`:

```
experiments/storage_demand/qualification_v1/
```

The archive has 249 files (5.38 MiB), including 240 manifest-bound records (5,343,269 stored bytes). Its offline checker and 37 tests pass without GPU, model, network or original local paths. Executables, weights and full source payloads remain external hash-pinned inputs, not archived or rechecked by that checker. Curation is not another run. Appendix A preserves prior failures.

Separate TP2 numerical evidence. The `native-whole-capture-full40-02` comparator reports 2,178 exact fields for each of Hello and Capital, covering prefill and first decode; max absolute error, RMSE, and field failures are zero, with matching greedy IDs. This is a separate full-boundary receipt from the six-row selection below. These are bounded captured/comparator fields, including oracle-only diagnostics that are not all physically observed device arrays, not a blanket bitwise claim over the model. The full-boundary audit freezes this scope. A separate multistep/lifecycle audit records the six cases in Table 7: 33 calls and 35,937 exact comparator metrics. This table deliberately selects the frozen `experiments/tp2_resident/model-runtime-quality.json` receipt and its `cuda-qwen35-mixed-bf16-tp2-local-shards-serial` profile, with test binary `a4977b85...e1cc`. A different frozen `production-opts-01` acceptance summary, referenced by the optimized release handoff, reports 37 calls and 40,293 exact metrics across its full40 and six-recipe selection. Its composition/model test binaries and source snapshot differ. That summary is not substituted for the six rows here: 33/35,937 is the selected denominator for this table only, not a total of all TP2 validation. Neither count covers the current nine-gate qualification above. The frozen greedy4-02 closure records supervisor/drain/release PASS; the older greedy4-01 supervisor failure remains a failure. These authenticated test binaries are not public serving binaries. No serving HTTP logits were observed, and this writing pass did not rerun payload comparisons.

The local frozen F32 process-EP8 full-model pass also remains scoped to its profile/snapshot, not a thread result or a capacity result. The independent single-GPU W8A16 strict oracle gate failed on logits, ordered routes, and internal states; argmax agreement does not reverse that failure. These records are kept separate from TP1's three limited canaries and from the resident timing supplement below. Neither quality nor capacity transfers between profiles merely because a checkpoint is shared.

Request	Logical read GiB	Uploads experts	Evictions experts	TTFT s	E2E s
Hello2	2.813	960	896	4.265	9.186
QA canary	10.777	3,678	3,678	8.332	18.309
Arithmetic canary	10.396	3,548	3,548	5.308	14.110
Chinese canary	9.766	3,333	3,333	2.605	9.932

Table 8: Four serialized TP1 `-O2` requests. Logical GiB ($1\,\text{GiB} = 2^{30}$ bytes) and HTTP timings are rounded to three decimals; expert counts are exact invocation deltas. Not physical storage, bus traffic, kernel timing, or comparative performance.

Experiment needed to close RQ1 (E0). Freeze representation/placement and complete-output tolerances; report covered budgets and measured working-set components without conflating them. If claiming physical insufficiency, additionally report each B_d, $W_{\text{full}}(r)$ and the compact-residency alternative. Measure physical working-set components and attribution across startup, requests, and cleanup, stating sampling coverage. Record full-model prefill/multistep decode for this text instance, repeated materialization/retirement, retained claims, swap, and declared cache conditions. Align request reads at the observable levels in RQ2; backing-disk measurements are required only for a backing-disk claim, not for the scoped bounded-residency result. Sweep configured budgets separately from physical capacity and retain failed/stalled points. Use a correct bounded disk-backed reference and a host-only memory-offload reference where feasible under the same quality/physical contract. An infeasible reference is infeasible, not zero time or infinite speedup.

7.3 RQ2: Request I/O – separate observability levels

For request q, the additive invocation observables are

$$\begin{aligned} L_q &= \sum_{v \in \mathcal{V}_q} \Delta_v \text{reader_bytes}, \\ H_q &= \sum_{v \in \mathcal{V}_q} \Delta_v \text{facade_H2D_bytes}, \qquad \Delta_q = H_q - L_q. \end{aligned} \tag{10}$$

L_q is logical checkpoint-reader demand and H_q is facade transfer activity. Their difference Δ_q compares two recorded counter layers; it is neither backing-storage traffic nor a physical bus measurement. In particular, L_q, H_q, Δ_q do not imply physical P_q^{net} or P_q^{blk}. Table 8 and Figure 4 apply these sums; gauge values in Table 5 are not summed.

The 30 invocations total 69,114 checkpoint-reader calls and 36,240,064,128 logical bytes, 88,763 facade H2D copies and 36,240,248,688 facade bytes, zero recorded expert hits, 11,519 expert uploads, 11,455 evictions, and 34,557 artifact uploads. Thus the four-request aggregate is $\Delta = 184{,}560$ bytes, or 180.234375 KiB; this small difference is an application-counter distinction, not evidence of disk traffic. Upload counts are not observed cache misses or measured I/O credits. These counters show repeated materialization activity, not the general registry's path integration or physical NAS/network/block-device reads. Startup reads are outside the sums.

For the expert-entry count, let I_j and E_j be the recorded install/upload and eviction deltas for completed invocation j, and let N_j be the entry gauge at that invocation boundary. The observed-boundary relation is

$$N_j = N_{j-1} + I_j - E_j, \qquad j = 2, \ldots, 30. \tag{11}$$

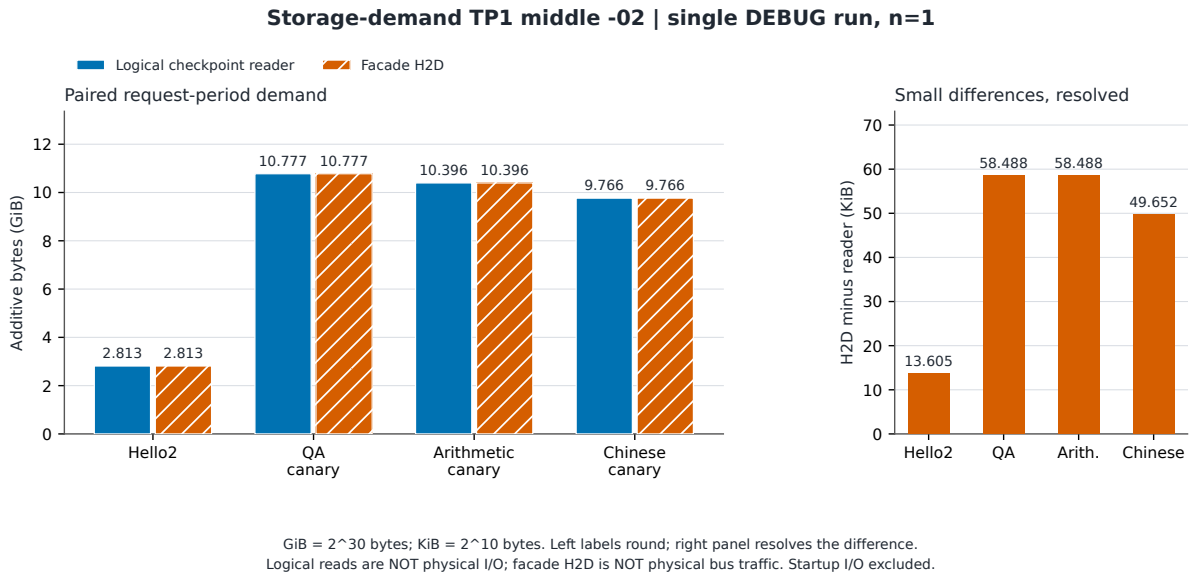


Figure 4: Paired logical checkpoint-reader and facade H2D byte deltas for four serialized requests in storage-demand TP1 middle -02. GiB means 2^{30} bytes; the right panel resolves the small H2D-minus-reader differences in KiB (2^{10} bytes). These are application counters, not physical storage I/O or physical bus traffic; startup I/O is outside invocation accounting (Equation 10).

For the first boundary, $N_0^{\text{implied}} = N_1 - I_1 + E_1 = 0$; this is an implied initial condition, not a measured pre-request gauge. Across the run, 11,519 uploads minus 11,455 evictions equal the last observed 64 resident entries; all subsequent observed boundaries have zero balance residual. This is entry-count reconciliation, not byte conservation, continuous occupancy, or a final owner-ledger proof.

Figures 5 and 6 show the aggregate and per-invocation forms of Equation 11. The timeline contains 3/9/9/9 invocation boundaries assigned by serialized request/log-byte anchors, not 30 output-token or elapsed-time samples.

For the cache-gauge subset, define

$$\begin{aligned} C_j^{\text{cache}} &= X_j + Q_j + S_j, \\ C_{\text{peak}}^{\text{cache}} &= \max_{u \in \mathcal{U}_{\text{owner}}} C_u^{\text{cache}} \le b_{\text{cache}}. \end{aligned} \tag{12}$$

where X_j is resident expert bytes, Q_j pending-upload bytes, and S_j scratch bytes at invocation boundary j. Unknown-quarantine bytes Z_j are an already charged retained subset of C_j^{cache}, not an additional term. This subset is distinct from general-ledger retained T_a. These are cache-owner counter semantics, not evidence that all physical allocations or grants were observed. This is an observed cache subset, not $M_d(t)$ and not the complete $F_a + R_a + I_a + L_a + T_a$ ledger. At the last invocation, the recorded values are $X_j = 201{,}351{,}168$ bytes, $Q_j = 0$, and $S_j = 8{,}388{,}608$ bytes; the reported Z_j is zero. Equation 12 adds compatible gauges at that one boundary, not maxima or observations across time. Figure 5 keeps the plotted gauges separate. $\mathcal{U}_{\text{owner}}$ denotes cache charged-state update points over the owner's lifetime, not just the 30 invocation samples. The recorded `cache.peak_bytes` is $C_{\text{peak}}^{\text{cache}} = 219{,}754{,}496$ bytes, an accounting high-water governed by `ExpertCacheLimits.max_bytes`, namely $b_{\text{cache}} = 335{,}544{,}320$ bytes. This is the cache admission/charging contract, not independent verification of every allocation or the general credit

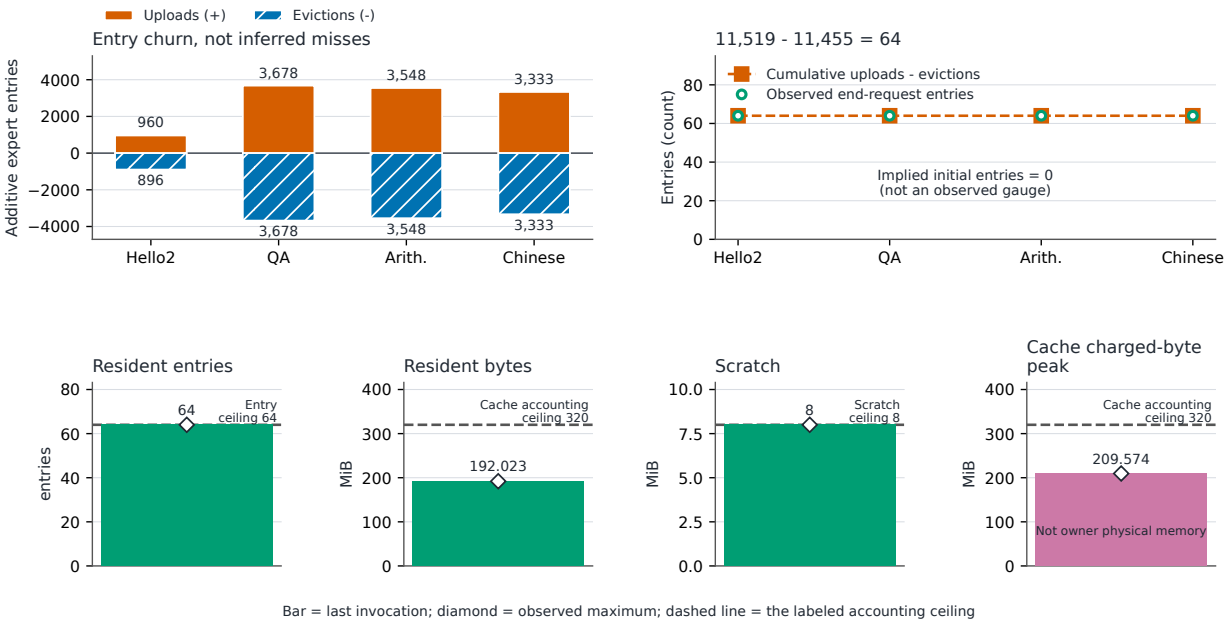


Figure 5: TP1 middle -02 entry reconciliation and cache-accounting gauges. Uploads minus evictions ($11{,}519 - 11{,}455 = 64$) match the last observed entries; initial zero is implied, not measured, and subsequent boundaries reconcile. Recorded `cache.peak_bytes`, 209.574 MiB, is the resident-plus-pending-plus-scratch high-water, bounded by `ExpertCacheLimits.max_bytes` = 320 MiB, not a physical-owner bound. Total physical-owner peak/ceiling are unobserved; no contradiction is implied. Invocation gauges and lifetime high-waters are not accumulated or a final ledger proof (Equations 11 and 12).

partition. Neither the peak nor b_{scratch} certifies total physical-owner memory, continuous $M_d(t)$, or shutdown custody closure. Unknown quarantine remains an already charged subset and does not add another copy of the charge.

Startup was 18.497 s. Four separate HTTP windows sum to 51.538 s, excluding startup and inter-request owner overhead; supervisor wall time was 75.874 s. The origins in Figure 7 are explicitly

$$\begin{aligned} T_{\text{start}} &= t_{\text{ready}} - t_{\text{launch}}, \\ T_q^{\text{first}} &= t_q^{\text{first content}} - t_q^{\text{send}}, \qquad T_q^{\text{end}} = t_q^{\text{EOF}} - t_q^{\text{send}}. \end{aligned} \tag{13}$$

These are distinct denominators. Startup has a launch origin; each request bar has its own HTTP send origin, not one cumulative timeline. It therefore does not imply an end-to-end throughput rate or an I/O bottleneck share. P_q^{net} and P_q^{blk} remain unknown. Logical reader/syscall bytes can be served by cache; facade transfers are not physical bus measurements.

Observability levels and a scoped closure. RQ2 does not require assuming that each logical byte reached the NAS disks. Table 9 fixes three non-interchangeable levels.

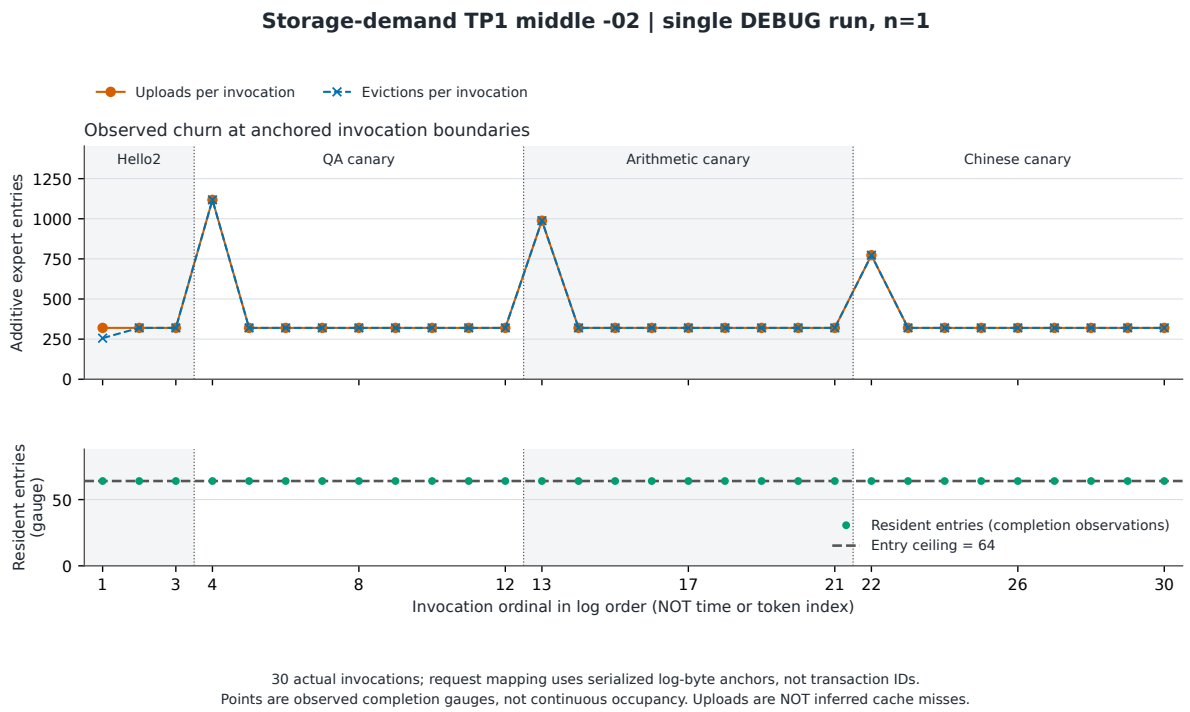


Figure 6: Per-invocation expert upload/eviction deltas and resident-entry observations for the 30 completed invocations in storage-demand TP1 middle -02. Serialized request/log-byte anchors assign the 3/9/9/9 invocations to requests. The x-axis is log-order invocation ordinal, not elapsed time, output-token index, or transaction-as-request ID. Entry occupancy is observed only at invocation completion, not continuously; uploads are not inferred cache misses (Equation 11).

Level	Observable / current status	Attribution limit
Logical reader	Historical HTTP `-02` counters and current factory verified reads actual	Distinct serving/test scopes; cache-served demand, not physical I/O
NFS client	Shared shutdown-only delta in stopped pilot; no Ready/request window	Shared client/mount scope, foreign traffic and reset/missing counters; not server-owned network bytes
NAS backing	Physical server network/block bytes and latency unknown	Client counters cannot establish server-cache misses or disk reads

Table 9: RQ2 observation hierarchy. Missing/reset fields stay unknown, not zero. No level certifies the unobserved level below it.

A scoped paper can close RQ2 by reporting same-run aligned logical demand, read timing, and available client observations with their explicit unknowns; it need not claim physical-storage necessity. Separate startup from request windows and runtime, OS/shared file, and storage-server caches. Fresh owners do not imply storage-cold runs. Disclose swap, uncontrolled caches, foreign traffic and clock/attribution limits. A physical disk-necessity or backing bandwidth claim would require additional attributed NAS measurements, not an inference from logical misses or shared-client deltas.

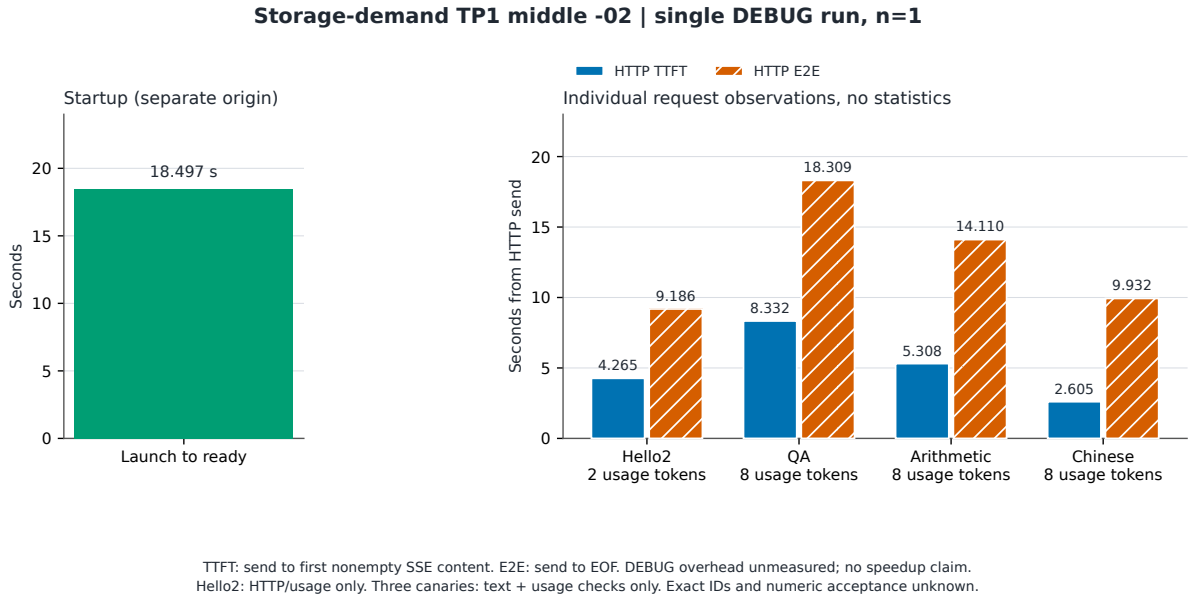


Figure 7: Launch-to-ready startup and per-request HTTP TTFT/E2E in the single successful DEBUG-instrumented storage-demand TP1 middle -02 run ($n = 1$). TTFT is send to first nonempty SSE content; E2E is send to EOF. Startup has a different origin and is not part of the request bars. Instrumentation overhead is unmeasured; there are no statistics, confidence intervals, kernel timings, or comparative speedups (Equation 13).

7.4 RQ3: Mechanism/numerics qualified; performance pending

The safe blocking control and bounded next-expert lookahead are implemented on the selected routed-expert path (Algorithm 3). Both opt in to the real worker; disabling the pipeline is not the control. Neither historical `-O2` HTTP latency nor the resident TP2/vLLM supplement is a blocking/lookahead comparison. Table 10 connects current code boundaries to the observations needed for E1/E4.

Implemented boundary	Observable	Paired test / evidence gate
Immutable exact read ticket	Operation, image/layer/expert, extent order, source-verified completion	Same-cut quality; no stale, duplicate-take or foreign-image consumption
Demand wait versus next read window	Submit/start/read/verify/take times, demand-wait duration, completed/wasted bytes	Same worker/limits; compare waiting and extra speculative I/O
Owner compute submission after upload	Host submission range, upload APIs, correlated kernel/copy activities	Nsight API-to-activity correlation and clock alignment; not host-span overlap
Quota and host scope end	Packed/prepared/alias charge, metadata/stack, join/Unknown status	Same ceilings; no refund on cancellation, poison, or unproved join

Table 10: Code → observable → RQ3 test. Both schedules have qualified factory mechanisms/numerics; performance and GPU overlap remain pending. This is a test mapping, not a timing-results table.

Stopped paired study and independent corrective plan. The registered five-block order is B–L, B–L, L–B, L–B, B–L (B: blocking; L: lookahead), a fixed 3:2 order balance. Each arm plans 4 initial requests, 1 excluded QA, 2 warmups and 5 measured QA: ten fresh lifecycles, 120 requests and 50 measured QA slots. P1-blocking failed before Ready with the prepared-source/TP2-accessor clash (Appendix A). The source-plan query fix is committed at `86c58a0`, with CPU/hidden-CUDA workspace checks passed, not new CLI/HTTP qualification. One lifecycle was attempted, none completed, and nine were unissued; all 120 requests and 50 QA slots remain unissued. All five failed/incomplete pair rows stay in the analysis, with null latencies, goodput, aggregate and intervals, not zero-valued performance. There is no retry or replacement. The separate baseline pilot remains five lifecycles/60 requests/25 QA slots, unexecuted.

Corrective protocol V3 is an independent CPU-only preregistration, not a resume or repaired result. It explicitly fixes decode-cohort target/maximum deferrals to 1/0 in both arms, matching the qualified scheduler rather than the stopped CLI's implicit 4/3 defaults. Its new candidate, same-cut nine numerical gates and exact-CLI Ready/HTTP/runtime binding remain unqualified; old `033519c` passes do not qualify the corrected cut. Both arms retain representation, arithmetic, ordered workload, resets, caches, logging, CUDA proofs and the 64/8/2 MiB host/metadata/requested-stack limits. Only scheduling differs. The prospective absolute wall envelope is 7,200 s, with 650 s per lifecycle and 100 s final reserve; STOP, no retry, censoring and all-five-pairs-or-null analysis are fixed before collection.

The primary estimand is the exponentiated median of five block log ratios of arm-median measured-QA E2E; TTFT and qualified serial-window goodput are secondary. Any prospective block-resampled interval is descriptive at $n = 5$, not a significance test or power guarantee; no interval exists for the stopped study. GPU overlap additionally requires correlated Nsight CUDA API/kernel/copy activity, exact owner/bucket identities and aligned clocks. A compute host span covers submission, not execution; its intersection with a worker read proves neither overlap nor a causal scheduling benefit. Profiling/observer costs are separate, and no performance/SOTA victory, nonblocking forward or starvation-freedom claim follows.

7.5 RQ4: Custody and failure – final resource contract unknown

The current nine-gate qualification retains physical release as unobserved despite authenticated drain, stable idle and lock release. Host quota zero or matched fingerprints cannot supply a GPU retirement fence. For historical TP1 `-O2`, authenticated descendant drain reached `ECHILD`; wrapper post-idle and independent all-eight-GPU observational idle checks passed. Nevertheless `resource_contract_complete=unknown` and the supervisor raw `gpu_quiescence=unverified` remain unchanged. Invocation pending-upload/quarantine gauges of zero are not a final owner ledger. Process drain and idle telemetry establish neither device fences nor source completion nor the partition in Equation 6.

Algorithm 2 therefore specifies the required behavior, not the measured closure of this run. E5 must inject delayed/stale completion, stage-wise cancellation, primary-plus-cleanup failure, owner loss, and lost ACKs on the integrated path. Track exact publication/reuse, original versus repeated submissions, ownership transfers, retained T_a, named completion proofs, drain, and post-failure service. Reconcile grants and allocations before and after proven quiescence; distinguish deterministic protocol injection from real hardware failure. Unknown should remain charged and unusable without fabricated rollback or recovery. A failed/stalled provider also tests liveness assumptions, not just safety. Integrated fault conservation, eventual reclamation, and final TP1 custody closure remain open.

Slot	Output tok/s	TTFT ms	E2E ms	Startup s	GPU GiB/card	PSS GiB
F1	14.317	176.248	523.399	349.585	21.04	61.05
V1	10.376	132.676	694.228	66.239	19.28	7.71
V2	10.330	134.651	699.213	62.814	19.28	7.32
F2	14.804	171.333	515.106	366.133	21.04	60.14

Table 11: Descriptive resident TP2 crossblocks ($n = 2$). TTFT and E2E are request means; GPU and server-tree PSS are sampled peaks, not exact peaks or per-owner GPU attribution. F denotes Ferrule; V denotes vLLM.

7.6 Supplement: resident observations, not an RQ1–RQ4 substitute

Supplementary resident TP2 observations. The frozen route-batch crossblocks record contains two independent blocks in F1–V1 and V2–F2 order, with four fresh server owners. Each slot excludes three QA8 requests and then measures five requests, each with five input and eight output tokens: 40 output tokens per measurement window. Table 11 transcribes these existing observations; no new measurement is introduced here.

Ferrule has a higher output rate and lower mean E2E in both short windows; vLLM has lower mean TTFT, shorter startup, and lower sampled GPU and host PSS. The descriptive F/V rate ratios are 1.380 and 1.433, not statistical or causal speedups. The record uses a pinned checkpoint and tokenizer with TP2, context 128, batch 8, sequence limit 2, and concurrency 1. Ferrule's FP8-checkpoint/BF16 activation, KV and convolution/F32 SSM path uses HF-eager-qualified GemmEx/routing, whereas vLLM uses BF16/FP8 Marlin. These are not bit-exact or arithmetic-equivalent implementations, and the record retains `fairmatch=false` and `exact_speedup=null`. Ferrule's 22 GiB/device admission budget, including a 4 GiB workspace with a 3 GiB arena, is not identical to vLLM's utilization setting of 0.8. The common sampled 22 GiB ceiling is not a hard physical partition or a matched total GPU/host working-set budget.

The measured vLLM deployment is `0.28.0+cu129`, with `Torch 2.13.0+cu129`; its installed identities are frozen under `target/vendor/vllm028-py312`, not inferred from a 0.18 source checkout. Visible-text/usage canaries and Ferrule's separate scoped oracle records do not establish cross-engine internal-state parity or universal task quality. These resident windows neither demonstrate request-time physical disk reads nor establish host-memory insufficiency, storage necessity, or E0 acceptance. The immutable source record is `papers/arxiv/experiments/tp2_resident/runs/route-batch-vllm-crossblocks-20261012-01/crossblocks-results.md`, with raw metrics in the adjacent `crossblocks-results.json` and frozen identities in `sealed-identities.json`. Historical pilots and failed attempts remain separate records, not pooled repetitions of this configuration.

Supplementary startup-metadata observations. The sealed release `36e6aaa5...70a58c` combines prepared admission and shared `Arc<Path>` checkpoint read-plan metadata. Its strict TP2 S1/S2 smoke lifecycles and ordinary F1 completed with source-bound executable and mapped-library checks. S1/S2 passed visible greedy canary, reuse, cancellation-policy and post-cancel gates, with authenticated descendant drain. These are scoped serving checks, not a fresh 14-million-logit numerical validation or proof of device quiescence from process exit. Table 12 separates smoke and performance records.

F1 excludes three QA8 requests and measures five, each with five input and eight output tokens: 40 usage output tokens over 2.877869 s. Its request TTFT mean/p50/p95 are 192.944/187.513/209.700 ms, descriptive summaries of five requests, not independent-trial confidence intervals. S1, S2, and

Record	Startup s	Peak tree PSS GiB	Usage rate tok/s
36e6 S1 (smoke)	222.079	42.716	—
36e6 S2 (smoke)	139.855	44.416	—
36e6 F1 (ordinary)	140.052	43.573	13.899
e674 F1 (historical ordinary)	351.114	60.149	14.227

Table 12: Startup and sampled PSS observations. Smoke slots have no QA8 throughput measurement. The historical ordinary row is context, not a paired before/after control.

F1 are not pooled, and the best startup is not treated as a stable constant. The historical ordinary e674 F1 is a different series from the crossblocks in Table 11; no fresh BEFORE or de288 control was run. Page-cache state is uncontrolled, not cold. Startup, sampled PSS, and output-rate tradeoffs therefore do not identify a causal optimization effect, live-heap reduction, freed OS memory, or a storage-capacity result.

Five F1 proc snapshots, in first-identification, ready, post-warmup, post-measurement, and post-exit order, cost 0.643/623.616/725.908/625.392/0.085 ms to read and validate. Their recorded bounds all lie outside the measured window; receipt writes are also outside, but persistence cost was not separately measured. This excludes those reads from the stated denominator, not observer overhead from the system. On the recorded Linux 4.19 deployment, `Pss_Anon`/`Pss_File` are unavailable; the post-exit snapshot is unavailable, not zero memory.

The source-relation proof records unchanged model mathematics, backend/native kernels, and runtime, while checkpoint representation differs. CPU IO-byte, ordering, and source-check tests support compatibility, not fresh GPU numerical acceptance; full-checkpoint payload hashing was not performed. Thus metadata representation and source-bound plans have implementation evidence, while the physical request-time storage frontier remains unvalidated.

The local reproducibility entry `REPORT.md` is in:

```
target/validation/startup-metadata-owner-20261013-01
```

Its `RESULTS.json` supplies the full binary hash and raw-pin manifest; `executed-source-pins.json` identifies executed Python/config revisions. The F1 raw HTTP and `proc-timeline.json` are under the repository's paper experiment runs directory, at:

```
papers/arxiv/experiments/tp2_resident/runs/
startup-metadata-info-after-20261013-01/
runs/F1/session
```

Historical ordinary raw pins are in `routebatch-performance.json`, under the sibling directory `route-batch-info-after-20261012-01`. These working-directory artifacts still require portable archival; they have not been moved or repackaged for publication.

7.7 Evidence gates and reproducibility

The fixed TP1 reports are relative to `papers/arxiv`, under `experiments/storage_demand/runs/`:

```
storage-demand-tp1-middle-20261013-01/report.json
storage-demand-tp1-middle-20261013-02/report.json
```

The latter has adjacent `case.json`, `server.log`, `samples.jsonl`, `supervisor.json`, `custody.json`, and `post-idle.json`. The independent closeout is under `../../target/validation/`:

Question	Available observation	Gate still open
RQ1	Nine same-cut factory gates PASS; qualified logical/allocator accounting point	Capacity frontier, physical working-set attribution and CLI/HTTP quality; necessity unclaimed
RQ2	Historical HTTP plots and factory read counts; stopped-pilot client snapshots	Aligned serving/client attribution; NAS backing remains unknown
RQ3	Same-contract mechanism/numerics qualified for both schedules	CLI pilot stopped pre-Ready; corrective V3 and correlated overlap pending
RQ4	Nine authenticated drains, stable-idle v2 and local lock release	Physical release unobserved; integrated faults/final custody ledger open
TP2 supplement	Separate numerical tests; two resident crossblocks	No TP1 transfer, arithmetic parity, fair speedup, or capacity proof

Table 13: Observable-to-claim map. An open gate is not filled by a result from another representation, binary, topology, or workload.

```
storage-demand-middle-gpu-20261013-02/closure.json
```

The deterministic offline receipt analysis is in `experiments/storage_demand/analysis/output/`; it reports source consistency PASS with 41 bound inputs, 39 consumed receipt pins, and zero divergences. That is a receipt-consistency check, not a new GPU or numerical acceptance result. The TP2 numerical sources are `experiments/tp2_resident/full-boundary-audit.json` and `experiments/tp2_resident/model-runtime-quality.json`. The separate 37-call selection and its handoff are under `../../target/validation/`:

```
tp2-mixed-production-opts-01/acceptance-summary.json
tp2-mixed-optimized-release-20261009/handoff.json
```

All paths in this paragraph are relative to `papers/arxiv`.

Publication requires portable frozen artifacts: source plus dirty-state manifest, binary, model, and input-processing identities, commands and environment, raw results, traces/clock-correlation method, retained failures, and deterministic plot/table generation. Local ignored `target/` files and private transcripts are not public artifacts. Do not silently omit failed trials. Before confirmatory collection, freeze trial/block units, repetition counts, order, stopping rules, contrasts, and interval methods. Windows in one server are not independent repetitions. The stopped five-pair study and independent corrective V3 are separate from the historical external-comparison minimum and baseline pilot. Failed/incomplete pairs retain null estimates; new candidate/CLI qualification remains a gate. No statistical victory or universal quality threshold is inferred from these exploratory records.

8 Related Work and Positioning

Storage-backed and offloaded inference predate Ferrule. The comparison here is by resource representation, physical domains, quality contract, and measured traffic—not by model name or an engine's feature list.

Storage-backed inference. FlexGen [2] places tensors across GPU, CPU, and disk and optimizes throughput-oriented schedules with compression and attention-cache management. DeepSpeed ZeRO-Inference's technical report [5] describes CPU/NVMe weight hosting, layer streaming, prefetch,

and release. PowerInfer-2 [4] studies smartphone inference beyond resident capacity with neuron clusters, segmented caching, and an I/O/computation pipeline. These systems establish that storage-backed model inference and transfer/computation scheduling are prior art. Their existence also rules out a first-storage-inference claim for Ferrule.

The relevant distinction is not whether a GPU cache is small. A comparison must say whether complete compact encoded weights and metadata can remain resident, whether expanded execution weights are streamed, which physical pools are disjoint, and whether request-time reads are physical backing reads or cache-served logical reads. A configured cache budget or an expanded-weight overflow does not establish storage necessity.

Resident-domain offload and serving. PowerInfer [3] uses activation locality to place hot neurons on a GPU and compute cold neurons on a CPU. This is a GPU–CPU hybrid and should not be conflated with PowerInfer-2's storage pipeline. Host/device movement, page migration, and file-backed I/O are different observables; shared physical pools must not be added twice. vLLM's PagedAttention [1] addresses serving KV memory and is a useful resident-state comparison, but KV paging is not request-time weight materialization. The frozen vLLM 0.28 resident observations in this paper are supplementary and arithmetic/resource-budget non-equivalent, although both selected crossblock engines use PP1/TP2; they are not a speedup proof. We make no capability-absence claim about other serving systems whose primary sources and exact profiles were not audited here.

Lifecycle and scheduling comparison. The design question in RQ3 is narrower than "which system is asynchronous?" The comparison requires submission/completion separation, dependency-ready work, progress ownership, bounded queues, and completion-driven physical lifetime. An asynchronous API does not prove event-driven progress, and worker offload does not prove nonblocking execution. Measured storage/transfer/compute overlap requires correlated provider or hardware traces. The current opt-in blocking and next-expert lookahead schedules share the same worker and bounds and pass selected same-cut GPU factory qualification. The matched performance study and physical overlap remain empirical questions, not results of the numerical gates.

Positioning. Ferrule contributes a profile-scoped contract connecting source-bound plans, stage-local admission, exact generations, physical custody, publication, and retirement. This is a protocol and measurement position, not a claim that prior systems lack equivalent mechanisms; their custody and failure semantics were not exhaustively audited here. The concrete difference remains an empirical question: whether a selected integrated profile can complete a quality-accepted workload under declared per-domain working-set limits, what physical traffic it causes, and whether its safety contracts survive faults. Registry coverage, OS-cache hard caps, durable ACID, cross-rank atomicity, NCCL/device-collective defaults, multimodal integration, and performance leadership are not implied by the design.

9 Limitations and Remaining Evidence Gates

Bounded accounting and physical attribution. Configured budgets cover selected ledgers, not the complete process/machine working set. Packed/prepared payload quota, metadata allowance and requested worker stack are bounded on the implemented path, but OS/TLS/allocator headroom and weak-reference-retained allocation backing are not physical RSS bounds. Device/PSS/RSS samples can miss transients and lack complete allocation attribution. Physical-memory insufficiency

Prior work	Resource/backing focus	Scheduling or state focus
FlexGen	GPU, CPU, disk	Tensor placement, batching, compression
ZeRO-Inference	CPU or NVMe to GPU	Layer streaming, release, prefetch
PowerInfer	GPU–CPU hybrid	Hot/cold neuron execution
PowerInfer-2	Smartphone memory/storage	Neuron clusters, cache, I/O pipeline
vLLM	Resident serving KV memory	PagedAttention

Table 14: Source-described comparison axes, not matched measurements or an exhaustive protocol audit. Unassessed guarantees are not marked unsupported.

and storage necessity are not claims of this scoped paper; they would additionally require compact-residency, page-cache/swap, co-tenant and per-pool evidence. Shared physical pools count once, and discrete CUDA records do not qualify shared/unified-memory deployment.

Logical demand and same-cut quality. Historical TP1 `-O2` records repeated application-level demand during four HTTP successes in one DEBUG run. Its 36,240,064,128 reader bytes are not physical NAS or backing-block traffic, and its plots are not current read- pipeline statistics. Three visible-text/usage reference canaries passed; Hello2 was HTTP/usage only, generated IDs were absent, and complete numerical quality was not observed. Current `033519c` factory qualification now passes all nine GPU gates, including both schedules' diagnostic, ordinary-fingerprint and current-primary cold checks. Ordinary routes/states remain unobserved, and stream denominators are not pooled. Historical initialization and strict-wrapper failures remain failures in Appendix A. Neither factory results nor common CLI source/runtime provenance establish CLI/HTTP numerical equivalence. Observer/allocator-getter cost is unmeasured, TP2 does not fill TP1 gates, and argmax matches do not waive W8A16 strict failure.

Scheduling, integration, and progress. Both opt-in blocking and bounded lookahead are implemented, with the same read worker and host/device budgets. Whole forwards remain synchronous; blocking syscalls, consumer waits, locks and synchronization remain possible. The read window is split from owner compute submission, not from every consumer completion. Nsight-correlated kernel/copy evidence and a completed same-cut paired study are still needed for GPU overlap or a scheduling benefit. Host-span intersections and a worker thread establish neither. The stopped five-block study supplies no estimates; corrective V3 remains unqualified, not a retry or a power guarantee. The nine gates provide one qualified factory/accounting point, not repeated performance data or a capacity frontier. Logical host high-water is not RSS; allocator high-water may be observed through a fence-polling getter, while passive snapshots leave allocator state unobserved.

The lifecycle algorithms describe named boundaries, not formal verification or universal registry integration. The prepared-KV, HostKV and specialized expert paths have distinct owners. Passive observations neither heal poison nor discharge custody. Host cancel/drain/join cannot prove GPU quiescence; Unknown may permanently reduce available capacity. Final historical TP1 resource closure remains unknown despite process drain and idle observations. No general reset/replay, device-reset recovery, bounded recovery time, availability, durable ACID or arbitrary cross-rank atomicity is provided. Mixed collectives retain their scoped host-default boundary. Other modalities and shared/unified memory need independent implementation and quality gates.

Comparisons and reproducibility. The two resident TP2 crossblocks are descriptive tradeoffs, not a fair, causal or statistically confirmed speedup. Arithmetic and total budgets differ; startup

observations have uncontrolled caches and unpaired context. No state-of-the-art or first-storage-inference claim is made. Submission requires portable source/binary/model/oracle identities, complete failure records, deterministic tables/plots and final archival/visual checks. A scoped logical/NFS-client report may retain unknown NAS backing traffic without claiming storage necessity; it may not invent that missing evidence.

10 Conclusion

Ferrule connects bounded storage-backed weight residency to source-bound planning, covered admission, exact completion, authorized publication, and custody through physical retirement. The concrete next-expert pipeline provides a same-worker blocking control and one bounded lookahead read after current upload and before compute submission. Its three-artifact plans, alias-aware host quota and sticky Unknown handling make the scheduling boundary and safe reuse conditions explicit without changing the synchronous forward.

The frozen `033519c` factory/profile now passes all nine exact-once GPU qualification gates for constructor/provider and both schedules' full40, fresh, ordinary and cold-owner paths. Ordinary fingerprints match their own schedule's diagnostic primary; streams and unobserved route/state fields keep their distinct scope. This is a qualified numerical/provider and accounting point, not CLI/HTTP equivalence, a capacity frontier or physical release. Aligned client I/O, integrated fault/custody evidence and the five-pair RQ3 performance study remain open after the pre-Ready CLI failure. GPU overlap requires correlated activity, not host submission spans; physical-memory necessity, general nonblocking execution and performance leadership require stronger separate evidence.

References

[1] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient Memory Management for Large Language Model Serving with PagedAttention, 2023. arXiv:2309.06180.

[2] Ying Sheng, Lianmin Zheng, Binhang Yuan, Zhuohan Li, Max Ryabinin, Daniel Y. Fu, Zhiqiang Xie, Beidi Chen, Clark Barrett, Joseph E. Gonzalez, Percy Liang, Christopher Ré, Ion Stoica, and Ce Zhang. FlexGen: High-Throughput Generative Inference of Large Language Models with a Single GPU, 2023. arXiv:2303.06865.

[3] Yixin Song, Zeyu Mi, Haotong Xie, and Haibo Chen. PowerInfer: Fast Large Language Model Serving with a Consumer-grade GPU, 2023. arXiv:2312.12456.

[4] Zhenliang Xue, Yixin Song, Zeyu Mi, Xinrui Zheng, Yubin Xia, and Haibo Chen. PowerInfer-2: Fast Large Language Model Inference on a Smartphone, 2024. arXiv:2406.06282.

[5] ZeRO-Inference: Democratizing massive model inference. DeepSpeed technical web page, September 2022. September 9, 2022; technical web page, not a conference paper.

A Qualification History and Stopped CLI Study

Table 15 separates historical test numerics from wrapper acceptance, not forecasts for enabled CLI serving. Full40 covers Hello/Capital prefill and one decode; fresh adds bounded history, feedback, reset and fork. Fixed atol = 0.00025, rtol = 0.012 is not bit-exactness or broad task quality.

Record	Metrics	Logit elements	Numerical / wrapper outcome
Middle quality `-01`	0	0	`cuInit (3)` before inference; fresh/cold unissued
Full40 `-02`	804	1,986,560	Numerical PASS; original strict post-idle wrapper FAIL
Retained fresh `-03`	5,427	9,187,840	Numerical PASS; original strict post-idle wrapper FAIL
Cold owner `-04`	1,809	3,228,160	Numerical and wrapper PASS under stable-idle v2; different source cut
Frozen CLI quality `-01`	0	0	`cuInit (3)` before inference; retained fresh unissued

Table 15: Later 2026-10-14 TP1 test-gate chronology. Middle quality `-01` and frozen CLI quality `-01` are distinct from the 2026-10-13 prelaunch `DEVNULL` wrapper failure. Numerical PASS does not override a rejected wrapper.

The strict full40/fresh wrappers rejected first post-idle samples of 1%/6% GPU0 utilization (zero MiB, no compute applications); later idle does not repair them. Cold `-04` accepts authenticated drain/stable-idle v2 on the 3,533-input `5662b7d2...` cut, not `fbf0c225...`. Its old primary records `same_source_cut=false` and wrapper rejection; these passes cannot be pooled. Frozen CLI quality fails at `cuInit (3)` with zero metrics, logits and token checks, not numerical mismatch. The later nine-gate `033519c` factory PASS changes none of these outcomes. Idle is not a GPU fence or final resource ledger.

Enabled CLI pilot: permanently stopped. On 2026-10-15, P1-blocking exits before Ready with:

```
invalid runtime request:  experimental qwen35-mixed-bf16 serving:
prepared source profile changed
```

The stopped `cmd_serve` called `mixed_tp2_cache_requirements` unconditionally; the non-TP2 accessor rejected enabled TP1's prepared-source payload before CUDA owner startup, not at `cuInit`. The source-plan query fix is now committed at `86c58a0`; this failure is unchanged. One lifecycle is attempted, zero complete, nine unissued; HTTP issued, terminal, success, failed and censored counts are all zero, with all 120 requests/50 measured QA slots unissued. All five incomplete pairs stay; quality metrics, request latencies, goodput, contrasts and intervals are unobserved/null, not zero. Ready maps, positive reads and the final owner ledger are unobserved/unknown despite authenticated `ECHILD`, absent owners, idle and released/reprobed locks. Shutdown-only shared NFS deltas are not request-attributable NAS traffic; physical release remains unproved.

Frozen `paired-final.json` and `closeout.json` are relative to `papers/arxiv`, under:

```
experiments/storage_demand/paired_read_v1/
artifacts-current-20261015-01/
```

No retry/replacement is allowed. Independent V3 fixes cohort target/deferrals to 1/0, not CLI defaults 4/3. CPU and hidden-CUDA workspace checks pass for the fix, not new CLI/HTTP qualification. New same-cut GPU gates and serving binding remain required; V3 does not resume this study or borrow the old factory PASS.