Completion evidence governs custody after a real routed-down fault

Mixed CUDA component · 4/4 cells PASS (2 selectors × 2 schedules) · width 16, GDN/GQA, 4 experts/top-3 · BF16; F32 SSM

a Transaction path (logical stages, not elapsed time)

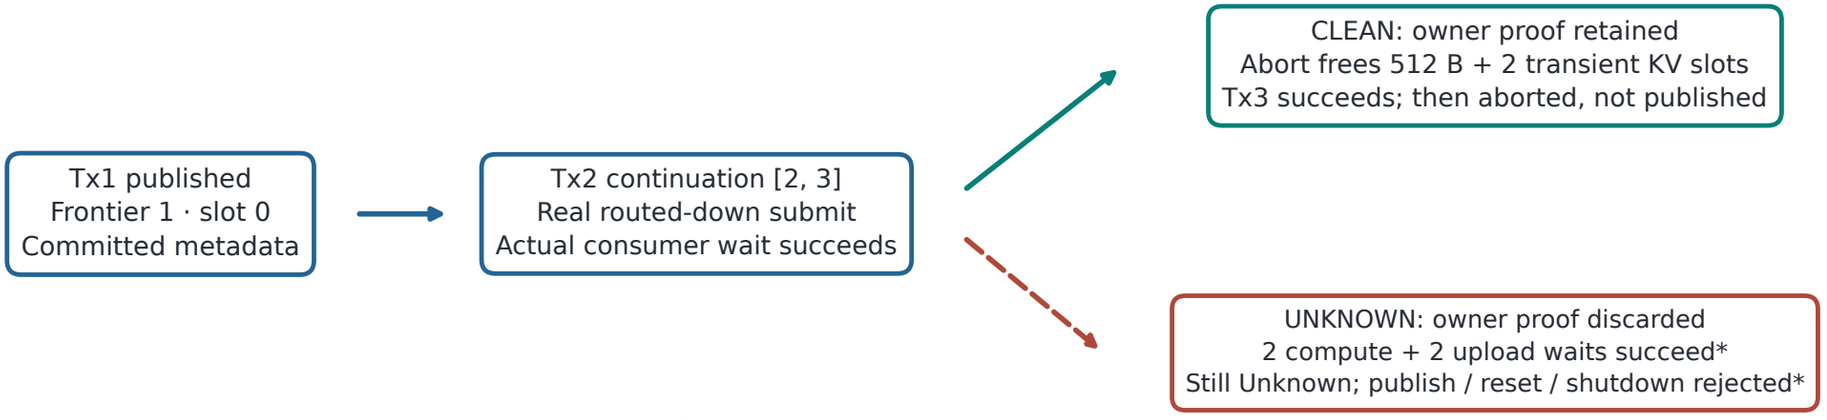

* Source + action log + PASS back these outcomes; no individual fence/error receipts were serialized.

b Retained sequence-state bytes

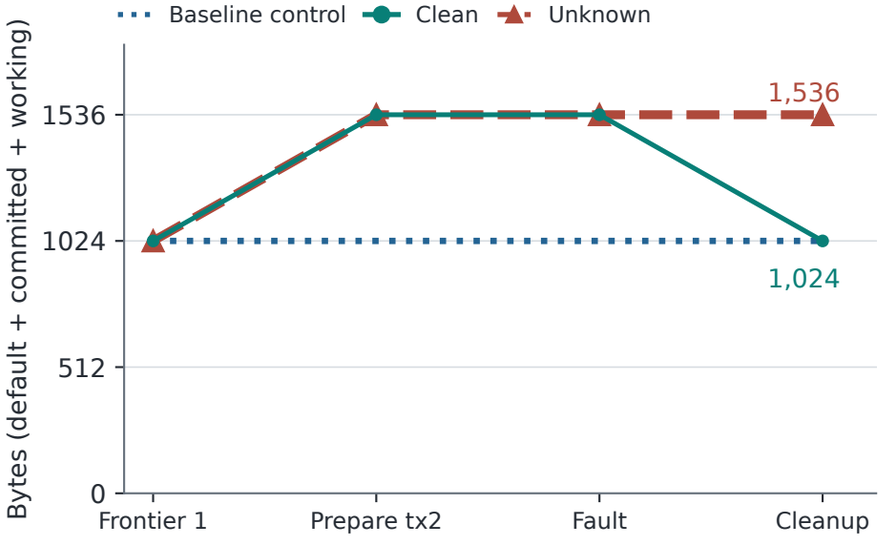

Clean: no active transaction; 31 free KV slots.
Clean KV / conv / SSM bitwise unchanged*.
Unknown: 1 active transaction; 29 free slots.
Both schedules agree; committed metadata stays at frontier 1.

c Cache budget after cleanup

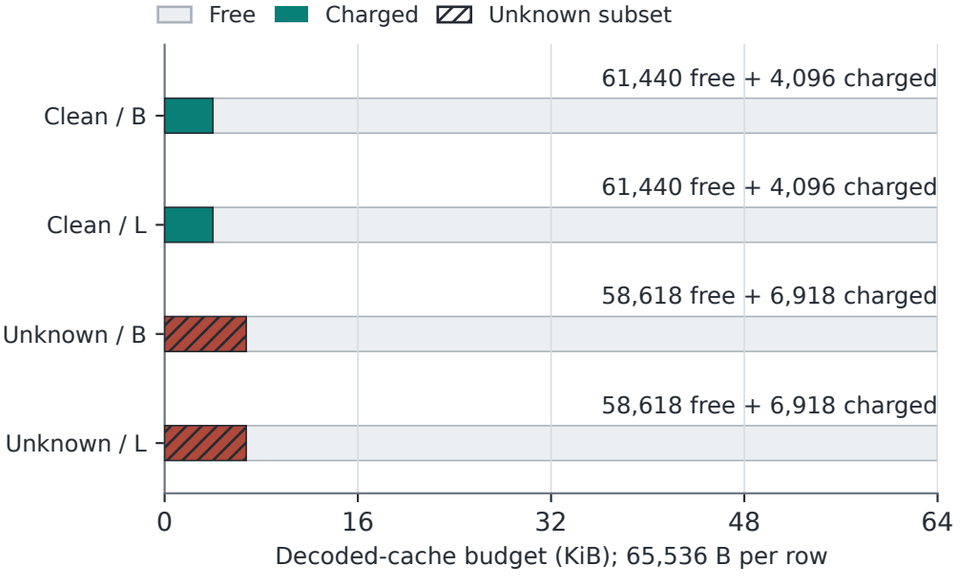

Hatch: 6,918 B Unknown is a subset, never added again.
Unknown = 774 pending + 6,144 scratch; clean = 4,096 scratch.
B = Blocking; L = Lookahead. Free + charged = 65,536 B.

Host reads close independently: charge = 0, host Unknown = 0 in all four cells.
Lookahead wastes 1 read (774 B) per cell; Blocking wastes 0. Host cleanup does not release GPU custody.

Post-Unknown KV / SSM CONTENTS UNOBSERVED — only committed metadata is unchanged.

The 318,720 B recurrent workspace is a separate, overlapping owner view: no additive physical-memory total.
Clean-only explicit owner drop + allocator trim returns fixture requested bytes 0 → 0; no device-wide release certificate.
Component contract only: not integrated broker / CLI E5, a 35B model fault test, or physical driver proof loss.