Lower read-demand wait accompanies lower measured E2E

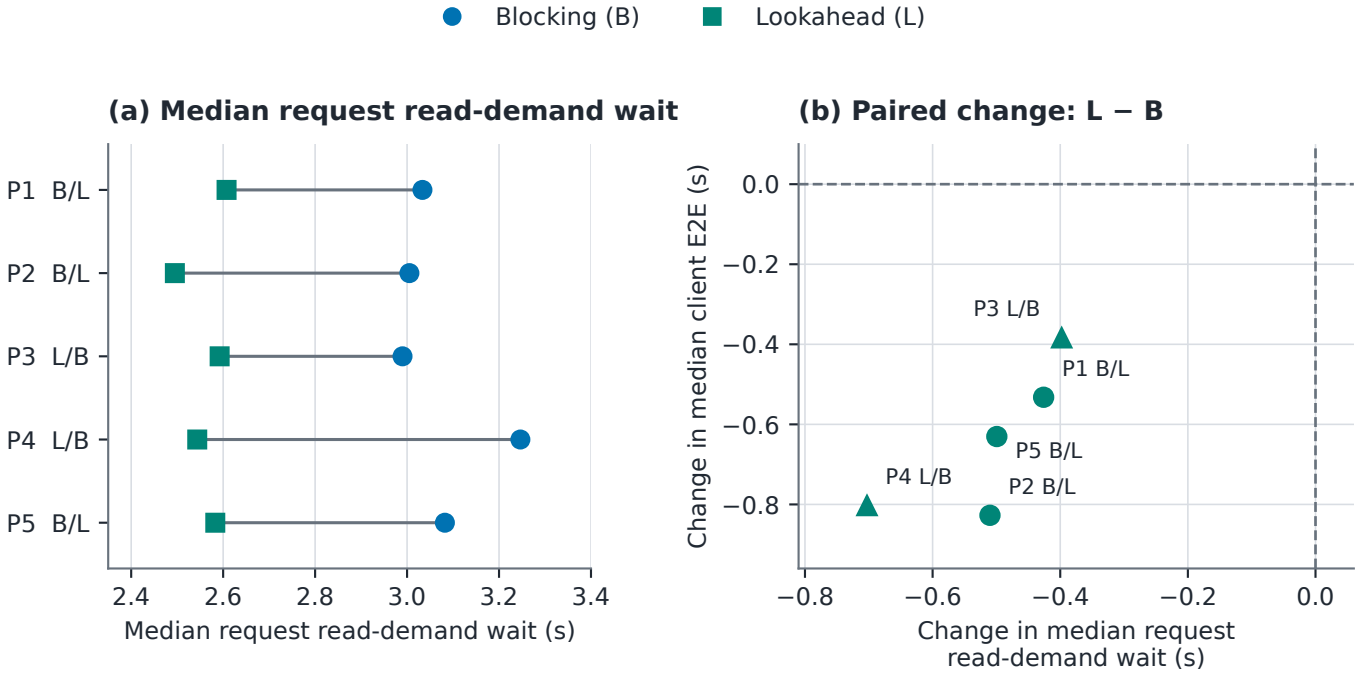

Each point is one paired block; each arm is the median of five measured-QA requests. No fit or causal coefficient.
Host demand-wait intervals are not CUDA overlap timelines and are not additive to client latency.
Association is consistent with less waiting, but does not establish critical-path causality; shared-cache conditions remain.